

Fundamental Data Structures, Algorithms and Complexity, and Computer Architecture

What are CSC211 and its partner CSC212 about?

- CSC211 concentrates on *problem solving* that leads to correct programs whose timing and memory requirements are well understood. Students *develop programs* to prove they understand the theory. They *time* them and try to isolate what is slow, and *speed up their code*. Students learn how to validate and *prove that their code is correct*. Students learn many *problem solving approaches*.
- CSC212 has two parts— (1) It continues the themes of problem solving and programming in a high level language and understanding the time and space complexity of algorithms. Students continue to *broaden their problem solving skills*.
(2) Hardware is covered and we learn some assembler programming to understand clearly what the underlying hardware does when it computes a loop or executes a function recursively. Different styles of computer architecture are compared and students gain an understanding of the interaction between high-level and low-level programming and hardware.

Fundamental Data Structures, Algorithms and Complexity

- This course concentrates on *problem solving* that leads to correct efficient programs.
- We will often first try *formulating problems* in English or pseudo code and then proceed to *code* them in Python, Java, C++, or even C# or whatever language you like the most.
- Representing data on computers is done using binary digits—called bits. We learn how eight, sixteen or more bits can be used to represent individual characters. How 32 or 64 bits can be used to represent integers and how to represent floating point numbers.
- We will revise some of the important ideas you should already have mastered.
- We learn data structures such as linear lists, queues, stacks, hash tables, binary search trees, priority queues.

Fundamental Data Structures

- We learn algorithms to store and retrieve data in dictionaries, trees.
- We study efficiency of algorithms in terms of their time and space complexity.
- Algorithms are implemented using Java or Python.
- You are, however, expected to *become an expert in Java programming*.
- An integral part of this course consists of your programming effort.
- Three hours are scheduled per week for programming in the Sun lab.
- Most practicals must be completed on the same day they are issued.
- Marks contributing to your semester mark are awarded for original work submitted. **Unoriginal submitted work gets negative marks.**

Fundamental Data Structures—prescribed books

Class notes for 2015

These slides form the core part of this course COS211/212.¹

Prescribed book for 2015

Algorithms FOURTH EDITION, by Robert Sedgewick and Kevin Wayne (2011), Addison-Wesley.

Books for reference in 2015

Data Structures and Algorithms in Java, 4th Edition, by Michael T. Goodrich and Roberto Tamassia (2006), John Wiley and Sons. Inc. The 5th edition is available.

Introduction to Algorithms, 3rd edition, Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein, MIT Press, 2009. ISBN 978-0-262-53305-8—paperback. This is a daunting book that scares most people off but is the best. Its second edition is freely available on the Internet.

¹The class notes can be found on the Sun systems in the directory /export/home/notes/ds/ in the files ds.ps, or ds.pdf—or 4-to-a-page: 4up-ds.ps, 4up-ds.pdf. Older versions of the notes are called ds-slides.*. Read the notes.

Fundamental Data Structures

For *Java programming* we recommend:
Savitch: the book you used last year.

Algorithms FOURTH EDITION, by Robert Sedgewick and Kevin Wayne is packed with examples of objects and methods in Java.

Chapter 1 of *Algorithms* FOURTH EDITION, by Robert Sedgewick and Kevin Wayne is available free on the Internet.

It is essential to have your own copy of the book.

For *Python* there are many good, free books available on the Internet. Stick to Python 3 and learn the differences between 2 and 3.

Fundamental Data Structures

Important *prerequisites* for this course are

- Proficiency in both Java and Python is expected before you enter this course. This is not a course about programming in Java or Python.
- Objects in Java.
- Translate pseudocode into runnable code
- Logarithms, summation of series, differential and integral calculus, linear algebra. Maths I is a prerequisite to this course.

At the conclusion of this course the aims are that you are expected:

- to program all the algorithms in this course proficiently in Java or in pseudo code.
- to explain any algorithm in this course coherently.
- to be able to understand, and discuss the time and space efficiency of each algorithm in this course.
- to know the problem-solving paradigms such as *divide-and-conquer, dynamic programming, brute force, breadth-first, depth-first or top-down design, recursion versus iteration, inductive evolution of programs, etc.*

An introductory example—Shuffling Cards

- We start by considering how to shuffle a deck of cards.
- The abstraction for the *unshuffled* deck of cards is the set

$\{1, 2, \dots, 52\}$, which is often written as $[1..52]$.

This can be built in Python using a list called **deck**.

```
1 deck = [0]*52
2 for card in range (1,53):
3     deck[card-1] = card
```

or more stylishly with

```
1 deck = [card for card in range (1,53)]
```

Notice that there is no extra unused zero at the start of the array.

TRY IT NOW.

An introductory example—Shuffling Cards

- An alternate abstraction to $\{1, 2, \dots, 52\}$, or $[1..52]$ is

$$\{0, 1, \dots, 51\}, \text{ or } [0..52) \equiv [0..51].$$

This is initialized very naturally in Python using

```
1 deck = [0]*52
2 for card in range (52):
3     deck[card] = card
```

or more stylishly with

```
1 deck = [card for card in range (52)]
```

It is hard to dictate which is preferable. The first way seems easy because we learn to count from 1 but the second way $[0..52)$ comes naturally to programmers, mathematicians and computer scientists.

Learn to use both abstractions.

Shuffling Cards—setting up the deck in Java

- In Java we need to declare the array `deck[]` with

```
1 int deck[] = new int [53];
```

and then initialize the array to $\{1, 2, \dots, 52\}$ with:

```
1 for (card = 1; card <= 52; card++)  
2     deck[card] = card;
```

- The element `deck[0]` is unused.
- We need another array to put our “shuffled” cards into.
- Create the array `shuffledDeck[]` in Java.

```
1 int shuffledDeck[] = new int [53];
```

A shuffled deck can be built by drawing cards one-by-one randomly and stacking them onto the ‘shuffled’ deck.

Shuffling Cards—in Java

- Set $k = 1$. A card is ‘drawn’ by randomly drawing from $\{1, 2, \dots, 52\}$ and put into the k th position of the shuffled deck as follows:

```
1  r = (int)(Math.random()*52) + 1;  
2  shuffledDeck[k++] = r;
```

- This *does not work*—duplicates are put into the shuffled deck. The solution is either:
 - (1) to deal with the duplicates, otherwise
 - (2) never generate a duplicate. *How do we do this?*

Deal with the duplicates

- (1) If the new card is already in the list then discard it and try again.

Never generate a duplicate

- (2) Maintain a list that contains numbers that have never been selected and then draw cards one-by-one from this list.

We will attempt both approaches.

Shuffling—make bad method work

- We must insert r into `shuffledDeck` if and only if r is not there.

```
1 k = 0;
2 while (k < 52) {
3     r = (int)(Math.random()*52) + 1;
4     for (j=1; j<=k; j++)
5         if (shuffledDeck[j] == r) break;
6     if (j == k+1)
7         shuffledDeck[k++] = r; }
```

- When r never matches any of the k values, then j becomes $k+1$, the `for` loop is exited and Line 7 is executed.
- Notice that $k < 52$. Once all the values in the array `shuffledDeck[1..51]` are filled in, getting the 52nd one by letting the loop run to $k = 52$ is very time consuming.
How long it takes on average?

Shuffling—entering the last item.

To insert the last value we need to compare the generated random number with each of the 51 numbers already in the list. This must be done until the missing number is found. On average we need to generate 52 numbers before we find the missing one.

$$\text{Since } 0 + 1 + 2 + 3 + \dots + 51 = \frac{51 \times 52}{2} = 1326,$$

the following strategy can be followed. Keep a running sum called **total** of each new card number as it is found. The value of the missing card is

$$1326 - \text{total}.$$

On average this calculation saves computing random numbers and comparing about 52 of them, 51 times.

There are some other ideas to improve matters.

Shuffling—make bad method work


```
1 k = 1;
2 while (k < 52) {
3     r = (int)(Math.random()*52) + 1;
4     for (j=1; j<=k; j++)
5         if (shuffledDeck[j] == r) break;
6     if (j == k+1)
7         shuffledDeck[k++] = r; }
```

- At first the shuffled deck starts filling rapidly, but as it gets more cards—values of r —matters slow down.
- How fast is this method?
- —It is slow.
- We will see how slow after looking at the same method in Python.

Shuffling—inserting into a shuffled deck in Python all the way


```
1 from random import random
2
3 def shuffle(N):
4     shuffledDeck = [0]*N
5     k = 0
6     while k < N:
7         r = int(random()*N)
8         for j in range (k+1):
9             if shuffledDeck[j] == r: break
10        if j == k:
11            shuffledDeck[k] = r
12            k += 1
13
14    return shuffledDeck
15
16 print (shuffle(52))
```

Shuffling—average number of probes




- Assume that there are already k cards in the list. The fraction $p = \frac{k}{n}$ contains the already shuffled cards in their final positions.
- The rest of the list, the fraction $1 - p$ of the list is open. Here $n = 52$.
- The $(k + 1)$ th card must now be randomly drawn by generating a random number $r \in [1..52]$.
- How long does it take to draw a valid card $= r$ that can be added to the list?

Shuffling—average number of probes



- Card = r is only valid if it does *not* appear in the list `shuffledDeck[1..k]`
- For $p = \frac{k}{52}$ of the time the card = r is already present in the first k .
- *How do we discover that a card has already been selected?*
- One method is to search through the list until it is found.

Shuffling—checking for the presence of $\mathbf{r}$



- Remember that we have assumed that we will find $\mathbf{r}$ in the list,
- In this case about half of the list of $\mathbf{k}$ elements must be searched to discover $\mathbf{r}$, so $\frac{k+1}{2}$ probes are done on average, for p of the time:

$$p^{\frac{k+1}{2}}$$

Shuffling—when r is not in the list

1	17
	7
	19
	5
	52
	⋮
	⋮
k	27
n	

$$p = \frac{k}{n}$$

$$1 - p = \frac{n-k}{n}$$

- For the rest of the time the whole list—now of length k —is searched, i.e. the r is compared with every one of the k items in the list, and is unequal to any of them. So k comparisons are made.
- This occurs when a new r is drawn that is not yet in the list.
- How often is there a new r , i.e. an r that is not in the first k elements?

$$\left(1 - \frac{k}{52}\right)k = (1 - p)k.$$

Shuffling—when r is not in the list

- This must be summed for all 52 cards:

Number of probes =

$$\sum_{k=1}^{52} p \frac{(k+1)}{2} + (1-p)k$$

- The first term is not quite correct, since if a card is already in the list we must search *again and again* until one is found that is not in the list, and that process may have to be repeated many times.
- Let us analyse this.

Shuffling—total number of probes

- So, how long does it take to draw the k th card?
- For $p = \frac{k}{52}$ of the time the card is already in the deck, so when a card *is not in the list*, then the entire list—currently of length k —must be searched, i.e. $(1 - p)k$.
- Suppose the generated card *is* in the list. It will be found on average after $\frac{k+1}{2}$ probes.
- But, if it *is* in the list another card must be generated, and this process must be repeated until finding a card not in the list.
- 1st probe costs $p \frac{k+1}{2}$,
- 2nd probe costs $(p + p^2) \frac{k+1}{2}$,
- 3rd probe costs $(p + p^2 + p^3) \frac{k+1}{2}$,

Shuffling—total number of probes

- 4th probe costs $(p + p^2 + p^3 + p^4) \frac{k+1}{2}$, etc.

- m th probe costs $(p + p^2 + \dots + p^m) \frac{k+1}{2}$

Probing forever guarantees card is found, so

Probe time = $\left[\sum_{i=1}^{\infty} p^i \right] \frac{k+1}{2}$, but putting

$$S_n = \sum_{i=1}^n p^i = p + p^2 + p^3 + \dots + p^n \text{ then}$$

$$pS_n = p^2 + p^3 + \dots + p^n + p^{n+1}$$

$$(1 - p)S_n = p - p^{n+1} \text{ giving } S_n = \frac{p - p^{n+1}}{1 - p}$$

and since $p = \frac{k}{n} < 1$, because $k < n$, it follows that

$\lim_{n \rightarrow \infty} p^{n+1} = 0$ and therefore

$$S_{\infty} = \sum_{i=1}^{\infty} p^i = \lim_{i \rightarrow \infty} S_n = \frac{p}{1 - p}$$

- Probe time = $\left[\frac{p}{1 - p} \right] \frac{k+1}{2}$ and

Shuffling Cards—improving the method

- Probe time $= \lceil \frac{p}{1-p} \rceil \frac{k+1}{2}$ and
Total number of probes, or $= \sum_{k=1}^{51} \frac{p}{1-p} \frac{(k+1)}{2} + (1-p)k$
- Why do we stop summing at 51?
- The time complexity is:
 $= \sum_{k=1}^{51} \frac{p}{1-p} \frac{(k+1)}{2} + (1-p)k$
- Improve by avoiding $\frac{p}{1-p} \frac{(k+1)}{2}$
- It mounts up to about 90% of the time.

Shuffling Cards—improving the method

- The time complexity is:

$$= \sum_{k=1}^{51} \frac{p}{1-p} \frac{(k+1)}{2} + (1-p)k. \text{ Improve by avoiding } \frac{p}{1-p} \frac{(k+1)}{2}$$

- It mounts up to about 90% of the time.
- The trick is to store a bit array—incorrectly called a bit vector—that indicates which cards have been generated:

```
1 boolean[]notInList = new boolean[53];
2   for (k=1; k<=52; k++)
3       notInList[k] = true;
4   while (k < 52) {
5       r = (int)(Math.random()*52) + 1;
6       if (notInList[r]) {
7           shuffledDeck[k++] = r;
8           notInList[k] = false; } }
```

Shuffling Cards—improving the method

- The time complexity is:

$$\sum_{k=1}^{51} (1-p)^k$$

- This is about 10% of the time taken by the original bad method.
- There is an efficient method that takes $O(n)$ time.

Shuffling Cards—much improved method

- Start with a deck of $N = 52$ cards, represented by an array of integers $\{1, 2, \dots, N\}$. Suppose the array is initialized using:

```
1 for (card = 1; card <= N; card++)  
2   deck[card] = card;
```

- Next, generate exactly N uniform random numbers drawn from $\{1, 2, \dots, N\}$.
- We hope that few duplicates do not affect the result significantly.

```
1 for (card = 1; card <= N; card++) {  
2   r = (int)(Math.random()*N) + 1;  
3   swap(deck[i], deck[r]); }
```

- There are some surprises. As in card games the deck is unshuffled only when it is new. Starting with a new deck each time will yield a biased shuffle. This becomes obvious when the deck $\{1, 2, 3\}$ is shuffled by hand.

Shuffling three cards—an experiment

- Start with the cards in order: $\{1, 2, 3\}$. What are the outcomes?
- Suppose that `card` = 1. Get a random number $r \in \{1, 2, 3\}$ and swap `deck[card]` with `deck[r]`. This yields three possibilities.

r	The deck	Effect
1	$\{1, 2, 3\}$	1 is swapped with 1
2	$\{2, 1, 3\}$	2 is swapped with 1
3	$\{3, 2, 1\}$	3 is swapped with 1

- Next, let `card` = 2. Get a random number $r \in \{1, 2, 3\}$ and swap `deck[card]` with `deck[r]`. This yields three possibilities for each of the above configurations—a total of nine possibilities.

The deck	r	Result of swap
$\{1, 2, 3\}$	1	$\{2, 1, 3\}$
$\{1, 2, 3\}$	2	$\{1, 2, 3\}$
$\{1, 2, 3\}$	3	$\{1, 3, 2\}$
$\{2, 1, 3\}$	1	$\{1, 2, 3\}$
$\{2, 1, 3\}$	2	$\{2, 1, 3\}$
$\{2, 1, 3\}$	3	$\{2, 3, 1\}$
$\{3, 2, 1\}$	1	$\{2, 3, 1\}$
$\{3, 2, 1\}$	2	$\{3, 2, 1\}$
$\{3, 2, 1\}$	3	$\{3, 1, 2\}$


- Starting with these nine orderings we next consider swapping where `card` = 3 `deck[card]` with `deck[r]` for $r \in \{1, 2, 3\}$.

Shuffling three cards—an experiment

Deck	r	After final swap
{2, 1, 3}	1	{3, 1, 2}
{2, 1, 3}	2	{2, 3, 1}
{2, 1, 3}	3	{2, 1, 3}
{1, 2, 3}	1	{3, 2, 1}
{1, 2, 3}	2	{1, 3, 2}
{1, 2, 3}	3	{1, 2, 3}
{1, 3, 2}	1	{2, 3, 1}
{1, 3, 2}	2	{1, 2, 3}
{1, 3, 2}	3	{1, 3, 2}
{1, 2, 3}	1	{3, 2, 1}
{1, 2, 3}	2	{1, 3, 2}
{1, 2, 3}	3	{1, 2, 3}
{2, 1, 3}	1	{3, 1, 2}
{2, 1, 3}	2	{2, 3, 1}
{2, 1, 3}	3	{2, 1, 3}
{2, 3, 1}	1	{1, 3, 2}
{2, 3, 1}	2	{2, 1, 3}
{2, 3, 1}	3	{2, 3, 1}
{2, 3, 1}	1	{1, 3, 2}
{2, 3, 1}	2	{2, 1, 3}
{2, 3, 1}	3	{2, 3, 1}
{3, 2, 1}	1	{1, 2, 3}
{3, 2, 1}	2	{3, 1, 2}
{3, 2, 1}	3	{3, 2, 1}
{3, 1, 2}	1	{2, 1, 3}
{3, 1, 2}	2	{3, 2, 1}
{3, 1, 2}	3	{3, 1, 2}

Permutation	Number of occurrences	Random trial of 10^7 biased shuffles	Random trial of 10^7 <i>unbiased</i> shuffles
{1, 2, 3}	4	8890048	10001750
{1, 3, 2}	5	11116476	10003651
{2, 1, 3}	5	11104844	9997698
{2, 3, 1}	5	11111553	9995874
{3, 1, 2}	4	8885890	10000476
{3, 2, 1}	4	8891189	10000551

The bias is quite clear and confirms the distribution of shuffles when r is chosen randomly. The histogram makes the bias very clear.



Shuffling three cards—shuffling bias

It is not surprising that this bias occurs. We have just shown that when starting from the unshuffled deck $\{1, 2, 3\}$ it turns out that there are $3^3 = 27$ ways of reaching one of six possible outcomes. This means that each of the six possibilities can be reached in an average of 4.5 ways. so it is reached in either 4 or 5 possible ways.

If we randomly start with any other ordering a similar distribution ensues.


Since the pattern of outcomes is always 4, 5, 5, 5, 4, 4 the distribution quickly evens out when the shuffle starts at the current ordering of the deck. Let us call the probability of getting the ordering s_k when starting with s_1 , $P_1(s_k)$
 $P_1(s_1) = 4/27$, $P_1(s_2) = 5/27$ $P_1(s_3) = 5/27$ $P_1(s_4) = 5/27$ $P_1(s_5) = 4/27$,
 $P_1(s_6) = 4/27$. Starting at ordering s_2 yields the same probabilities tied to other orderings, namely $P_2(s_1) = 5/27$, $P_2(s_2) = 4/27$ $P_2(s_3) = 4/27$
 $P_2(s_4) = 4/27$ $P_2(s_5) = 5/27$, $P_2(s_6) = 5/27$.

Shuffling three cards—shuffling bias

Call the orderings $s_1 = \{1, 2, 3\}$, $s_2 = \{1, 3, 2\}$, $s_3 = \{2, 1, 3\}$, $s_4 = \{2, 3, 1\}$, $s_5 = \{3, 1, 2\}$, $s_6 = \{3, 2, 1\}$. Let $P(s_k)$ for $k \in [1..6]$ be the probability of getting ordering s_k . The table below shows the distribution of the expected orderings.

Starting at ordering	$P(s_1)$	$P(s_2)$	$P(s_3)$	$P(s_4)$	$P(s_5)$	$P(s_6)$
s_1	4/27	5/27	5/27	5/27	4/27	4/27
s_2	5/27	4/27	4/27	4/27	5/27	5/27
s_3	5/27	5/27	4/27	5/27	4/27	4/27
s_4	4/27	4/27	5/27	4/27	5/27	5/27
s_5	5/27	5/27	4/27	4/27	4/27	5/27
s_6	4/27	4/27	5/27	5/27	5/27	4/27

It can be seen that the expected probability of getting any ordering is equal to that of any other as is evidenced by the histogram.



Shuffling Cards—much improved method—complexity

- The timing is simple: if there are n cards, there are n operations;
- each operation takes $O(1) + O(1) = O(2)$, but since 2 is a constant that is independent of n , we write $O(2)$ as $O(1)$.
- But the expression $n \times O(1)$ is dependent on n , so $n \times O(1) = O(n)$
- so it is of the order of n and is written as $O(n)$.
- It is much faster than the previous two methods.
- The bad news is that **it does not generate the ideal shuffle.**
- If you use it in your online poker game you are going to lose money.
- We will first show why sorting three card cards with this method yields a biased result.
- We then illustrate experimentally how badly it is biased.

Shuffling Cards—by drawing randomly from $\{1, 2, \dots, N\}$.

How does the ideal shuffle work?

Suppose that $N = 3$. There are $N! = 6$ ways to shuffle 3 cards.

The six possible orderings are

$\{1, 2, 3\}, \{1, 3, 2\}, \{2, 1, 3\}, \{2, 3, 1\}, \{3, 2, 1\}$ and $\{3, 1, 2\}$.

Randomly draw a card from the deck $\{1, 2, 3\}$, i.e., draw 1, 2 or 3—there are three possibilities—leaving two cards.

Next we draw another card—this is done in two possible ways.

When this card is drawn there is no further randomization.

There are exactly six— $3!$ —ways of doing the shuffle.

Shuffling Cards—ideal shuffle

- Start with a deck cards, this time represented by an array of integers $\{0, 2, \dots, 51\}$. Initialize the array using:

```
1 for (card = 0; card < 52; card++)  
2   deck[card] = card;
```

- Next, generate exactly 51 uniform random numbers drawn from $\{0, 1, 2, \dots, 51\}$.
- Duplicates are never generated—there can be no bias.

```
1 for (card = 0; card < 52; card++) {  
2   r = (int)(Math.random()*(52 - card)) + card;  
3   swap(deck[card], deck[r]); }
```

- We prove that this method is not biased by running it a large enough number of times and counting the number of times each shuffle permutation appears. It turns out that the shuffle patterns appear to have an equal probability of appearing.

Shuffling Cards—ideal shuffle

- Start by initializing the deck to $\{c_1, c_2, \dots, c_n\}$.
- Select an item c_j from the list $\{c_1, c_2, \dots, c_{n-1}\}$ randomly.
- Swap c_j with c_n , turning the list into $\{c_1, c_2, \dots, c_n, \dots, c_{n-1}, c_j\}$.
- The final item in the list will not be moved again. So we will agree to call it c_n , and next we will consider the list ending at c_{n-1} : $\{c_1, c_2, \dots, c_{n-1}\}$ and again select an item c_j from the list randomly.
That c_j is swapped with c_{n-1} which will never be moved again.
- In the following round we consider the list $\{c_1, c_2, \dots, c_{n-2}\}$ and choose an item called c_j that we swap with c_{n-2} that will never be moved again.
- We carry on like this until only one element is left, so there is no choice left, and it remains where it is.
- Next we try to program the selection and replacement.

Shuffling Cards—ideal shuffle

- Next we generalize the selection steps.
- Considering the lists $\{c_1, c_2, \dots, c_{n-i}\}$ for values of $i = 0, 1, \dots, n - 1$.
- We select an item called c_j randomly from each list and then swap c_j with c_{n-i} .

```
1 for (i=1; i < N; i++) {  
2     j = i + (int)(Math.random()*(N-i) + 1;  
3     swap(card[N-i+1], card[j]); }
```

Shuffling Cards—another attempt

- Start by initializing the deck

```
1 for (i=1; i <=n; i++)  
2   card[i] = i;
```

- Generate a random number $r \in \{1, 2, \dots, n\}$.
- Replace the old array `cards` with the new array:
 $card[r + 1], \dots, card[n - 1], card[n], card[1], \dots, card[r]$
- Repeated this m times.
- What is the time complexity?
- Each iteration costs $O(n)$, and this must be repeated n times, so the time complexity is $mO(n) = O(mn)$.
- Does this method work? Of course it merely rearranges the position of the first element and does not shuffle the deck.

Shuffling Cards—riffing

- Start by initializing the deck

```
1 for (i=1; i <=n; i++)  
2   card[i] = i;
```

- Generate a random number $r \in \{1, 2, \dots, n\}$.
- Then cut the cards at $card[r]$ to create two more or less half packs:
 $\{card[1], \dots, card[r]\}$ and $\{card[r + 1], \dots, card[n - 1], card[n]\}$
- Now *riffle* the cards together m times.
- How do you riffle programmatically? Think about it. It's almost like merging.
- Merging is when two sorted arrays are put together in sorted order.

Shuffling Cards—riffing

- In our example the two arrays to merge would be $card[1..r]$ and $card[r + 1..n]$, to be merged into $target[1..n]$ in sorted order.
- An index i must be run over $[1..r]$ and another index j must be run through $[r + 1..n]$ and a third index k must be run over $[1..n]$ to take the merged item.
- When MERGING $card[i]$ is compared with $card[j]$ and the smallest (biggest) one of these is moved to $target[k]$. i or j is updated by 1 and k is always updated by 1.
- When we RIFFLE instead of copying only 1 card a small positive random integer number of cards r are copied from $card[i]$ or $card[j]$ where the i or j is randomly chosen, and k is incremented by the r the riffing batch.
Riffing is not a really good method of shuffling because it does not disturb the cards at the end points much.

The cards also need to CUT more or less in the middle and ROTATED.

Shuffling Cards—riffing

- It can be shown that if there are n cards then $3/2 \log_2 n$ riffles will shuffle the cards randomly.
- So for a pack of 52 cards you need to riffle $8.56 \approx 9$ times.
- What is the time complexity?
- Each iteration costs $O(n)$, and this must be repeated $m = 3/2 \log_2 n$ times, so the time complexity is really $mO(n) = O(n \log n)$.
- Riffing is a $O(n \log n)$ method.

How long does linear search take?

- Linear search is applied in an N entry unsorted list.
- Compare each entry in the list with the `item` being looked up.
- The best case is when `item` lies at `entry[1]`. Only one probe is required.
- The worst case is when `item` lies at `entry[N]`—needs N probes.
- The average may be estimated by looking up each `entry[i]` for $i \in [1..N]$: the 1st entry needs one probe, the 2nd entry needs two probes, ..., the N th entry needs N probes.
- In total there are $\sum_{i=1}^N i = 1 + 2 + \dots + N - 1 + N = \frac{N(N+1)}{2}$ probes.
- So, the expected number of probes is $\frac{N+1}{2}$.
- We can instead say that that average number of probes is $O(N)$.

The exact number of bits needed to store a digit

Bits in N $\lceil \log_2 N \rceil$	N	N as $2^i - 1$	digits exact $\lfloor \log_{10} N \rfloor$	$\log_{10} N$
1	1	$2^1 - 1$	0	0.000
2	3	$2^2 - 1$	0	0.301
3	7	$2^3 - 1$	0	0.845
4	15	$2^4 - 1$	1	1.176
5	31	$2^5 - 1$	1	1.491
6	63	$2^6 - 1$	1	1.799
7	127	$2^7 - 1$	2	2.104
8	255	$2^8 - 1$	2	2.407
9	511	$2^9 - 1$	2	2.708
10	1023	$2^{10} - 1$	3	3.010
...	...	...	...	...
13	8191	$2^{13} - 1$	3	3.913
14	16383	$2^{14} - 1$	4	4.214
15	32767	$2^{15} - 1$	4	4.515
16	65535	$2^{16} - 1$	4	4.816
17	131071	$2^{17} - 1$	5	5.118
19	524287	$2^{19} - 1$	5	5.720
20	1048575	$2^{20} - 1$	6	6.021
21	2097151	$2^{21} - 1$	6	6.322
22	4194303	$2^{22} - 1$	6	6.623
23	8388607	$2^{23} - 1$	6	6.924
24	16777215	$2^{24} - 1$	7	7.225
...	...	...	...	...
29	536870911	$2^{29} - 1$	8	8.730
30	1073741823	$2^{30} - 1$	9	9.031
31	2147483647	$2^{31} - 1$	9	9.332
32	4294967295	$2^{32} - 1$	9	9.633

The decimal number 1048575 is represented exactly with 20 binary digits, so *any* 6-digit decimal number also needs 20 bits for exact representation. A 24-bit number has enough bits to store 7 digits—the numbers for doing geodesy require 7 decimal digits. Why is the 32-bit number interesting?

The bits needed to store a digit

- $\lceil \log_{10} n \rceil$ yields the number of exact decimal digits needed to store n .
- $\lceil \log_2 n \rceil$ gives the number of bits required to store any number n .
- Given a word with n bits, what is the number of *decimal* digits that it can store precisely?

Since the biggest n -bit binary number is $2^n - 1$:

$$\begin{aligned} \text{Number of digits in } n\text{-bit binary} &= \lceil \log_{10}(2^n - 1) \rceil \\ &\doteq \lceil \log_{10} 2^n \rceil \\ &= \lceil n \times \log_{10} 2 \rceil \\ &\doteq \lceil n \times 0.3010299957 \rceil \end{aligned}$$

- i.e. number of *bits* $\times 0.3$,
i.e. $n \times 0.3$ gives the accuracy in *digits*.
- To find the number of *bits* needed to store m *digits*, simply calculate $\lceil m/0.3 \rceil$, e.g. to store 7 digits, use $\lceil 7/0.3 \rceil = \lceil 23.333333 \rceil = 24$ bits.

The bits needed to store decimal number

- Another approach is to think of $2^{10} = 1024$ as almost 1000 which can exactly store exactly up to 3 digits. So, 2^{21} which is $2 \times 2^{10} \times 2^{10}$ and since each 10-bits can store 3 digits, so 21 bits can store a 6-digit number precisely.
- In fact if the decimal value of the number starts with a 1 then the 21-bit binary can store 7 digit positive integers without any loss of precision.
- Note that 20 bits are necessary for 6 decimal digits.
- A 21-bit number can represent a sign and 6 decimal digits.

How many times, k , can n be divided by b so that $n/b^k < 1$?

- Let n and b be positive numbers. If the number n is represented in decimal digits, then each division by $b = 10$ removes one digit.
- We are looking for the smallest k such that $\frac{n}{b^k} < 1$, i.e. $n < b^k$, taking logarithms to the base b , $\log_b n < \log_b b^k = k$. So $k > \log_b n$, but in turn $k \leq \lceil \log_b n \rceil$, i.e. we have $\lceil \log_b n \rceil \geq k > \log_b n$.
- The number of times that n can be divided by 10 is thus $\lceil \log_{10} n \rceil$.
- The same holds for repeated halving—dividing by 2—of a binary number: each division by 2 removes one *bit* from the least significant end of the binary number, so it can only be done $\lceil \log_2 n \rceil$ times before the result becomes less than 1.
- The number of times a list of length n can be halved until its length is 0 is thus $\lceil \log_2 n \rceil$.

How long does binary search take?

- Suppose we are looking for *item* in the ordered, i.e. sorted, list $d[1], d[2], \dots, d[n]$.

```
1 def binarySearch(d, item):
2     left = 1; right = len(d) # which is n
3     while right >= left:
4         middle = (left + right)/2
5         if item == d[middle]:
6             return middle
7         else:
8             if item < d[middle]:
9                 right = middle-1
10            else:
11                left = middle+1
12    return -left # the item was not found
```

- Each iteration of the **while** loop halves the length of the list that is being searched. This can only happen $\lceil \log_2 n \rceil$ times. So, the time complexity of binary search is thus $O(\log_2 n)$.

Searching, Insertion and Deletion in lists

- Let there be n persons and a table $d[m+1]$ of positions $[1..m]$. For linear and binary search $n == m$
- If a list is not sorted then it takes $\frac{n+1}{2} = O(n)$ to search for a value in the list.
- Inserting an item into an unsorted list is extremely fast—it takes $O(1)$ time: `data[++n] = item;`
- For sorted lists, binary search takes $O(\log n)$, but inserting `item` in its correct position—maintaining the sorted order—requires: (1) a look up to find where the item must be inserted in $O(\log n)$ time, and (2) on average half of the list must be moved down one place to make place for the new item, in $\frac{n+1}{2} = O(n)$ time.

Searching, Insertion and Deletion in lists

Repeating some of the previous slide:

- Inserting an item into an *unsorted* list is extremely fast—it takes $O(1)$ time: `data[++n] = item;`
- For sorted lists, binary search takes $O(\log n)$, but inserting `item` in its correct position—maintaining the sorted order—requires: (1) a look up to find where the item must be inserted in $O(\log n)$ time, and (2) on average half of the list must be moved down one place to make place for the new item, in $\frac{n+1}{2} = O(n)$ time.
- So while *binary search* is excellent for looking items up, it is tardy when new items are inserted, as opposed to *linear search* which is slow for looking up, but extremely fast, i.e. $O(1)$ for appending new items at the end of the list.

$O(1)$ Searching, Insertion and Deletion— $\text{id} \in [1..m]$

- Assume we assign a unique identifier, $\text{id} \in [1..m]$ to each of the n persons in the class—in the order in which persons arrive in the class, or any other order.
- If the data for person with identifier id is stored at $d[\text{id}]$, then any person's data may be retrieved given their id in $O(1)$ time.
- Inserting a new arrival's data entails first assigning the next available id and by inserting his data into $d[\text{id}]$.
- Insertion and deletion can obviously be done in $O(1)$ time, and $n == m$.
- Deletion is similarly efficient. To delete an item: (1) it must be marked as deleted, and (2) its id must be added to the list of available ids , so that this id is later available for re-use.
- The method is not very practical if the identifiers have already been allocated—we will *hash* the ids .

$O(1)$ Searching, Insertion and Deletion— $id \in \mathbb{K}, |\mathbb{K}| \neq m$

- Suppose there are n persons or entries and we have a table with m positions or slots. Now n is at most m .
- A solution is to devise a function $h(id)$ where $h(\)$ is called a *scatter* function or more often a *hash* function.
- Insertion is done as follows—assuming id is not in list:

```
1  i = hash(id)
2  while d[i].isFilled:
3      i = (i + p) % m + 1
4  d[i].key = id
5  d[i].isFilled = true
```


$O(1)$ Insertion in a hash table— $id \in \mathbb{K}, |\mathbb{K}| \neq m$

Insert an item into the hash table as follows:

Insert item directly when the *slot is open*

When the *slot has an entry* but is not filled with the item's key, try another index $i = (i + p) \bmod m + 1$ where p is some prime number so that the new i lies in $[1..m]$.

If $id = d[i].key$, then the key is already present in the table, and the item most likely needs updating.

```
1  i = hash(id)
2  while d[i].isFilled and d[i].key != key:
3      i = (i + p) % m + 1
4  if d[i].key == id:
5      update(d[i], item)
6  else:
7      d[i].key = id
8      d[i].isFilled = true
```

$O(1)$ Searching— $\text{id} \in \mathbb{K}, |\mathbb{K}| \neq m$

- Searching for id is similar:

```
1  i = hash(id)
2  while d[i].isFilled and d[i].key != id:
3      i = (i + p) % m + 1
4  if !d[i].isFilled:
5      return notFound
6  else:
7      return i
```

A *hash* function $h : \mathbb{K} \rightarrow [1..m]$, which maps every $\text{key} \in \mathbb{K}$ into the set of slots $[1..m]$ where $n == m$, is called *perfect* when every key maps uniquely onto its own slot.

Hash functions— $id \in \mathbb{K}, |\mathbb{K}| \neq m$

Let $i_j = \text{hash}(id_j)$ and $i_k = \text{hash}(id_k)$. When $id_j \neq id_k$ but $i_j = i_k$, this is known as a *collision*.

1. Design the hash function to minimize collisions.
 - Attempt to spread the range of $h(id)$ uniformly over $[1..m]$.
 - $h(id)$'s values must be as random as possible.
2. Since collisions are often unavoidable methods of handling, collisions must be devised.
 - Make use of linear probing: suppose id_j is already in the table and $i_j = i_k$, then map i_k onto some multiple r of a suitable prime p , i.e.
 $i_k = (i_j + r \times p) \bmod m + 1$.
The $\bmod m + 1$ at the end ensures that
 $i_k \in [1..m]$.

A hash function with no collisions is called a *perfect hash* function.

The time complexity of *open* hashing

- In the literature this method has been given the confusing name of *open* hashing. We will consider *closed* hashing later.
- The speed of the table depends on how “full” it is. The fullness of a table is defined in terms of its *load factor*, $\alpha = \frac{n}{m}$.
- For *open* hash tables $\alpha = \frac{n}{m} \leq 1$, i.e., $n \leq m$.
- Assume that there are now k entries in the hash table, then the load factor is $\alpha = \frac{k}{m}$.
- When no entries have been made, then a single probe is enough to make the first entry.
- Suppose k entries have been made, then the probability of a collision is α . The probability of a collision on the i^{th} probe is α^i . The probability of finding an entry in precisely i probes is $\alpha^{i-1}(1 - \alpha)$.

Search length in *open* hash table


- Suppose k entries have been made, then the probability of a collision is α . The probability of a collision on the i^{th} probe is α^i . The probability of finding an entry in precisely i probes is $\alpha^{i-1}(1 - \alpha)$.

*Average Search Length*_{hash}

$$\begin{aligned} &= \sum_{i=0}^n i\alpha^{i-1}(1 - \alpha) \\ &\leq (1 - \alpha) \sum_{i=0}^{\infty} i\alpha^{i-1} \\ &= \frac{(1 - \alpha)}{(1 - \alpha)^2} \\ &= \frac{1}{1 - \alpha} \end{aligned}$$

Loading as a %	α	<i>Average Search Length</i> $\frac{1}{1-\alpha}$
50%	0.50	$1/0.50 = 2.000$
75%	0.75	$1/0.25 = 4.000$
85%	0.85	$1/0.15 = 6.667$
95%	0.95	$1/0.05 = 20$
99%	0.99	$1/0.01 = 100$
99.9%	0.999	$1/0.001 = 1000$
99.99%	0.9999	$1/0.0001 = 10000$

Chained hash tables




- The table $d[m+1]$ has m positions and n entries. In chained tables n may exceed m . The performance of a *chained* hash table depends on the average length of the chains, $\frac{n}{m}$. So the performance may be tuned.
- One probe is required, to get to the head of the chain. This is followed by a linear search of average length $\frac{n}{m} = \alpha$ is required.
- This amounts to an average of $1 + \frac{\alpha+1}{2}$ probes to find any item.
- So $ASL_{\text{chained hash}} = O(\alpha)$.

Chained hash tables can be tuned to $O(1)$


- The performance of a *chained* hash table depends on the average length of the chains, $\frac{n}{m}$. So the performance may be tuned.
- One probe is always required to get to the head of the chain, and then a linear search is done in a list of average length $\frac{n}{m} = \alpha$ is required.
- This amounts to an average of $1 + \frac{\alpha+1}{2}$ probes to find any item.
- So $ASL_{\text{chained hash}} = O(\alpha)$.
- But since α can be made independent of the number of elements n in the list, e.g. putting $6m = n$, makes $\alpha = \frac{n}{m} = 6$,
- i.e. hashing can be made to be $O(1)$.

Chained hashing—example



- Set up the hash table by entering keys: **two**, **five**, **three**, **four**, **six**, **one**, **seven**, **nine**, **eight** in this order. Other orders will give logically equivalent tables, but their links will be different.
- Initially there are $n = 0$ items in the store. Let us enter **two**, assume $id = \text{"two"}$, so, first hash it: $i = \text{hash}(id)$ giving $i = 2$. Enter id at $++n$, in the store, by: $\text{store}[n] = \langle id, 0 \rangle$. Check if the table entry 2 is open or filled by examining $d[i]$. If $d[i] \neq 0$, i.e. $d[i]$ is already occupied, the link 0 in the node $\langle id, 0 \rangle$ must be replaced by $d[i]$ as in $\langle id, d[i] \rangle$, and $d[i] == n$.

Chained hashing—example




Set up the hash table by entering keys: **two**, **five**, **three**, **four**, **six**, **one**, **seven**, **nine**, **eight** in this order.

- Now enter **five**.

There is now $n = 1$ item in the store. Entering **five**, assume $id = \text{"five"}$, so, first hash it: $i = \text{hash}(id)$ For this table $i = 10$. The data, id , may be placed at the next open place, $++n$, making $n = 2$ in the store, by: $\text{store}[n] = \langle id, 0 \rangle$ Check if the table entry i is open or filled by examining $d[i]$. Since $d[i] == 0$, i.e. $d[i]$ is *not* occupied, the link 0 in the node $\langle id, 0 \rangle$ is used as in $\langle id, 0 \rangle$, and $d[i] = n == 2$.


Chained hashing—example



Set up the hash table by entering keys: **two**, **five**, **three**, **four**, **six**, **one**, **seven**, **nine**, **eight** in this order.

- Next enter **three**.
- There are $n = 2$ items in the store. Entering **three**, put $id = \text{"three"}$, so, first hash it: $i = \text{hash}(id)$; For this table $i = 7$. The data, id , may be placed at the next open place, $++n$, making $n = 3$ in the store, by: $\text{store}[n] = \langle id, \emptyset \rangle$; The table entry $d[i]$ is still open, so the link $\emptyset$ in the node $\langle id, \emptyset \rangle$ is used and $d[i] = n == 3$;


Chained hashing—example: Enter four



Set up the hash table by entering keys: two, five, three, four, six, one, seven, nine, eight in this order.

- Now enter **four**.
- There are now $n = 3$ item in the store. Entering **four**, assume $id = \text{"four"}$, so, first hash it: $i = \text{hash}(id)$; For this table $i = 7$. The data, id , may be placed at the next open place, $++n$, makes $n = 4$ in the store, by: $\text{store}[n] = \langle id, \emptyset \rangle$; Check if the table entry i is open or filled by examining $d[i]$. Since $d[i] \neq \emptyset$, i.e. $d[i]$ is *occupied*, the link $\emptyset$ in the node $\langle id, \emptyset \rangle$ must be replaced by $d[i]$ as in $\langle id, d[i] \rangle$;, and $d[i] == n$;

Chained hashing—example: Enter **six**




Set up the hash table by entering keys: two, five, three, four, six, one, seven, nine, eight in this order.

- Now enter **six**.
- There are now $n = 4$ item in the store. Entering **six**, assume $id = \text{"six"}$, so, first hash it: $i = \text{hash}(id)$; For this table $i = 7$. The data, id , may be placed at the next open place, $++n$, makes $n = 5$ in the store, by: $\text{store}[n] = \langle id, \emptyset \rangle$; Check if the table entry i is open or filled by examining $d[i]$. Since $d[i] \neq \emptyset$, i.e. $d[i]$ is *occupied*, the link $\emptyset$ in the node $\langle id, \emptyset \rangle$ must be replaced by $d[i]$ as in $\langle id, d[i] \rangle$; and $d[i] == n$;

Chained hashing—hashInsert—array implementation

The hash table has m entries or slots that each contains the index of the key of the data that has been hashed to that slot.


The data itself is stored in a storage pool of nodes.



Java-like pseudocode to create the table and the storage pool.

```
1 int d[] = new int [m+1];  
2 node store[] = new node [6*m +1];  
3 int n = 0;
```

Chained hashing—hashInsert—array implementation



```
1 int d[] = new int [m+1]; // Java-like pseudocode
2 node store[] = new node [6*m +1];
3
4 int n = 0;
5
6 void hashInsert(string id) {
7     ++n;
8     store[n].key = id;
9     i = hash(id);
10    store[n].link = d[i];
11    d[i] = n; }
```

Chained hashing—hashSearch

This method is called as in the example:

```
1 boolean found = hashSearch("ten");
```

The method follows

```
1 boolean hashSearch(string id) {  
2   int i = hash(id);  
3   i = d[i];  
4   while (i != 0 && store[i].key != id) {  
5     i = store[i].link; }  
6   return store[i].key == id; }
```


Chained hashing into an array of nodes—Java for the main class

The coding of some of the methods clarify some of our doubtful pseudo code constructs.

```
1 public class tryHash {
2 public static void main (String args[]) {
3 String data [] =
4 {"zero", "one", "two", "three", "four", "five",
5   "six", "seven", "eight", "nine", "ten",
6   "eleven", "twelve", "thirteen", "fourteen", "fifteen",
7   "sixteen", "seventeen"};
8 int k;
9 tableHash H;
10
11   H = new tableHash();
12
13   for (k = 0; k < data.length; k++)
14     H.hashInsert(data[k]);
15   H.display(); }}
```

Here we have declared a reference or variable **H** to a hash table, and then instantiated or created it using the constructor **tableHash()**. Next, we insert the data items one-by-one into the **H**, and finally display the table.

Chained hashing into an array of nodes—the **node** class in Java



```
1 public class node {  
2  
3     long key;  
4     String id;  
5     int next;  
6  
7     public node () {  
8         key = 0L;  
9         id = "";  
10        next = 0; }  
11  
12    public node (long k, String s) {  
13        key = k;  
14        id = s;  
15        next = 0; } }
```

Chained hashing into an array—the `tableHash()` constructor

```
1 public class tableHash {
2   int m = 10;
3   int d[] = new int [m+1];
4   node store[] = new node [6*m + 1];
5   int n = 0;
6
7   public tableHash () { // the table constructor
8
9     for (int i=0; i < m + 1; i++)
10       d[i] = 0;
11
12     for (int j = 0; j < 6*m +1; j++)
13       store[j] = new node(); }
```

First `m` is defined as 10, next space is reserved for the hash table of integers `d[]`. Space is made for `store[]`, and `n` is set to zero.

In the `tableHash()` constructor the hash table is set to zeros and the pool of nodes, called `store[]`, is filled with instantiated nodes using the `node()` constructor.

A hashing function $h(id)$ in Java

```
1 int h (String id) {  
2   int hash = 0;  
3  
4   for (int ch = 0; ch < id.length(); ch++)  
5     hash = (hash + (int)id.charAt(ch)) % m;  
6  
7   return hash + 1; }
```

The function $h(id)$ simply sums the numerical or ASCII values of each character and gets $\text{mod } m$ at each step. The value of the running hash will lie in the range $[0 \dots m-1]$. The addition of 1 puts the result back into the range $[1 \dots m]$.

A method to display the entire hash table in Java

```
1
2 void display() {
3     int i, j;
4     System.out.printf("The_hash_table:\n");
5     for (i = 1; i <= m; i++) {
6         j = d[i];
7         System.out.printf("%d:_", j);
8         while (j > 0) {
9             System.out.printf("-->%s", store[j].id);
10            j = store[j].next; }
11        System.out.printf("-->%d\n", j); }
12    System.out.printf("-----\n"); }
```


The **display()** method runs through the “pointers” in **d[]** one-by-one and follows the “pointers” until the end of the chain is reached.

The `hashInsert(id)` and `hashSearch(id)` methods in Java

```
1 void hashInsert(String id) {
2   int i;
3   ++n;
4   store[n].key = (long)n;
5   i = h(id);
6   store[n].next = d[i];
7   store[n].id = id;
8   System.out.printf("insert_%s_at_%d\n" , id, n);
9   System.out.printf("%s\n" , store[n].id);
10  d[i] = n; }
11
12 boolean hashSearch(String id) {
13  int i = h(id);
14  i = d[i];
15  while (i != 0 && store[i].id != id) {
16    i = store[i].next; }
17  return store[i].id == id; } }
```

Chained hashing—hashInsert—list implementation

Now we alter the **node** so that it has a **next** pointer that is a **node** reference instead of an **integer**.



```
1 node d[] = new node [m+1];
2 int n = 0;
3
4 void hashInsert(long key, String id) {
5     node n = new node (key, id);
6     i = hash(id);
7     n.next = d[i];
8     d[i] = n; }
```

Note that the array `store[]` is no longer needed.

The program creates new `nodes` on-the-fly as they are required.

The array `d[]` appears to be a list of nodes, but it is actually a list of references to nodes.

Chained hashing—hashSearch—using lists

This method is called as in the example:

```
1 boolean found = hashSearch(10L);
```

The method follows

```
1 boolean hashSearch(long key, String id) {  
2   int i = hash(key);  
3   node here = d[i];  
4   while (here != null && here.key != key)  
5       here = here.next;  
6   return here != null && here.key == key; }
```


A class for a node

Suppose we have a **node** class, the declaration:


```
1 node n;
```

It creates the variable **n** of type **node**

We depict the instance of a **node** graphically using



When the method **n = new node();** is invoked an instance of the node is created and a pointer that points to that instance is placed into the variable **n**.



class node

```
1 public class node{
2     long key;
3     String data;
4     node next;
5
6     public node(long k, String d){
7         key = k;
8         data = d;
9         next = null; }
10
11    public node(){
12        key = 0;
13        data = "";
14        next = null; }
15
16    public String toString() {
17        return "<" + key + ",_" + data +
18                ",_" + next + ">" } }
```

class linkedList

```
1 public class linkedList{
2
3     protected static int length;
4     protected node head;
5
6     public linkedList(){
7         head = new node (0, "head_node");
8         length=0; }
9
10    public void insert(node n){
11        n.next = head.next;
12        head.next = n;
13        length++; }
14
15    public int length(){
16        return length; }
```

class LinkedList continued

```
1 public boolean isEmpty(){
2     return length==0; }
3
4 public void clear(){
5     head.next = null;
6     length = 0; }
7
8 public void display(){
9     node here = head.next;
10    while (here != null) {
11        System.out.println(
12            "Key_=" + here.key +
13            "_Data_=" + here.data);
14        here = here.next; }
15    System.out.println("length=" + length); } }
```

class tryList: Get going

```
1 import java.lang.Math.*;
2 import java.io.*;
3
4 public class tryList {
5     public int n = 2;
6     public static void main(String args[]) {
7         linkedList L = new linkedList();
8         System.out.println(
9             "L.Head.Key=_ " + L.head.key +
10            "_L.Head.Data=_ " + L.head.data +
11            "_L.Length=_ " + L.length +
12            "_L.Head.Next=_ " + L.head.next);
13         node item;
14         long id;
15         String studentName;
16         L.display();
17         L.head.next = null; // L.clear(); is better
```

class tryList Check the arguments

```
1 String theFileName = args[0];
2 if (args.length == 1) {
3     try {
4         FileInputStream inFileName
5             = new FileInputStream (theFileName);
6         DataInputStream inFile
7             = new DataInputStream (inFileName);
8         System.out.println(
9             "File:_" + theFileName
10             + "_successfully_opened");
11         // proceed with processing
```

class tryList Do the processing

```
1  int lineNumber = 0;
2  // Process first line
3  //---the header---differently
4  System.out.println(inFile.readLine());
5  // // readLine() deprecated
6  // Process rest of file
7  while (inFile.available() != 0){
8      String inString = inFile.readLine();
9      String outString;// readLine() deprecated
10     int commaPos = inString.indexOf(",");
11     // One space: handle Da Costa, Van Wyk, Le Roux, ...
12     // Not treated: van der Merwe, de la Querra, etc.
13     // Before processing
14     //2410832 VAN ZITTERS, ST ... 000000
15     // Final output for this record:
16     // Key = 2410832 Data = Van Zitters, ST
17     int spacePos = inString.substring
18         (10, commaPos).indexOf("_") + 10;
19     int minusPos = inString.substring
20         (10, commaPos).indexOf("-") + 10; }
```

class tryList Continue processing

```
1 if (spacePos > 10)
2   outString = inString.substring(0, 10)
3   + inString.substring(10, spacePos).toLowerCase()
4   + inString.substring(spacePos, spacePos+2)
5   + inString.substring(spacePos+2, commaPos).toLowerCase()
6   + inString.substring(commaPos, inString.length());
7 else if (minusPos > 10) {
8   System.out.println(minusPos);
9   outString = inString.substring(0, 10)
10  + inString.substring(10, minusPos+1).toLowerCase()
11  + inString.substring(minusPos+1, minusPos+2)
12  + inString.substring(minusPos+2, commaPos).toLowerCase()
13  + inString.substring(commaPos, inString.length()); }
14 else outString = inString.substring(0, 10)
15   + inString.substring(10, commaPos).toLowerCase()
16   + inString.substring(commaPos, inString.length());
17 studentName = outString.substring(9, 71).trim();
18 System.out.println(outString + "␣" + outString.length());
```

class tryList Build linked list

```
1  // Build the linked list:
2  id = (long)Integer.parseInt(
3      outString.substring(1,8));
4  // first 7 characters: the student number
5  item = new node (id, studentName);
6  L.insert (item);
7  }
8  inFile.close();
9  L.display(); }
10 catch (Exception e) {
11     System.err.println("File:_" + theFileName
12         + "_error_occurred"); } }
13 else {
14     System.err.println("There_are_" + args.length
15         + "parameters,_there_should_be_1"); } } }
```


The `lookup` method

```
1 public node lookup (long key) { //tiny
2   node here = this.head.next,
3     prev = here;
4   while (here != null && here.key != key) {
5     prev = here;
6     here = here.next; }
7   return prev; }
```

The Makefile

```
1 % make $ mean \$ in lstlisting
2 OBJ = tryList
3 LIST = linkedList
4 DEP = node
5 DATA = cos224.list
6
7 all:$(DEP).class, $(LIST).class
8     -rm errors
9     javac -deprecation $(OBJ).java 2> errors
10    -rm output
11    java $(OBJ) $(DATA) > output
12
13 $(DEP).class, $(LIST).class:
14     javac -deprecation $(LIST).java
15     javac -deprecation $(DEP).java
16
17 clean:
18     -rm *~ *.class errors output
```

Improved `class linkedList`

```
1 public class linkedList{
2
3     protected static int length;
4     protected node head;
5
6     public linkedList(){
7         head = new node (0, "head_node");
8         length=0; }
9
10    public void insert(node n){
11        n.next = head.next;
12        head.next = n;
13        length++; }
14
15    public int length(){
16        return length; }
```

Improved `class LinkedList` continued

```
1 public boolean isEmpty(){
2     return length==0; }
3
4 public void clear(){
5     head.next = null;
6     this.length = 0; }
7
8 public void display(){
9     node here = this.head.next;
10    while (here != null) {
11        System.out.println(
12            "Key_=_ " + here.key +
13            "_Data_=_ " + here.data);
14        here = here.next; }
15    System.out.println(
16        "length=" + this.length); } }
```

The `class node` revised

This version of `class node` is not properly encapsulated—one can access parts of the class that should be invisible outside of the `node` class.

```
1 public class node{
2     long key;
3     String data;
4     node next;
5
6     public node(long k, String d){
7         key = k;
8         data = d;
9         next = null; }
10
11    public node(){
12        key = 0;
13        data = "";
14        next = null; }
15
16    public String toString() {
17        return "<" + key + ",_" + data + ",_" + next + ">" } }
```

For example the `insert` method in the `linkedList` class can access the `next` field directly.

```
1 public void insert(node n){
2     n.next = head.next;
3     head.next = n;
4     length++;
5 }
```

The `class node` revised

The fields inside `node` should be made `private` or `protected` and methods to `get` and `set` them should be supplied.

```
1 public class node{
2     private long key;
3     private String data;
4     private node next;
5
6     public node(long k, String d){
7         key = k;
8         data = d;
9         next = null; }
10
11     public node(){
12         key = 0;
13         data = "";
14         next = null; }
15
16     public void setKey(long k) {
17         key = k; }
18
19     public long getKey() {
20         return key; } . . . }
```

```
1     public long getKey() {
2         return key; }
3
4     public String getData() {
5         return data; }
6
7     public node getNext() {
8         return next; }
9
10    public void setKey(long k) {
11        key = k; }
12
13    public void setData(String d) {
14        data = d; }
15
16    public void setNext(node n) {
17        next = n; }
```

The `class node` revised

Now the `LinkedList` method `insert` becomes

```
1 public void insert(node n){  
2     \ n.next = head.next;  
3     n.setNext(head.getNext());  
4  
5     \ head.next = n;  
6     head.setNext(n);  
7  
8     length++; }
```

Linked list—insert (node n)

Start with an instantiated list

```
1 linkedList L;  
2 L = new linkedList();
```




Insert a "first" node

```
1 node n = new node(1L, "first")  
2 L.insert(n);
```

Insert a "second" node


```
1 node n = new node(2L, "second")  
2 L.insert(n);
```

```
1 public void insert(node n){  
2   if (head == null)  
3     head == n;  
4   else {  
5     n.next = head;  
6     head = n;  
7     length++; }  
}
```



Linked list—no dummy head: remove (long Key)


Start with a list of two elements. Consider removing either of the nodes.



```
1 public node remove(long Key){
2   node here = head, previous = here;
3   while (here != null && here.key != Key) {
4     previous = here;
5     here = here.next;}
6   if (here != null) {
7     if (here == head)
8       head = here.next;
9     else
10      previous.next = here.next;
11    length--; }
12   return here;}
```

Linked list with dummy head—remove (long Key)

Start with a list of two elements. Consider removing either of the nodes.



```
1 public node remove(long Key){
2   node here = head.next, previous = here;
3   while (here != null && here.key != Key) {
4     previous = here;
5     here = here.next;}
6   if (here != null) {
7     previous.next = here.next;
8     length--; }
9   return here;}
```

The code is much simpler.

Stacks


A stack behaves like a linked list where operations are allowed only at the head of the list.

Suppose that a **stack** variable called **S** is created by `public static stack S;`

Our stack will hold items of type **node**.


If **n** is of type **node**, then a new instance of a **stack** with **n** placed on the stack is created by

`S = new stack(n);`



Stacks—a similar stack created differently

If **m** is a **node**, then it may be pushed onto the stack with `S.push(m);`



Stacks—a similar stack created differently


Suppose that a **stack** variable called **S** is created by

```
public static stack S;
```

If **n** is of type **node**, then a new instance of a stack with *nothing* placed on the stack is created by

```
S = new stack();
```


If **n** is a **node**, then it may be pushed onto the stack with **S.push(n)**;



Stacks—`S.push(l)` and `item = S.pop()`;

If `m` is a **node**, then it may be pushed onto the stack with `S.push(m)`;

If `l` is a **node**, then it may be pushed onto the stack with `S.push(l)`;




Stacks—`S.push(l)` and `item = S.pop();`

If the stack is not empty, the function method

`node pop()`

returns the value of uppermost item on the stack and also deletes it off the stack. Now create a variable called `item` of type `node` that points to the node that was removed from the top of the stack.

`node item = S.pop();`



Stacks

- Letting
`S = new stack();`
creates an instance of an empty stack to which **S** points.
- Further items of data, **data**, must first be put into a **node** before being pushed onto our stack by first creating an instance of a node
`node n = new node(data);`
- **n**, may then be pushed onto the stack by
`S.push(n);`
- The top item of stack **stack** is retrieved by:
`item = S.pop();`
which deletes it from the top of the stack.

Stacks implemented with a fixed array

- A stack can also be made using an *array* of nodes.
- The integer **top** points to the top of the stack.
- Set the size of the array `static int m = 1<<20;`
- Create an array to hold the stack
`static node [] d = new node [m+1];`
- Put methods `void push(node n)` and `node pop()` into the **class**.
A useful method is `node top()` which undestructively reads the the uppermost element of the stack. The class would not be complete without the functions `int getLength()` and `boolean isEmpty()`.
- The `boolean isFull()` method is questionable because when a push is being done the node being stacked has already been allocated memory.

A stack implemented with a fixed array

Build a stack of nodes.

```
1 public class stack {
2   public static node [] d = new node [1<<20+1];
3   public static int top;
4
5   public stack() {
6     top = 0; }
7
8   public void push (node n) {
9     d[++top] = n; }
10
11  public node pop() {
12    if (top > 0)
13      return d[top--];
14    else
15      return null; } }
```

A stack implemented as linked nodes

Use the deprecated non Object Oriented **node** class.

```
1 public class stack {  
2  
3     public static node top;  
4  
5     public stack() {  
6         top = null; }  
7  
8     public void push (node n) {  
9         n.next = top;  
10        top = n; }  
11  
12    public node pop() {  
13        node value = top;  
14        top = top.next;  
15        return value; } }
```

A stack implemented as linked nodes


Use the Object Oriented **node** class.

```
1 public class stack {  
2  
3     public static node top;  
4  
5     public stack() {  
6         top = null; }  
7  
8     public void push (node n) {  
9         n.setNext(top); // n.next = top;  
10        top = n; }  
11  
12    public node pop() {  
13        node value = top;  
14        top = top.getNext(); // top = top.next;  
15        return value; } }
```


Queues implemented with a fixed array

- A queue is a first-in-first-out (FIFO) data structure whose elements are served when they reach the front but they enter the queue at the rear.
- It may be implemented with a fixed length array with k elements.
- A queue is empty when its **length** is 0 (zero).

Note: We use the counter **length**.




- A queue with only the node 23 entered.




Queues implemented with fixed arrays

- A queue with the nodes 23, 45, 67 and 89 entered in this order.




- The queue where the nodes 23, 45 have been dequeued or “served”.




Queues implemented with fixed arrays

- The previous queue where the nodes 02, 24 and 46 have also been entered.




- When k nodes have been *entered* using—*enqueue*—the last node of the array becomes occupied and the next free place is at 0 (zero). There may be no elements in the queue. How is **rear** calculated?



Fixed-array queues

—computing **rear** and **front**

- When k nodes have been *entered*—the last node of the array becomes occupied and the next free place is at 0 (zero). How should **rear** be calculated?




- Simply adding 1 as in **rear++** will crash when the value of **rear** reaches k —the maximum subscript of the array is $k - 1$. Use:
rear = (rear + 1) % k; and also **length++;**
- Updating **front** when doing a *dequeue* is similar:
front = (front + 1) % k; and also **length--;**
- Note that a 1 is *added* in both cases.

Fixed-array queues

—computing **rear** and **front**

- The queue after enqueueing the item 00 :




- If the “fullness” of the queue is implemented using a **length** counter and the relation, **length** $\leq k$, the array may be completely filled with **rear** pointing to the same place as **front**.
- Note: When the comparison **rear** $==$ **front** is used to test if the array-implemented queue is filled without using the queue’s **length** counter then it is impossible to say whether the array is *full-up* or *empty*, so in this case the queue is regarded as being full when its length is $k - 1$.
- Here **length** $= (\text{rear} - \text{front} + k) \% k$; calculates length when there is no counter.


A queue implemented with linked nodes

- Queues may be implemented with a linked list of nodes with a pointer to the serving or leaving end called **front**, and with another pointer to the joining end named **rear**.

- A queue's **node**:




- An empty **queue**




A queue with linked nodes


- An queue with one node:



- A queue with two linked nodes.




- An queue with three linked nodes.



Queue's node class—elements private

A queue's node



```
1 public class node{
2
3     private long key;
4     private String data;
5     private node next;
6
7     public node(){
8         key = 0;
9         data = "";
10        next = null; }
11
12    public node(long k, String d){
13        key = k;
14        data = d;
15        next = null; }
16
17    public long getKey() {
18        return key }
19        ...
```

Queue's node—methods public

```
1 public void setKey(long k) {  
2     key = k; }  
3  
4 public String getData() {  
5     return data; }  
6  
7 public void setData(String d) {  
8     data = d; }  
9  
10 public node getNext() {  
11     return next; }  
12  
13 public void setNext(node n) {  
14     next = n; }  
15  
16 public String toString() {  
17     return "<" + key + ", " + data + ", " + next + ">" } }
```

The queue itself

```
1 public class queue{
2
3 private int length;
4 private node front;
5 private node rear;
6
7 public queue(){
8     rear = new node (0, "head_node");
9     front = rear;
10    length=0; }
11
12 public queue(long k, String d){
13     rear = new node (k, d);
14     front = rear;
15     length=0; }
16
17 public int size(){
18     return length; }
19
20 public boolean isEmpty(){
21     return length==0; }
```

The basic manipulations of a queue

```
1 public boolean isFull(){
2     return length==N; }
3
4 public void enqueue(node n){
5     rear.setNext(n);
6     rear = n;
7     length++; }
8
9 public node dequeue(){
10 node n;
11     if (isEmpty())
12         return null;
13     else {
14         n = front;
15         front = front.getNext();
16         length--;
17         return n; } }
18
19 public node topQueue(){
20     if (isEmpty())
21         return null;
22     else {
23         return front; } }
```


void display() is useful for debugging

```
1 public void display(){
2     node here = front;
3     int listed = 0;
4     while (here != rear && listed++ < 6) {
5         System.out.println(here);
6         here = here.getNext(); }
7     if (length > 0)
8         System.out.println("Queue's_length=_ " + length);
9     else
10        System.out.println("Queue_is_Empty"); } }
```

The necessity of a **void display()** method is questionable, but we find it useful for debugging our **queues**.


The node of a deque

A node for a deque



The node points backwards with its **prev** pointer and forwards with its **next** pointer.

A deque




The “sentinel” node’s **next** pointer points to the front of the deque and the sentinel’s **prev** pointer points to the the rear of the deque.


The deque has methods for adding and removing elements at both ends.

Insert a node to the front of a deque

A deque node




A deque



```
1 public void insertFirst(node n){  
2     n.setNext(header.getNext());  
3     n.setPrev(header);  
4     (header.getNext()).setPrev(n);  
5     header.setNext(n);  
6     length++; }
```

Insert a **node** to the rear of a deque


A deque



```
1 public void insertLast(node n){  
2     n.setPrev(header.getPrev());  
3     n.setNext(header);  
4     (header.getPrev()).setNext(n);  
5     header.setPrev(n);  
6     length++; }
```

Remove a **node** from the rear of a **deque**


A deque



```
1 public node removeLast(){
2   node n = null;
3   if (!isEmpty()) {
4     n = header.getPrev(); // 1
5     header.setPrev(n.getPrev()); // 2
6     (n.getPrev()).setNext(header); // 3
7     n.setPrev(null);
8     n.setNext(null);
9     length--; }
10  return n; }
```

Remove a node from the rear of a deque

- A deque




-

```
1 public node removeLast(){
2   node n = null;
3   if (!isEmpty()) {
4     n = header.getPrev(); // 1
5     header.getPrev(n.getPrev(n)); // 2
6     (n.getPrev()).setNext(header); // 3
7     n.setPrev(null);
8     n.setNext(null);
9     length--; }
10  return n; }
```

Binary search trees (BST)


- A binary search tree (BST) is an efficient data structure for storing and retrieving data.
- A search in a balanced or proper BST is $O(\log n)$ at worst.
- A binary tree is *proper* or *complete* when the nodes at the lowest level are all *leaves* and these leaves are the only nodes with null pointers.
- Numbering the level with only one node in it, *level-0*, there are $n = 2^k$ elements in the k th level of a proper binary tree.
- The bottom level of a binary tree has one more node than the rest of the tree above it.
- The number of nodes in a complete binary tree up to but excluding *level-k* is $n = 2^{k+1} - 1 - 2^k = 2^k - 1$.
- Since the height of a proper binary tree is thus k , it follows that the maximum search length is k , but $k = \log n + 1$, i.e. $O(\log n)$.

The nodes of a tree



- A tree node
- The nodes in a binary **tree** have a **left** and a **right** forward pointer, a back link to the **parent** is useful. The element also should contain a **key** and **data** field.
- In a *binary search tree* each node's **left** pointer points to a node such that **left.key** < **key** when **left** $\neq$ **null** and **right** and a pointer that points to a node such that **right.key** $\geq$ **key** when **right** $\neq$ **null**.

A proper binary search tree



```
1 public class tryTree{
2 public static void main(String args[]){
3 tNode tItem;
4 tree T = new tree ();
5
6 T.display ();
7 T.insert (new tNode (23, "twenty_three"));
8 T.insert (new tNode (17, "seventeen"));
9 T.insert (new tNode (7, "seven"));
10 T.insert (new tNode (35, "thirty_five"));
11 tItem = new tNode (41, "fourty_one");
12 T.insert (new tNode (27, "twenty_seven"));
13 T.insert (tItem);
14 T.insert (19, "nineteen");
15 T.display (); } }
```


A class for a tNode

```
1 public class tNode{
2
3 private long key; private String data;
4 private tNode left; private tNode parent;
5 private tNode right;
6
7 public tNode() {
8     key = 0;
9     data = "";
10    parent = null;
11    left = null;
12    right = null; }
13
14 public tNode(long k, String d) {
15     key = k;
16     data = d;
17     parent = null;
18     left = null;
19     right = null; }
20
21 public long getKey() {
22     return key; }
23
24 public void setKey(long k) {
25     key = k; }
26 // continued on next slide
```

The tNode class (continued)

```
1 public String getData() {  
2     return data; }  
3 public void setData(String d) {  
4     data = d; }  
5  
6 public tNode getLeft() {  
7     return left; }  
8 public void setLeft(tNode n) {  
9     left = n; }  
10  
11 public tNode getRight() {  
12     return right; }  
13 public void setRight(tNode n) {  
14     right = n; }  
15  
16 public tNode getParent() {  
17     return parent; }  
18 public void setParent(tNode n) {  
19     parent = n; } }
```

Defining a tree

- An empty **tree** has a **root** set to **null** and a counter **length** set to 0 by the constructor.

root:

null

length:

0

```
1 public class tree {  
2     private int length;  
3     private tNode root;  
4  
5     public tree() {  
6         root = null;  
7         length=0; } }
```

- An instance of the tree is defined in the user program by:

```
1 tree T = new tree();
```

- The first **node** in a tree of type **tNode** is pointed to by **root**.
- A subsequent node **n** is added to the left subtree if **n.key < root.key** or otherwise to its right.

Inserting a node into an empty tree

- An empty **tree** has a **root** set to **null** and a counter **length** set to 0 by the constructor.

```
1 tree T = new tree();
```

root:

null

length:


0

- The tree after inserting a node of type **tNode** with **key = 23** using:

```
1 T.insert(new tNode(23, "twenty_three"));
```

```
1 public class tree {  
2   ...  
3   public void insert(tNode n) {  
4     if (root == null) {  
5       root = n;  
6       root.setParent(null); }  
7   else  
8     root.insert(n); }  
}
```

root:




length:

1

- **root.insert(tNode)** in the **tNode** class inserts the **tNode** **n** into the “tree” starting at **root**.


Inserting a second node into the tree

- Starting with a tree with a single node at the root



- and after inserting a node with `key = 17` using:

```
1 T.insert(new tNode(17, "seventeen"));
```




- The code that does this lies in the `tNode` class.

```
1 public void insert(tNode n) {  
2     if (n.key < key) // n into left subtree  
3         if (left == null) {  
4             n.parent = this;  
5             left = n; }  
6     else  
7         left.insert(n); ... }
```

- If `left != null`, then `left.insert(n)` is called, as happens when a node with `key = 7` is entered.


Adding a node with **key = 7**

- Start with the two node tree:



- Insert a node with **key = 7** using:

```
1 T.insert(new tNode(7, "seven"));
```




The pointers down the left are followed until a **null** is found.

```
1 public void insert(tNode n){
2     if (n.key < key) // n into left subtree
3         if (left == null){
4             n.parent = this;
5             left = n;
6         }
7     else
8         left.insert(n);
```

Adding the node with key = 35

- Insert a node with key = 35 using:

```
1 T.insert(new tNode(35, "thirty_five"));
```




The **right** pointer of the root node is **null** and **n** is inserted there.

```
1 public void insert(tNode n) {  
2     if (n.key < key) // n into left subtree  
3         if (left == null) {  
4             n.parent = this;  
5             left = n; }  
6     else  
7         left.insert(n);  
8     else // n into rightsubtree  
9         if (right == null) {  
10            n.parent = this;  
11            right = n; }  
12    else  
13        right.insert(n); }
```

Adding the node with key = 41

- Insert a node with key = 41 using:

```
1 T.insert(new tNode(41, "fourty_one"));
```




The **right** pointer of the root node is not **null** and the pointer is followed to **35** whose right pointer is **null**. where **n** is inserted.

```
1 public void insert(tNode n) {
2     if (n.key < key) // n into left subtree
3         if (left == null) {
4             n.parent = this;
5             left = n; }
6     else
7         left.insert(n);
8 else // n into right subtree
9     if (right == null) {
10        n.parent = this;
11        right = n; }
12    else
13        right.insert(n); }
```


Adding the node with key = 27

- Insert a node with key = 27 using:

```
1 T.insert(new tNode(27, "twenty_seven"));
```




The **right** pointer of the root node is not **null** and the pointer is followed to **35** whose left pointer is **null**. where **n** is inserted.

```
1 public void insert(tNode n) {
2     if (n.key < key) // n into left subtree
3         if (left == null) {
4             n.parent = this;
5             left = n; }
6     else
7         left.insert(n);
8 else // n into right subtree
9     if (right == null) {
10        n.parent = this;
11        right = n; }
12    else
13        right.insert(n); }
```

Adding the node with key = 19

- Insert a node with key = 19 using:

```
1 T.insert(new tNode(19, "nineteen"));
```



The **left** pointer of the root node is not **null** and is followed to **17** whose right pointer is **null** where **n** is inserted.

```
1 public void insert(tNode n) {
2     if (n.key < key) // n into left subtree
3         if (left == null) {
4             n.parent = this;
5             left = n; }
6     else
7         left.insert(n);
8     else // n into right subtree
9         if (right == null) {
10            n.parent = this;
11            right = n; }
12    else
13        right.insert(n); }
```

The tree class

```
1 public class tree{
2 private static final int N = 100000;
3 private static int listed = 0;
4
5 private static int length;
6 private tNode root;
7
8 public tree() {
9     root = null;
10    length=0; }
11
12 public void setRoot(tNode n) {
13     root = n; }
14 public tNode getRoot() {
15     return root; }
16
17 public void insert(tNode n) {
18     if (root == null) {
19         root = n;
20         trace(n, "entered_at_root"); }
21     else
22         insert(root, n); }
23 // two more insert methods on next slide
```

insert methods entirely in tree class

```
1 public void insert(long k, String d){
2     tNode n = new tNode(k, d);
3     if (root == null){
4         root = n;
5         trace(n, "entered_at_root"); }
6     else
7         insert(root, n);
8     length++; // Why not at the start? }
9
10 public void insert(tNode here, tNode n) {
11     if (n.getKey() < here.getKey()) // n goes left
12         // n must go into left subtree
13         if (here.getLeft() == null) {
14             n.setParent(here);
15             here.setLeft(n); }
16     else insert(here.getLeft(), n);
17     else if (here.getRight() == null) { // n goes right
18         n.setParent(here);
19         here.setRight(n); }
20     else insert(here.getRight(), n); }
```

Methods in class tree

```
1 \begin{verbatim}
2 public void tDelete(long key){
3     . . . follows later
4     length--; }
5
6 public int size(){
7     return length; }
8
9 public boolean isEmpty(){
10    return length==0; }
11
12 public boolean isFull(){
13    return length==N; }
```

The *lNr* display method for class tree

```
1 public void display() {
2     if (this == null)
3         System.out.println("Tree_is_empty.");
4     else
5         display(root); }
6
7 public void display(tNode n) { // left Node right
8                                 // i.e. lNr
9     if (n != null) {
10        display(n.getLeft()); // left
11        System.out.print( // Node
12            "In_Tree: _Node_number_" + listed++ +
13            "_Key_=" + n.getKey() +
14            "_Data_=" + n.getData());
15        if (n.getParent() == null)
16            System.out.println("_Parent.key=_null");
17        else
18            System.out.println("_Parent.key_="
19                                + (n.getParent()).getKey());
20        display(n.getRight()); } } } // right
```

A tree class where nodes are trees

```
1 public class tree {
2     private static final int N = 100000;
3     private static int length = 0;
4     private static tree root = null;
5     // tree nodes
6     private long key; private String data;
7     private tree parent;
8     private tree left; private tree right;
9
10    public tree() {
11        key = 0;
12        data = "";
13        parent = null;
14        left = null;
15        right = null; }
16
17    public tree(long k, String d) {
18        key = k;
19        data = d;
20        parent = null;
21        left = null;
22        right = null; } }
```

A tree class where nodes are trees

```
1 public void insertNode(tree n) {
2     if (root == null){
3         root = n;
4         trace(n, "entered_at_root"); }
5     else
6         root.insert(n);
7     length++; }
8
9 public void insertNode(long key, String data) {
10    tree n = new tree(key, data);
11    if (root == null) {
12        root = n;
13        trace(n, "entered_at_root"); }
14    else
15        root.insert(n);
16    length++; }
```


A tree class where nodes are trees

```
1 public void setRoot(tree n) {  
2     root = n; }  
3 public tree getRoot() {  
4     return root; }  
5  
6 public void tDelete(long key) {  
7     length--; }  
8  
9 public int size() {  
10    return length; }  
11  
12 public boolean isEmpty() {  
13    return length==0; }  
14 public boolean isFull() {  
15    return length==N; }
```

A tree class where nodes are trees

```
1 public long getKey() {  
2     return key; }  
3 public void setKey(long k) {  
4     key = k; }  
5  
6 public String getData() {  
7     return data; }  
8 public void setData(String d) {  
9     data = d; }  
10  
11 public tree getLeft() {  
12     return left; }  
13 public void setLeft(tree n) {  
14     left = n; }
```

A tree class where nodes are trees

```
1 public tree getRight() {  
2     return right; }  
3 public void setRight(tree n) {  
4     right = n; }  
5  
6 public tree getParent() {  
7     return parent; }  
8 public void setParent(tree n) {  
9     parent = n; }
```

A tree class where nodes are trees

```
1 public void insert(tree n) {
2     if (n.key < key)
3         // n must go into left subtree
4         if (left == null){
5             n.parent = this;
6             trace(n, "entered_at_here.L");
7             left = n; }
8     else {
9         trace(n, "insert(here.L,_n)");
10        left.insert(n); }
11    else
12        // n must go into rightsubtree
13        if (right == null){
14            n.parent = this;
15            trace(n, "entered_at_here.R");
16            right = n; }
17        else {
18            trace(n, "insert(here.R,_n)");
19            right.insert(n); } }
```

A tree class where nodes are trees

```
1 public void displayTree() {
2     if (root == null)
3         System.out.println("Tree_is_empty.");
4     else
5         root.display();
6     System.out.println(listed + "_nodes_in_tree."); }
7
8 public void display(){// Node left right
9     // i.e. Nlr
10    if (left != null) left.display(); // left
11    System.out.print("// N
12        "In_Tree:_Node_number_" + listed++ +
13        "_Key=_ " + key +
14        "_Data=_ " + data);
15    if (parent == null)
16        System.out.println("_Parent.key=_null");
17    else
18        System.out.println("_Parent.key=_ "
19            + parent.key);
20    if (right != null) right.display(); } // right
21
22 public void trace(tree n, String message) {
23     System.out.println("tree:_ "
24         + n.key
25         + " " + message); } }
```

A method for removing nodes from a tree

```
1 public void Remove(long k) {
2     tree n = find(k);
3     tree here;
4     if (n != null)
5         if (n.left == null) // one right child
6             if (n.parent == null){
7                 root = n.right;
8                 if (root != null)
9                     root.parent = null;
10            }
11            else if (n.key < (n.parent).key)
12                (n.parent).left = n.right;
13            else
14                (n.parent).right = n.right;
15        else if (n.right == null) // one left child
16            if (n.parent == null){
17                root = n.left;
18                if (root != null)
19                    root.parent = null;
20            }
21            else if (n.key < (n.parent).key)
22                (n.parent).left = n.left;
23            else
24                (n.parent).right = n.left;
```

A method for removing nodes from a tree


```
1  else { // both children not null
2      here = n.left;
3      while (here.right != null)
4          here = here.right;
5      if (here.parent == null){
6          root = here.left;
7          if (root != null)
8              root.parent = null;
9      }
10     else if ( here.key < (here.parent).key)
11         (here.parent).left = here.left;
12     else
13         (here.parent).right = here.left;
14     n.key = here.key;
15     n.data = here.data; }
16 else
17     System.out.println("Node_" + k + "_not_in_tree"); }
```

Remove node from a tree—more readable

```
1 public void remove(long k) {
2     tree n = find(k);
3     tree here;
4     if (n != null)
5         if (n.left == null) // one right child
6             updateTree(n, n.right);
7         else if (n.right == null) // one left child
8             updateTree(n, n.left);
9         else { // both children not null
10            here = n.left;
11            while (here.right != null)
12                here = here.right;
13            updateTree(here, here.left);
14            n.key = here.key;
15            n.data = here.data; }
16     else
17         System.out.println("Node_" + k + "_not_in_tree"); }
18
19 public void updateTree (tree n, tree p){
20     if (n.parent == null){
21         root = p;
22         if (root != null)
23             root.parent = null; }
24     else if (n.key < (n.parent).key)
25         (n.parent).left = p;
26     else
27         (n.parent).right = p; }
```


Explanation of removal of nodes from a tree



Consider the tree




1. Removing elements with zero children is the easiest: Simply replace the pointer to the leaf with a null. We will see later that *this case can be ignored* because it is subsumed by the case where any one of the pointers is a null. Initially, the nodes with keys 11, 19, 27 and 41 are such nodes.
2. To remove a node with only one child: The non-null pointer emanating from the node to be deleted replaces the pointer pointing to the node that must be removed. Initially, the nodes with keys 7, 9, and 10 are such nodes.
3. To remove a node with two children: Find its biggest (smallest) child in its left (right) subtree, copy that biggest (smallest) child over the node to be deleted and then remove the node that was copied as usual. The copied node cannot have two children.

Remove using `update(n, p)`

- The method `update(n, p)` is given a non-null node `n` to be removed and a pointer `p` to the node to replace it.
- If `n.parent == null` then `n` is the root node, i.e. it has no parent:



- Then `root = p;` suffices, but `p` could become `null`. This would make the `root == null`,





otherwise the new node, now pointed to by `root`, must have its `parent` set to `null`, i.e.

```
1 if (root != null)
2   root.parent = null;
```

Remove using `update(n, p)`



- The parent may have its `left` pointing to `n`. Test this using:

```
1  else if (n.key < (n.parent).key)
2      (n.parent).left = p;
3  else
4      (n.parent).right = p;
```



Remove using update(n, p)

```
1 public void updateTree (tree n, tree p){
2     if (n.parent == null){
3         root = p;
4         if (root != null)
5             root.parent = null;
6     }
7     else if (n.key < (n.parent).key)
8         (n.parent).left = p;
9     else
10        (n.parent).right = p;
11 }
```



Traversing trees—depth-first

- A tree may be traversed in ℓNr order or any of its other 5 combinations, $\ell rN, N\ell r, Nr\ell, rN\ell, r\ell N$.

```
1 public void NlrTraverse (tree t) {  
2     visit(t);  
3     if (t.left != null)  
4         left.NlrTraverse();  
5     if (t.right != null)  
6         right.NlrTraverse(); }
```

- $N\ell r$ is also called ‘depth-first’ order.
- A **visit** to a node could entail printing its value or processing it somehow.

Traversing trees—breadth-first

- A tree may be traversed in ‘breadth-first’ order, where all the nodes on one level are dealt with before going down to the next level.
- Start by queueing the root.
- Service the head of the queue and put the children of each node into the queue and repeat until the queue empties.

```
1 queue Q = new queue();
2 public void BFS(tree t) {
3     Q.enqueue(t);
4     while (!Q.isEmpty()) {
5         n = Q.dequeue();
6         Q.enqueue(n.left);
7         Q.enqueue(n.right);
8         visit(n); } }
```

- A **visit** to a node could entail printing or processing it.

Heaps—priority trees

- A *heap* has a tree structure without pointers,
- “most important” item is at the top of the tree,
- and is also called a *priority queue*.
- Store the *heap* in an array, **A**, of type **node** **[]**.
- $\text{leftChild}(i) \equiv 2 \times i$
- $\text{rightChild}(i) \equiv 2 \times i + 1$
- $\text{parent}(i) \equiv \lfloor i/2 \rfloor$
- The *heap* property:
 $A[\text{parent}(i)] \geq A[i]$ for a heap with its greatest element on top, *or*
 $A[\text{parent}(i)] \leq A[i]$ with the minimum element on top.

The *heap* property stated in words

- In a maximizing *heap*:
For every node its key is greater or equal to the key of any of its descendants.
- In a minimizing *heap*:
Every node's key is less than or equal to the key of any of its descendants.

Heaps—the methods

- Constructors, gets and sets.
- *node removeMax()*; or *node removeMin()*;
- *boolean isEmpty()*;
- *boolean isFull()*;
- *insert(node n)*;
- *buildUp(node X[], int sizeX)*;
- *heapify(int i)*;
- *sort(node X[], int sizeX)*;

Heaps—Constructors, gets and sets

```
1 public class Heap() {  
2     private int final maxSize = 1 << 20;  
3     private int n = 0; // n is the size of the heap  
4     private node[] A = new node[maxSize+1];  
5  
6     public Heap(){ }  
7  
8     int getSize(){  
9         return n; }  
10    void setSize(int i){  
11        n = i; }  
12  
13    ... }
```

Heaps—*heapify(int i)*

```
1 public node heapify(int i) {  
2  
3     l = left(i); r = right(i);  
4     if (l <= n && A[l].key > A[i].key)  
5         top = l;  
6     else  
7         top = i;  
8  
9     if (r <= n && A[r].key > A[top].key)  
10         top = r;  
11  
12     if (top != i) {  
13         swap(A[i], A[top]); // this is not Java  
14         heapify(top); } }
```

Heaps—*heapify(int i)*—a bit faster

```
1 public node heapify(int i) {  
2  
3      $\ell = 2 \times i$  ;  $r = \ell + 1$ ;  
4     if ( $\ell \leq n$  &&  $A[\ell].key > A[i].key$ )  
5         top =  $\ell$ ;  
6     else  
7         top = i;  
8  
9     if ( $r \leq n$   
10        &&  $A[r].key > A[top].key$ )  
11        top = r;  
12  
13     if (top != i) {  
14         swap( $A[i]$ ,  $A[top]$ );  
15         if (left(top) <= n)  
16             heapify(top); } }
```

Heaps—*heapify(int i)*—iterative

First, we see if `child = left(i)` is within the heap. Next, make `child += 1` if the right child is bigger than the left child, then `break` if `A[i] >= A[child]` or carry on by swapping `A[i]` and `A[child]` and repeat another level.

```
1 public node heapify(int i) {
2     int child = left(i);
3     while (child <= n)
4         if (child < n
5             && A[child] < A[child+1])
6             child++;
7     if (A[child] < A[i]){
8         swap(A[i], A[child]); // not Java
9         i = child;
10        child = left(i); }
11    else
12        break } }
```

Heaps—*insert(node N)*

```
1 public node insert(node N) {  
2  
3     n++;  
4     i = n;  
5     while (i > 1 && A[parent(i)].key < N.key) {  
6         A[i] = A[parent(i)];  
7         i = parent(i); }  
8     A[i] = N; }
```

Heaps—*node removeMax()*

```
1 public node remove() {  
2     if (n < 1)  
3         return null;  
4     else {  
5         top = A[1];  
6         A[1] = A[n];  
7         n--;  
8         heapify(1);  
9         return top; } }
```

Heaps—*node buildUp(node X[], int sizeX)*

Pseudocode:

```
1 // pre: given an array X of sizeX nodes
2 // post: node X[] is copied into the
3 // heap array, node A[] which is
4 // turned int a heap of size n.
5 public void buildUp(node X[], int sizeX) {
6
7     for (i = 1; i <= sizeX; i++)
8         A[i] = X[i];
9
10    n = sizeX;
11    for (i = n/2; i > 0; i--)
12        heapify(i); }
```


Heaps—*node [] sort(node X[], int sizeX)*

```
1 // pre: Given array X[] of sizeX elements
2 // n becomes the size of the heap.
3 // post: The array node A[] is sorted.
4
5 public node [] sort(node X[], int sizeX) {
6     buildUp(X, sizeX);
7     for (i = n; i>1; i--){
8         swap(A[1],A[i]);
9         n--;
10        heapify(1); }
11    n = sizeX;
12    return A;}
```

Note that the heap is both created and “destroyed” in the process of sorting using **buildUp** and **sort**.

The procedure invocation **swap(A[1], A[i])** needs to be adapted to your programming language. Try **swap(A, 1, i);** in Java.

Heaps—*node sort(node X[], int sizeX)*

```
1 // pre: Given array X[] of sizeX elements
2 // n becomes the size of the heap.
3 // post: The array node A[] is sorted.
4
5 public void sort(node X[], int sizeX) {
6     buildUp(X, sizeX);
7     for (i = n; i>1; i--){
8         swap(A[1],A[i]);
9         n--;
10        heapify(1); }
11    copy A back to X;}
```

Note that the heap is both created and destroyed in the process of sorting using **buildUp** and **sort**.

The procedure invocation **swap(A[1], A[i])** needs to be adapted to your programming language.

Heaps—complexity of `buildUp`

- In a heap of n nodes call the total time to run `buildUp`, T .
- The height of the heap tree, $H = \log_2 n$.
- For the purposes of this derivation the height of a vertex is the number of edges that separate the vertex from the bottom level of the tree.
- At most $\lceil \frac{n}{2^{h+1}} \rceil$ vertices are of height h .
- The heights of vertices run from $0, 1, 2, \dots, H$.
- So, by running over all heights:

$$T \leq \sum_{h=0}^H h \frac{n}{2^{h+1}} = n \sum_{h=1}^H h 2^{-h-1}$$

Heaps—complexity of `buildUp`

- The heights of vertices run from $0, 1, 2, \dots, H$.
- So, by running over all heights:

$$T \leq \sum_{h=0}^H h \frac{n}{2^{h+1}} = n \sum_{h=1}^H h 2^{-h-1}$$

- It is easily shown that $\sum_{h=1}^H h 2^{-h-1} < 1$,
- and it follows that $T = O(n)$.

$$S_H = \sum_{h=1}^H h2^{-h-1} < 1$$

$$\begin{aligned} S_H &= \sum_{h=1}^H h2^{-h-1} \\ &= 1.2^{-2} + 2.2^{-3} + \dots + (H-1).2^{-H} + H.2^{-H-1} \\ 2S_H &= 1.2^{-1} + 2.2^{-2} + 3.2^{-3} + \dots + H.2^{-H} \end{aligned}$$

Subtracting S_H from $2S_H$ gives

$$\begin{aligned} S_H &= \underbrace{2^{-1} + 2^{-2} + 2^{-3} + \dots + 2^{-H}}_{0.111\dots 11_2} - H.2^{-H-1} \\ &= 0.111\dots 11_2 - H.2^{-H-1} \\ &= 1 - 2^{-H} - H.2^{-H-1} \\ &< 1 \end{aligned}$$

An n -node heap has height $H = \lfloor \log_2 n \rfloor$

Suppose that the root of a heap has height 0.

n	height	$\lfloor \log_2 n \rfloor$
1	0	0
2	1	1
3	1	1
4	2	2
5	2	2
6	2	2
7	2	2
8	3	3
9	3	3
10	3	3
11	3	3
12	3	3
13	3	3
14	3	3
15	3	3
16	4	4
31	4	4
32	5	5
64	6	6
127	6	6
255	7	7
256	8	8
511	8	8
1023	9	9
1024	10	10
2048	11	11

- A heap is *full* when it has exactly $2^{h+1} - 1$ nodes, where h is the height of the heap.
- Adding a node to a full heap increases its height by 1.
- The height of the i^{th} node in a heap is $\lfloor \log_2 i \rfloor$.

At most $\lceil \frac{n}{2^{h+1}} \rceil$ vertices are of height h

Suppose that the root of a heap has height 0.

n	height	$\lceil \frac{n}{2^{h+1}} \rceil$
1	0	1
2	1	1
3	1	2
4	1	2
5	1	3
6	1	3
7	1	4
8	1	4
9	1	5
10	1	5
11	1	6
12	1	6
13	1	7
14	1	7
15	1	8
16	1	8
31	1	16
32	1	16
64	1	32
127	1	64
255	1	128
256	1	128
511	1	256
1023	1	512
1024	1	512
2048	1	1024

- A heap is *full* when it has exactly $2^{h+1} - 1$ nodes, where h is the height of the heap.
- Adding a node to a full heap increases its height by 1.
- The height of the i^{th} node in a heap is $\lfloor \log_2 i \rfloor$.
- The heights run from $1, 2, \dots, H$.
- At most $\lceil \frac{n}{2^{h+1}} \rceil$ vertices are of height h .

Heaps—applications

Heaps have many applications in other algorithms.²

Sorting	Heap sort is optimal $\sim 2N \log N$
Data compression	Combine smallest two elements in Huffman coding
Event-driven simulation	Put most important request first
Physics simulation	Next particle to collide
Dijkstra's shortest path	Finds closest node
Prim's MST algorithm	Choosing edge with minimal weight
Number theory	Sum of powers
Artificial intelligence	A^* search
Statistics	Maintain largest K values in a sequence
Finding the median	Compare using quicksort for median
Operating systems	Prioritize interrupts, load balancing
Discrete optimization	Bin packing, scheduling
Bayesian spam filtering	Filter messages with highest spam weight
Reducing roundoff error	How is heap used?

²The list was inspired by [Sedgewick and Wayne, 2011]

At most $\lceil \frac{n}{2^{h+1}} \rceil$ vertices are of height h

Now we vary the heights for a fixed value of n .

Suppose that the root of a heap has height 0.


n	height	$\lceil \frac{n}{2^{h+1}} \rceil$
4096	0	2048
4096	1	1024
4096	2	512
4096	3	256
4096	4	128
4096	5	64
4096	6	32
4096	7	16
4096	8	8
4096	9	4
4096	10	2
4096	11	1
4096	12	0

$$\sum_{h=1}^H h 2^{-h-1} = 4083.$$

Which is clearly *less* than 4096.

- A heap is *full* when it has exactly $2^{h+1} - 1$ nodes, where h is the height of the heap.
- Adding a node to a full heap increases its height by 1.
- The height of the i^{th} node in a heap is $\lfloor \log_2 i \rfloor$.
- The heights run from $1, 2, \dots, H$.
- At most $\lceil \frac{n}{2^{h+1}} \rceil$ vertices are of height h .

A complete binary search tree




A tree with h levels, has $N = 2^h - 1$ nodes, and $h = \log_2(N + 1)$.

The number of probes to visit each node once in a complete BST is

$$T = \sum_{r=1}^h r 2^{r-1}, \text{ and } 2T = \sum_{r=1}^h r 2^r$$

$$T = 2T - T = h 2^h - \sum_{r=0}^{h-1} 2^r = h 2^h - 1 + 2^h$$

A complete binary search tree




A tree with h levels, has $N = 2^h - 1$ nodes, and $h = \log_2(N + 1)$.

The average search time, S , in a full, balanced, i.e. a complete BST is:

$$\begin{aligned}
 S &= \frac{\text{total number of probes}}{\text{number of nodes}} \\
 &= \frac{1 - 2^{h-1} + h2^h}{2^h - 1} \text{ in terms of } h, \\
 &= \frac{-N + (N + 1) \log_2(N + 1)}{N} \text{ in terms of } N.
 \end{aligned}$$

A complete binary search tree




A tree with h levels, has $N = 2^h - 1$ nodes, and $h = \log_2(N + 1)$.


The average search time, S , in a full, balanced, i.e. a complete BST is:

$$\begin{aligned}
 S &= \frac{\text{total number of probes}}{\text{number of nodes}} \\
 &= \frac{1 - 2^{h-1} + h2^h}{2^h - 1} \text{ in terms of } h, \\
 &= \frac{-N + (N + 1) \log_2(N + 1)}{N} \text{ in terms of } N.
 \end{aligned}$$

Search trees—a hierarchy




Search trees



Search trees implement the dictionary abstract data structure (ADT) and have the methods

- `boolean isFull()` and `boolean isEmpty()`
- `node find(k)` as well as `iterator findAll(k)`
- `boolean insert(k, data)` or `boolean insert(n)`
- `node remove(k)`.


Binary search trees (BSTs)



Binary search trees (BSTs):

- Unbalanced BSTs may behave badly, e.g. a tree built up from a sorted list presents $O(n)$ times.
- *AVL* trees are balanced BSTs with $O(\log n)$ times.
- Splay trees are BSTs that dynamically move the most accessed nodes to the top of the tree. They have an amortized running time: After m operations the average time to access item k —which has been accessed $f(k)$ times— is $O(\log(m/f(k)))$.

Multi-way search trees



Multi-way search trees:


- 2 - 4 trees or 2 - 3 - 4 trees put all the leaf nodes at the same level.
- *Red-black* trees correspond to 2 - 4 trees, use $O(n)$ space and have $O(n \log n)$ worst case insertion, deletion and lookup times.
- *B-trees* are multi-way trees that optimize input/output operations for large data bases on disks.

- The most difficult method to program is `delete(k)` which removes the node with key `k` from a BST:
 - Removing node `p` with at least one node that has a `null` child, entails moving the other child pointer up to `p.parent`'s child that points to `p` and so doing discarding `p`.
 - To remove a node `p` with two children: Find its biggest (smallest) child in its left (right) subtree, copy that biggest (smallest) child over the node to be deleted and then remove the node that was copied as usual. The copied node cannot have two children.
- The other methods are pretty elementary.
- Retrieval in a balanced BST takes $O(\log n)$ time and in an unbalanced BST retrieval can be as bad as $O(n)$.
- So, much effort is made to keep BSTs balanced.


AVL trees

- The two mathematicians G.M. Adel'son-Vel'skii and Y.M. Landis invented AVL trees in 1962. Georgii Maksimovich Adelson-Velskii and Yevgenii Michailovich Landis.
- An AVL tree has the *height-balanced property* or *AVL property*:
The children of every internal node n have heights that differ by at most 1.

root:



root:



- Since AVL trees are balanced the retrieval of nodes is at worst $\log_2 n$


AVL trees—height of nodes

Consider a complete BST.

Level	Nodes on level	Total
1	1	1
2	2	3
3	4	7
4	8	15
5	16	31
6	32	63
...	...	...
$h - 1$	2^{h-2}	$2^{h-1} - 1$
h	2^{h-1}	$2^h - 1$


- A complete height- h BST has $n = 2^h - 1$ nodes, so that $h = \log_2(n + 1)$.
- Note that one more than half of its nodes are at level h .
Since $2^h - 1 = 2^{h-1} + [2^{h-1} - 1]$.
- The maximum height of a complete binary tree is $h = \log_2(n + 1)$,
i.e. height $= O(\log n)$.
- An AVL tree is “almost” complete so its maximum height is $O(\log n)$.

Turning a non-AVL tree into an AVL tree



Node *B* does not violate the AVL property, and neither does node *C*, but node *A* has a child *B* with height $h+2$ and a child *C* with height h . The difference in the heights is 2 which violates the AVL property.

After a single rotation, in which *A* becomes *B*'s right child and *T*₂ becomes *A*'s left child; the tree stays a BST and it becomes an AVL tree:




Given only *A*, the code is:

```
B = A.left;  
A.left = B.right;  
B.right = A;
```


None of the nodes now violates the AVL property.

Turning a non-AVL tree into an AVL tree—mirror



Node B does not violate the AVL property, and neither does node C , but node A now has a child C with height $h + 2$ and a child B with height h . The difference in the heights is 2 which violates the AVL property.

After a single rotation, A becomes C 's *left* child and T_3 becomes A 's *right* child; the tree stays a BST and it becomes an AVL tree:




Given only **A**, the code is:

```
1  C = A.right;  
2  A.right = C.left;  
3  C.left = A;
```


None of the nodes now violates the AVL property.

AVL trees—double rotations



Node *A* violates the AVL property. The other internal nodes do not. A double rotation is required to turn this into an AVL tree.

After a double rotation, in which *D* becomes the root, the tree stays a BST and becomes an AVL tree:




Given only *A*, the code is:

```
1  B = A.left;
2  D = B.right;
3  B.right = D.left;
4  D.left = B;
5  A.left = D.right;
6  D.right = A;
```


D becomes the root and none of the nodes now violates the AVL property.

AVL trees—double rotations: same code



Insertion at T_3 instead of at T_2 results in: Node A violates the AVL property. The other internal nodes do not. A double rotation is required to turn this into an AVL tree.

After a double rotation, in which D becomes the root, the tree stays a BST and becomes an AVL tree:




Given only A , the code is:

```
1  B = A.left;  
2  D = B.right;  
3  B.right = D.left;  
4  D.left = B;  
5  A.left = D.right;  
6  D.right = A;
```


D becomes the root and none of the nodes now violates the AVL property.

AVL trees—double rotations: mirror



Insertion at T_2 now results in: Node A violating the AVL property. The other internal nodes do not. A double rotation is required to turn this into an AVL tree.

After a double rotation, in which D becomes the root, the tree stays a BST and becomes an AVL tree:




Given only A , the code is:

```
1  C = A.right;  
2  D = C.left;  
3  A.right = D.left;  
4  D.left = A;  
5  C.left = D.right;  
6  D.right = C;
```


The tree with D as root is now an AVL tree.

AVL trees—double rotations: mirror



An insertion at T_3 now results in node A violating the AVL property. The other internal nodes do not. A double rotation is required to turn this into an AVL tree.

After a double rotation, in which D becomes the root, the tree stays a BST and becomes an AVL tree:



Given only A , the code is:

```
1  C = A.right;  
2  D = C.left;  
3  A.right = D.left;  
4  D.left = A;  
5  C.left = D.right;  
6  D.right = C;
```

The tree with D as root is now an AVL tree.

Insertion into an AVL tree

- Suppose the new node is called **n**
- **insert** proceeds as usual, but the balance factor **bf** of each traversed node is updated in the appropriate direction:

```
1  if (here.key < (here.parent).key)
2      ++bf;
3  else if (here.key > (here.parent).key)
4      --bf;
```

- Remember the last node with a non-zero balance factor—the critical node.
- Update until a leaf is found for the new node **n**.
- From this position work bottom-up to the critical node readjusting balance factors and perform a rotation if necessary.

Deletion from an AVL tree

- Deletion is more complex.
- The tree can sometimes only be rebalanced after $O(\log n)$ rotations, where n is the number of nodes in the tree.
- The worst case is $O(\log n)$ rotations.


Summary

- AVL trees are efficient. It has been shown that they require $1.45 \times$ comparisons of optimal trees.
- Experiments have shown that AVL trees yield an average search time of $\log_2 n + 0.25$ comparisons.
- A disadvantage is the extra space, and bookkeeping required for handling the balance factors.

Splay trees

- Splay trees do not force balance.
- After each access a *move-to-root* operation is applied.
- This *splaying* operation is applied at the bottom-most node n reached during insertion, deletion or lookup.
- Splaying guarantees $O(\log n)$ amortized behaviour.
- No balance, height or colour labels are required.
- Splay trees were first described by Sleator and Tarjan in 1983.

Splay tree rotations—right rotation



Rotation of the edge joining node x and node y .


Each triangle represents a subtree.

Note that the subtrees are in the same $\ell N r$ order before and after a rotation.

Rotation does not affect the ordering of the nodes in the subtrees, A , B , and C . Assuming the figure on the *left* and pointers to nodes x and y , the code to execute a *right* rotation is:

```
1 subtree B = x.right;
2   x.right = y;
3   y.left = B;
4   y.parent.(left or right) = x;
5   x.parent = y.parent;
6   y.parent = x;
```

Splay tree rotations—left rotation




Rotation of the edge joining node x and node y . Triangles represent subtrees. Subtrees are in the same $\ell N r$ order before and after a rotation. Rotation preserves the ordering of the nodes in the subtrees, A , B , and C . Assuming the figure on the *left* and pointers to nodes x and y , the code to execute a *right* rotation is:

```
1 subtree B = x.left;  
2   x.left = y;  
3   y.right = B;  
4   y.parent.(left or right) = x;  
5   x.parent = y.parent;  
6   y.parent = x;
```


Splay tree rotations

- A single rotation around $x - y$ embedded inside the tree.




Splay trees

Before the *zig* operation:



After the *zig* operation:



- Splay trees restructure the tree so that each time a node x is accessed it is “*moved to the root*” by doing rotations bottom-up along the path of access by one three splaying steps until x reaches the root.
 1. *zig*, which reduces x ’s height by one.
 2. *zig-zig*, and
 3. *zig-zag* and their mirrored symmetries.
- *Zig-zig* and *zig-zag* reduce x ’s height by two, i.e. each splay x moves 1 or 2 closer to the root.

The Splaying steps *zig*, *zig-zig* and *zig-zag*

- To move x to the root:

Case 1: *zig*: If $y = \text{parent}(x)$ is the root, rotate along the edge $x - y$.


Case 2: *zig-zig*: If $y = \text{parent}(x)$ is not the root and x and y are *both* left children, or *both* right children: let $z = \text{parent}(y)$, first rotate the edge $y - z$ and then rotate the edge $x - y$.

Case 3: *zig-zag*: If $y = \text{parent}(x)$ is not the root and x is a left child and y a right child, or vice versa, letting $z = \text{parent}(y)$, first rotate the edge $x - y$ and afterwards rotate $x - z$.

- Splaying a node x at depth d takes $\Theta(d)$ time which is proportional to the time that it takes to access x .
- Note that each splay step preserves the inorder properties of the tree.

Splay trees—a left-left *zig-zig* step-by-step


A *zig-zig* is done using two rotations: If $y = \text{parent}(x)$ is not the root and x and y are *both* left³ children: let $z = \text{parent}(y)$, first left rotate the edge $y - z$ and then *left* rotate the edge $x - y$.



³Or *both* children are right. See the next slide.

Splay trees—a right-right *zig-zig* in 2 steps


A *zig-zig* is done using two rotations: If $y = \text{parent}(x)$ is not the root and x and y are *both* right⁴ children: let $z = \text{parent}(y)$, first rotate the edge $y - z$ and then rotate the edge $x - y$.



⁴Or *both* children are left. See the previous slide.

Splay trees—a left-right *zig-zag* in 2 steps


A *zig-zag* employs two rotations: If $y = \text{parent}(x)$ is not the root and x is a left child and y a right child⁵ letting $z = \text{parent}(y)$, first rotate the edge $x - y$ and afterwards rotate $x - z$.



⁵Or vice versa, see next slide.

Splay trees—a right-left *zig-zag* in 2 steps


A *zig-zag* is done in two rotations: If $y = \text{parent}(x)$ is not the root and x is a right child and y a left child⁶ letting $z = \text{parent}(y)$, first rotate the edge $x - y$ and afterwards rotate $x - z$.



⁶Or vice versa, see previous slide.

Splay trees—a left-left *zig-zig* step-by-step

A direct *zig-zig* is done using two rotations: If $y = \text{parent}(x)$ is not the root and x and y are *both* left⁷ children: let $z = \text{parent}(y)$, first rotate the edge $y - z$ around y and then rotate the edge $x - y$ around x .




Assume that we know x , y and z . Why can we assume this? The code to do a zig-zig:

```
1 rightRotate(y, z);  
2 rightRotate(x, y);
```

⁷Or *both* children are right. See the next slide.

Splay trees—code for left rotation

Consider the figure




Assume that we know x and y . The code for a left rotation:

```
1 rotateLeft(tree x, tree y) {
2   // pre: y == x.parent != null && x != null
3   // && x.key >= y.key
4   tree C = x.left;
5   x.left = y;
6   y.right = C;
7   x.parent = y.parent;
8   y.parent = x;
9   if (x.parent != null)
10      if ((x.parent).key > x.key)
11         (x.parent).left = x;
12      else
13         (x.parent).right = x; }
```

Splay trees—code for right rotation

Consider the figure



Assume that we know x and y . The code for a right rotation:

```
1 rotateRight(tree x, tree y) {
2   // pre: x != null
3   // && y == x.parent != null
4   // && x.key < y.key
5   tree B = x.right
6   x.right = y;
7   y.left = B;
8   x.parent = y.parent;
9   y.parent = x;
10  if ((x.parent).key > x.key)
11    (x.parent).left = x;
12  else
13    (x.parent).right = x; }
```


Splay trees—code for automatic rotation

Do a *left* rotation when $x.key \geq (x.parent).key$, i.e. y is a *right* child—otherwise do a *right* rotation.

Left rotation:

```
1 rotateLeft(tree x, tree y) {
2 //pre: && x != null
3 // && y == x.parent != null
4 // && x.key >= y.key
5 tree C = x.left;
6   x.left = y;
7   y.right = C;
8   x.parent = y.parent;
9   y.parent = x;
10  if ((x.parent).key > x.key)
11    (x.parent).left = x;
12  else (x.parent).right = x; }
```

Right rotation:

```
1 rotateRight(tree x, tree y) {
2 // pre: x != null
3 // && y == x.parent != null
4 // && x.key < y.key
5 tree B = x.right
6   x.right = y;
7   y.left = B;
8   x.parent = y.parent;
9   y.parent = x;
10  if ((x.parent).key > x.key)
11    (x.parent).left = x;
12  else (x.parent).right = x; }
```

Only one argument is needed: the pivot of the rotation, x .

$y = x.parent$; // assuming that $x \neq \text{null}$

The temporary nodes **B** or **C** are not needed. Rotation is determined by $\text{arity}(x)$, i.e. the *leftness* or *rightness* of x , given by $x.key < y.key$.

zig, zig-zig and zig-zag with rotate(x)

- Splaying a node x is a sequence of rotations that moves it to the root.
- Move x to the root by repeating:

Case 1: *zig*: If $\text{parent}(x)$ is the root, $\text{rotate}(x)$.


Case 2: *zig-zig*: If $\text{parent}(x)$ is not the root and x and $\text{parent}(x)$ are both *left* children, or both *right* children: first the $\text{rotate}(\text{parent}(x))$ and then $\text{rotate}(x)$.

Case 3: *zig-zag*: If $\text{parent}(x)$ is not the root and x is a left child and $\text{parent}(x)$ a right child, or vice versa, first $\text{rotate}(x)$ and afterwards $\text{rotate}(x)$.

- Splaying a node x at depth d takes $\Theta(d)$ time which is proportional to the time that it takes to access x .
- Note that each splay step preserves the inorder properties of the tree.

Splay trees—a left-left *zig-zig* step-by-step

A *zig-zig* is done using two rotations: If $\text{parent}(x)$ is not the root and x and $\text{parent}(x)$ are both *right*⁸ children: first $\text{rotate}(x.\text{parent})$ and then $\text{rotate}(x)$.




It is called a “left-left” *zig-zig* because both rotations are to the left, i.e. counter-clockwise or anti-clockwise.

⁸Or both are *left* children. See the next slide.

Splay trees—a right-right *zig-zig* in 2 steps

A *zig-zig* is done using two *clockwise* rotations: If $\text{parent}(x)$ is not the root and x and $\text{parent}(x)$ are both *left*⁹ children: first $\text{rotate}(y)$ and then $\text{rotate}(x)$ again.




Notice that the code for a “right-right” *zig-zig* is the same as that for a “left-left” one. $\text{rotate}(x)$ decides the direction.

⁹Or both are *right* children. See the previous slide.

Splay trees—a left-right *zig-zag* in 2 steps


A *zig-zag* employs two rotations: If $\text{parent}(x)$ is not the root and x is a *left* child and $\text{parent}(x)$ a *right* child¹⁰: first $\text{rotate}(x)$ and then $\text{rotate}(x)$ again.



¹⁰Or vice versa, see next slide.

Splay trees—a right-left *zig-zag* in 2 steps


A *zig-zag* is done in two rotations: If $\text{parent}(x)$ is not the root and x is a *right* child and y a *left* child¹¹ first $\text{rotate}(x)$ and afterwards $\text{rotate}(x)$ again.



¹¹Or vice versa, see previous slide.

Splay trees—a left-left *zig-zig* step-by-step

A direct *zig-zig* is done using two rotations: If $y = \text{parent}(x)$ is not the root and x and y are *both* left¹² children: let $z = \text{parent}(y)$, first rotate the edge $y - z$ and then rotate the edge $x - y$.



Assume that we know x , y and z . Why can we assume this? The code to do a zig-zig:

```
1 leftRotate(y, z);
2 leftRotate(x, y);
```

This becomes

```
1 rotate(x.parent);
2 rotate(x);
```

¹²Or *both* children are right. See the next slide.

Splay trees—rotation around the pivot

```
1 rotate(tree x) {
2 // pre: x != null
3 tree y = x.parent;
4 if (y != null) // x is not the root
5     if (x.key < y.key) { // right rotation
6         y.left = x.right;
7         x.right = y;
8     }
9     else { // x is right child:
10        left rotation
11        y.right = x.left;
12        x.left = y;
13    }
14    if (y.parent != null)
15        if (y.key < (y.parent).key)
16            (y.parent).left = x;
17        else
18            (y.parent).right = x;
19    x.parent = y.parent;
20    y.parent = x; }
```


- In our pseudo-code we have used the expression `x.parent` instead of `getParent(x)` to clarify our coding.
- Next we tackle the code for `splay(x)` given only the tree node `x`.

Splay trees—splay(x)—1st version

```
1 void splay(tree x) {
2     while (x.parent != null) { // x is root
3         if ((x.parent).parent == null)
4             rotate(x); // x.parent is root: zig
5         // x.parent has a non-root parent
6         else if (x.key < (x.parent).key) // x is left child
7             if ((x.parent).key < ((x.parent).parent).key)
8                 rotate(x.parent); // left-left: zig-zig
9             else
10                rotate(x); // left-right: zig-zag
11        // x is a right child
12        else if ((x.parent).key < ((x.parent).parent).key)
13            rotate(x); // right-left: zig-zag
14        else { // right-right: zig-zig
15            rotate(x.parent); }
16        rotate(x); } }
```

Take care this has some errors. Can you spot them?

Splay trees—splay(x)—2nd version


```
1 void splay(tree x) {
2     // repeat while x is not the root node
3     while (x.parent != null) { // y -- x
4         if ((x.parent).parent != null) { // z -- y -- x
5             // x.parent has a non-root parent
6             if (x.key < (x.parent).key)
7                 // x is a left child
8                 if ((x.parent).key < ((x.parent).parent).key)
9                     // x.parent is a left child: left-left zig-zig
10                    rotate(x.parent);
11                else // left-right: zig-zag
12                    rotate(x);
13            else // x is a right child
14                if ((x.parent).key < ((x.parent).parent).key)
15                    // x.parent is a left child: right-left zig-zig
16                    rotate(x.parent);
17                else // right-right: zig-gig
18                    rotate(x);
19            rotate(x); } } }
```

Still error ridden. Can you spot the errors?

Splay trees—**splay(x)**—tidier


```
1 void splay(tree x) {
2     // repeat while x is not the root node
3     while (x.parent != null) { // y -- x
4         if ((x.parent).parent != null) // z -- y -- x
5             if (x.key < (x.parent).key)
6                 if ((x.parent).key < ((x.parent).parent).key)
7                     rotate(x.parent);
8                 else // left-right: zig-zag
9                     rotate(x);
10            else // right
11                if ((x.parent).key < ((x.parent).parent).key)
12                    rotate(x.parent);
13                else // right-left: zig-gig
14                    rotate(x);
15            rotate(x); } }
```

Splay trees—**splay(x)**—flow chart



delete(k) from a *splay tree*

- First `find(k)` and splay it.



- Next replace the root, k , with the *largest* key in the *left subtree*. T_2 or replace the root with the *smallest* key in the *right subtree*.
- Note that all the keys in T_1 are less than k and all the keys in T_2 are greater than k .
- Notice that the *greatest* key in T_1 is a *right* child but has no *right* child: so deleting or removing the *greatest* key from T_1 is simple: move its *left* key to its parent's *right* key.

Splay trees—complexity

- **Theorem:** Any sequence of m operations on an initially empty splay tree that never contains more than n nodes takes $O(m \log n)$
- Running time of each operation is proportional to time for splaying.
- Define $rank(v)$ as the $\log_2 n$, where n is the number of nodes in subtree rooted at v .
- Actual Costs: $zig = \$1$, $zig-zig = \$2$, $zig-zag = \$2$.
- Thus, cost for splaying a node at depth $d = \$d$.
- Assume that $rank(v)$ dollars is stored at each node v of the splay tree
- *Invariant.* For each node $v \geq \$rank(v)$.
- **Lemma 1:** Each splay operation requires us to invest at most $3 \log_2 n + 1$ new dollars, to maintain the invariant.


- See Goodrich and Tamassia for more slides in the SunLab at `.../.../notes/ds/SplayTrees.pdf` when you are in your home directory.

More up-to-date versions are on the web site that Goodrich and Tamassia maintain for the book.

These notes are superficial. Consult the book, pages 440–550 (4th edition).

Multiway trees

Before discussing $(2 - 4)$ trees we first consider multiway trees. The figure below is a multiway tree.



An n entry multiway tree has $n + 1$ leaves—external nodes. It is proved by induction.

$P[1]$: A multiway tree with one node has 2 leaves.


$P[k] \rightarrow P[k + 1]$: Assume $P[k]$, i.e. T_k is a multiway tree with k entries and has $k + 1$ leaves. Inserting a node into T_k at any leaf will replace one leaf and add two, i.e. the new tree T_{k+1} will have $k + 2$ leaves. So, $\forall n P[n]$.

Red-black trees


- A $(2 - 4)$ tree has an update—search, insert and remove—time of $O(\log n)$.
- A *red-black tree* has an update time of $O(\log n)$, i.e. it needs a constant number of changes to do an update.
- A *red-black tree* is a BST with nodes coloured *red* or *black* using the following rules:
 1. The root is black.
 2. Every external node is black.
 3. The children of a red node are black.
 4. External nodes have the same *black depth*. i.e. the number of black ancestors minus one.
- The following slides shows how to convert nodes from a $2 - 4$ tree into a *red-black tree*

Converting a 2 – 4 tree into a red-black tree



The figures below illustrate the conversion of nodes from a 2 – 4 tree into those for a red-black tree.




$\Rightarrow$



$\Rightarrow$




$\Rightarrow$




Converting a 2 – 4 tree into a red-black tree

We convert the following into a *red-black* tree:




The *red-black* tree follows.




Converting a 2 – 4 tree into a *vertical-horizontal-red-black* tree

We convert the following into a *red-black* tree:




The vertical-horizontal *red-black* tree looks more like the original (2 – 4) tree.




Deletion from a *red-black* tree

Consider removing 27 from the *red-black* tree below.




The resultant tree below has lost its black-depth property. Since 25's right child—a *double black*—now has a black depth of 2, it is no longer a *red-black* tree.




It is easy to repair the *double-black* error by restructuring 23 – 25 – 24 as three separate black 2-nodes with 24 on top and the black-depth property is restored.

Repairing the *double black* at 25

It is easy to repair the *double-black* error by restructuring $23 - 25 - 24$ first into the 4-node $23 - 24 - 25$




The leaves of 23 and 24 now have the incorrect *black depth* of 2. Subsequently, this 4-node $23 - 24 - 25$ is recoloured into three separate black 2-nodes with 24 on top. All the leaves now have a *black depth* of 3.




Non-trivial deletions from this *red-black* tree

Deletion of either 23 or 25 from this *red-black* tree cause a double black to hang from node 24. Let us remove 23. The double black hangs from node 24.




The red-black tree is repaired fixed by first amalgamating 24 and 24 into a 3-node and then the next step becomes obvious.




Easy deletions from this *red-black* tree

Amalgamating 24 and 24 into a 3-node makes the next step seem obvious.



Move the root 22 down to form the 4-node 5 – 10 – 22.



Since they are all red, removing 3, 8, 11, 17 or 24, has no effect on the depths of the leaves of the tree.

(a, b) Trees

- A tree where all the nodes have *at least* a children and *at most* b children is a B-tree when
 $2 \leq a \leq (b + 1)/2$.
All the leaves have the same depth.
- The root may have 0 children when the tree is empty.
- Otherwise the root must have at least 2 children.
- If there are less than a entries, they all lie in the root node, such that the root may have less than a children.

The height of an (a, b) tree T with n entries is $\Omega(\log_b n)$ and $O(\log_a n)$

- The height of an (a, b) tree is bounded from above by $O(\log_a n)$ and from below by $\Omega(\log_b n)$.
- The number of external nodes n' of T is at least $2a^{h-1}$: If the root has only one entry and the height, $h = 1$, then there are 2 leaves. If $h = 2$ and the root has at least a children and each child on the second level has a single entry, then there are at least $2a$ children on the second level. The third level will have at least $2a^2$ leaves.
- The number of external nodes is *at most* b^h where h is the height of the tree.
- An n -entry mutiway tree has $n + 1$ leaves, so

$$\begin{aligned} 2a^{h-1} &\leq n + 1 \leq b^h \\ \log_2(2a^{h-1}) &\leq \log_2(n + 1) \leq \log_2(b^h) \\ 1 + (h - 1) \log_2 a &\leq \log_2(n + 1) \leq h \log_2 b \end{aligned}$$

Upper and lower bounds for the height of an (a, b) tree

Since

$$1 + (h - 1) \log_2 a \leq \log_2(n + 1) \leq h \log_2 b$$

it follows that

$$\begin{aligned} (h - 1) &\leq \frac{\log_2(n + 1) - 1}{\log_2 a} \\ h &\leq \frac{\log_2(n + 1) - 1}{\log_2 a} + 1 \\ &\doteq \log_a(n + 1) \\ h &\text{ is } O(\log_a n) \end{aligned}$$

and

$$\begin{aligned} h &\geq \frac{\log_2(n + 1)}{\log_2 b} \\ &= \log_b(n + 1) \\ h &\text{ is } \Omega(\log_b n) \end{aligned}$$

Searching in an (a, b) tree

- Each node of an (a, b) tree has at least $a - 1$ entries and at most $b - 1$ entries.
- Looking for an entry in a *node* using *any* search method will take $f(b)$ time.
- Searching for an entry in an (a, b) *tree* T will take at worst $O(f(b) \log_a n)$ time.
- Since b is constant, a is also, and the search time is $O(\log n)$, which is independent of the specific search method.

A B -tree of order d

- A B -tree of order d is a $(\lceil d/2 \rceil, d)$ tree.
- d is chosen such that $d - 1$ keys and d references can be stored in a single block of size B .
- So each time a node is accessed a single disk transfer is done.
- The most transfers that ever need to be done is $O(\log_a n) = O(\log_{\lceil d/2 \rceil} n)$ and it will occupy $O(n/B)$ blocks.

Overflow and underflow in a *B*-tree

- During **insertion** *overflow* might occur.
- When the number of entries in a node in which an insertion must be done is already d , an *overflow* occurs. The node is then split into two halves with $\lfloor (d+1)/2 \rfloor$ and $\lceil (d+1)/2 \rceil$ children respectively.
- During **deletion** *underflow* could occur.
- When the number of entries in a node from which a deletion must be done is only $\lceil d/2 \rceil$, an underflow occurs and some housekeeping must be done in which the entries of the underfull node is moved to one or more of its siblings.

Sorting—an $O(n^3)$ method

- Given a list of n objects $A[1], A[2], \dots, A[n]$ that can be totally ordered, the problem is to arrange them such that $A[1] \leq A[2] \leq A[3] \leq \dots \leq A[n]$.
- Suppose that the first $i - 1$ elements are in order, i.e. $A[1] \leq A[2] \leq A[3] \leq \dots \leq A[i - 1]$ insert the i -th one, $A[i]$ in its correct position relative to the elements we have already considered. Carry on until $i = n$.
- How do we start? Start with $i = 1$: the list, $A[1]$, of 1 element is already sorted.
- Proceed to $i = 2$ by inserting $A[2]$ into its correct position before or after $A[1]$. This may entail moving $A[1]$ down one position to $A'[2]$ to vacate $A'[1]$ for $A[2]$ and, finally, rename the elements giving the sorted list: $A'[1] \leq A'[2]$.
- The following step is to insert $A[3]$ into the sorted list $A'[1] \leq A'[2]$.

Sorting—an $O(n^3)$ method

- Insert $A[3]$ into the sorted list $A'[1] \leq A'[2]$.
- First make a copy: $now \leftarrow A[3]$, and then compare now one-by-one with $A'[1]$ and $A'[2]$.
- If $now \leq A'[j]$, where $j \in [1..i-1]$, where $i = 3$, then $A[i] \equiv now$ has found its place in the list. Starting at $A'[j]$, move all the elements in the list down one position and then copy: $A'[j] \leftarrow now$.
- Since now is a true copy of the element that was at $A'[3]$ it is not destroyed when $A'[2]$ overwrites $A[3]$. Throw away the primes ('s).
- At this stage: $A[1] \leq A[2] \leq A[3]$.
- Instead of proceeding with $A[4]$ we rather see how to handle $A[i]$.
- In a proof by mathematical induction show that a basis proposition $\mathcal{P}[0]$ holds, assume that $\mathcal{P}[k]$ is true and then if $\mathcal{P}[k+1]$ follows, the proposition $\mathcal{P}[i]$ holds for all integers i .

Sorting—an $O(n^3)$ method

- Assume $\mathcal{P}[i - 1]$: $A[1] \leq A[2] \leq A[3] \leq \dots \leq A[i - 1]$ insert the i -th element in its correct position in the list.
- Copy: $now \leftarrow A[i]$, and then compare now one-by-one with $A[j]$ where $j \in [1..i - 1]$. Find the first position j where $now \leq A[j]$.
- Move all the elements $A[j], A[j + 1], \dots, A[i - 1]$ down one position. Copy *from the back to the front* and then place now at the vacated $A[j]$.
- Why must we copy *back-to-front*?
- Now $A[1] \leq A[2] \leq A[3] \leq \dots \leq A[i]$.
- This completes the induction: we assumed $\mathcal{P}[i - 1]$ and showed how to do $\mathcal{P}[i]$ using the assumption. What was $\mathcal{P}[0]$?
- Carry on until $i = n$. The n element list will then be sorted. We now code the algorithm and analyse it.

Sorting—an $O(n^3)$ method—the code

- Set up a loop for the variable i . Start at 2. Why?

```
1  for (i = 2; i <= n; i++) {
```

- Create now and find $j \in [1..i - 1]$ for which $now \leq A[j]$.

```
1      now = A[i];  
2      for (j = 1; j <= i-1; j++)  
3          if (now <= A[j])  
4              insert(A, i, j); }
```

- The **void** procedure `insert(A, i, j)` inserts `A[i]` into the array `A[1..i]` at `j`, by moving each element down one place and inserting `now` at `A[j]`:

```
1 void insert (int A[], int i, int j) {  
2     int now = A[i], k;  
3     for (k = i; k >= j; k--)  
4         A[k] = A[k-1];  
5     A[j] = now; }
```

Sorting—an $O(n^3)$ method—the analysis

- The length of the `for (k ...` loop inside `insert` varies from 1 to $n - 1$. So it takes $O(n)$ time.
- The length of the `for (j ...` loop inside the `i` loop executes once, then twice, etc, up to almost n times. So it also done $O(n)$ times. The $O(1)$ stuff inside the `insert` procedure is done at most n times so it takes $O(n)$, thus the `for (j ...` loop will consume $O(n^2)$ time.
- Finally, the outermost loop is itself repeated about n times, so the complete sort algorithm requires $O(n^3)$ time to execute.
- This algorithm is a lesson in how *not* to sort.
- There are worse ways to sort. Can you think of sorting method that has exponential time complexity?
- We will next show another bad algorithm that takes $O(n^2)$ time.

Sorting—an $O(n^2)$ method

- Now we grow our sorted list by placing new elements one-by-one into their *final* position after one sweep of the remainder of the array.
- Given a list of n objects $A[1], A[2], \dots, A[n]$ that can be totally ordered, sorting is the problem to arrange them such that $A[1] \leq A[2] \leq A[3] \leq \dots \leq A[n]$.
- Suppose that the first $i - 1$ elements are in their *correct* positions, i.e. $A[1] \leq A[2] \leq A[3] \leq \dots \leq A[i - 1]$. Find the i -th element in the rest of the array so that, $A[i]$ will be in its final position when it is placed. Carry on until $i = n$.
- This is easy to do: Simply find the smallest element in the list $A[i..n]$ and when it is found swap it with whatever is at $A[i]$.
- After the swap $A[1..i]$ is completely sorted.
- Note this is again a disguised proof by mathematical induction that this sorting method is actually correct.

Sorting in $O(n^2)$ —the code

- The $i \in [1..n)$ in the outer loop points to the final position of the smallest element. It sets **smallest** to the $A[i]$, and preserves the index. The j -loop then iterates through all of the elements in $A[i+1..n]$ to get the smallest element.

```
1  for (i = 1; i <= n-1; i++) {  
2      smallest = A[i];  
3      indexSmallest = i;  
4      for (j = i+1; j <= n; j++) {  
5          if (smallest > A[j]){  
6              smallest = A[j];  
7              indexSmallest = j;} }  
8      A[i] = A[indexSmallest];  
9      A[indexSmallest] = smallest; }
```

- This time there are two nested loops of similar complexity to those of the previous sorting method. All the other lines are $O(1)$, so the algorithm has a complexity of $O(n^2)$.
- Although it improves the previous algorithm by an ‘order’ it is not good enough to earn our respect.

Merge sort—an $O(n \log n)$ algorithm

- Given a list of $to - from + 1$ objects:
 $A[from], A[from + 1], A[from + 2], \dots, A[to]$ that can be totally ordered, arrange them such that
 $A[from] \leq A[from + 1] \leq \dots \leq A[to]$.
- Divide the given list into two almost equal parts: $A[from..middle - 1]$ and $A[middle..to]$ and sort each part separately.
- After these sorting each of these two lists, *merge* them into one list $A[from..to]$.
- Suppose that sorting n elements takes $T(n)$, then $T(n) = 2T(n/2) + cn$ where cn is the cost of merging two lists.
- A one element list is already sorted, so $T(1) = 0$. When there are only two elements $T(2) = 2c$.
- Solving this, we will find that $T(n) = O(n \log n)$.

Merge sort is an $O(n \log n)$ algorithm

- Suppose that sorting n elements takes $T(n)$, then $T(n) = 2T(n/2) + cn$ where cn is the cost of merging two lists.
- Now find a closed form for $T(n)$ by *telescoping*.

$$\begin{aligned}T(n) &= 2T(n/2) + cn \\&= 2(2T(n/2^2) + cn/2) + cn = 2^2T(n/2^2) + 2cn \\&= 2^3(2T(n/2^3) + cn/2^2) + 2cn = 2^3T(n/2^3) + 3cn \\&= 2^4(2T(n/2^4) + cn/2^3) + 3cn = 2^4T(n/2^4) + 4cn \\&= \dots \\&= 2^i T(n/2^i) + icn\end{aligned}$$

$$\begin{aligned}T(n) &= 2^i T(1) + icn, && \text{when } n = 2^i, \text{ so } i = \log_2 n. \\&= icn, && \text{since } T(1) = 0. \\&= cn \log n\end{aligned}$$

When $2^i = n$, i.e. $i = \log_2 n$, then $T(n) = icn$ and $T(n) = cn \log_2 n$,
i.e. $T(n)$ is $O(n \log n)$.

mergesort—the $O(n \log n)$ algorithm

- Given a list $A[from..to]$ Divide it into two almost equal parts at the midpoint p : $A[from..p]$ and $A[p + 1..to]$ and sort each part separately.
- After these sorting each of these two lists, *merge* them into one list $A[from..to]$.

```
1 public void mergeSort(int A[], int from, int to) {  
2     int p;  
3     if (from + 1 >= to) {  
4         if (A[from] <= A[to])  
5             return;  
6         swap(A, from, to); }  
7     else {  
8         p = (from+to)/2;  
9         mergeSort (A, from, p);  
10        mergeSort (A, p+1, to);  
11        merge (A, from, p, to); } }
```

- We will now discuss the merging process.

The merging process—start merging

- Put $i \leftarrow from$, $j \leftarrow to$, and put $k \leftarrow 0$. Compare $A[i]$ with $A[j]$ and store the smallest into $merged[k]$. Update i or j , and k appropriately.

```
1 public void merge(int A[], int from, int p, int to) {  
2   int merged[] = new int [to - from + 1];  
3   int i = from, j = p+1, k = 0;  
4   while (i <= p && j <= to)  
5     if (A[i] < A[j])  
6       merged[k++] = A[i++];  
7   else  
8     merged[k++] = A[j++]; ...}
```


- At some stage either $i > p$, or $j > to$.
- We will first discuss the case where $i > p$, then all of $A[from..p]$ and only some of $A[p + 1..to]$ has been copied into $merged[0..k]$.
- The tail end of $A[]$, namely $A[k + 1..to]$ is in place and only $merged[0..k]$ needs to be copied into its right place at $A[from..from + k]$.

Merging— i runs out first, i.e. $i > p$


- At some stage either: $i > p$, or $j > to$.

Assume $i > p$,

then all of $A[from..p]$ and only some of $A[p+1..to]$ has been copied into $merged[0..k]$.




- Notice that, here, j coincides with $from + k + 1$.
- To complete the sorting of $A[from..to]$, $merged[0..k]$ must be copied back into $A[from..j-1]$ since $A[j..to]$ is already in place.




The entire $A[from..to]$ is now in sorted order.

Merging—the code when i runs out first

- We will only copy the contents of $merged[0..k]$ back into $A[from..j-1]$ since $A[j..to]$ is already in place.




```
1 // Copy merged[] and A[i..p] or A[j..to] back to A[]
2 if (i == p+1) { // rest of A[j..to] already in place
3     r = from; s = 0;
4     while (r < j) // OR while (s <= k)
5         A[r++] = merged[s++]; }
6 else ...
```




The entire $A[from..to]$ is now in sorted order.

Merging— j runs out first

- Put $i \leftarrow from$, $j \leftarrow to$, and put $k \leftarrow 0$. Compare $A[i]$ with $A[j]$ and store the smallest into $merged[k]$. Update i or j , and k appropriately.
- At some stage $i > p$, or $j > to$. Assume $j > to$, then all of $A[p+1..to]$ and some of $A[from..p]$ has been copied into $merged[0..k]$.



- First copy the contents of $A[i..p]$ into $A[k+1..to]$:




Only $A[from+k+1..to]$ is now in its final position.

Completing the merge when $j > to$

```
1  . . .
2  else { // first park A[i..p] into A[from]k+1..to]
3      r = to; s = p;
4      while (s >= i)
5          A[r--] = A[s--];
6      // move merged[0..k] to A[from..from+k]
7      r = from; s = 0;
8      while (s <= k)
9          A[r++] = merged[s++]; }
```

- To complete the sorting of $A[from..to]$ the contents of $merged[0..k]$ must be copied back into $A[from..from + k]$ since $A[from + k + 1..to]$ is already in place.



The entire $A[from..to]$ is now in sorted order.

mergeSort—the code

```
1 public void mergeSort(int A[], int from, int to) {  
2   // in : A[from..to] array of elements in any order  
3   // int from, and int to, within bounds  
4   //  
5   // out: Array[from..to] in sorted order  
6   int p;  
7   if (from + 1 >= to) {  
8     if (A[from] <= A[to])  
9       return;  
10    swap(A, from, to);  
11    return; }  
12   else {  
13     p = (from+to)/2;  
14     mergeSort (A, from, p);  
15     mergeSort (A, p+1, to);  
16     merge (A, from, p, to);  
17     return; } }
```

Code for merge

```
1 public void merge(int A[], int from, int p, int to) {
2 // in : A[from..p] array of elements in sorted order
3 // A[p+1..to] array of elements in sorted order
4 // int from, int p, and int to, within bounds
5 // out: returns array A[from..to] in sorted order
6 //
7 int merged[] = new int [to -from + 1], k = 0,
8     i = from, j = p+1, r, s;
9
10 while (i <= p && j <= to)
11     if (A[i] < A[j])
12         merged[k++] = A[i++];
13     else
14         merged[k++] = A[j++];
15 // Copy merged[] and A[i..p] or A[j..to] back to A[]
16 if (i == p+1) { // rest of A[j..to] already in place
17     r = from; s = 0;
18     while (s < k)
19         A[r++] = merged[s++]; }
20 else { // first park A[i..p] into A[..to]
21     r = to; s = p;
22     while (s >= i)
23         A[r--] = A[s--];
24     r = from; s = 0;
25     while (s < k)
26         A[r++] = merged[s++]; } }
```


Code for swap

The code for **swap**:

```
1 public void swap (int A[], int x, int y) {
2   int keep = A[x];
3   A[x] = A[y];
4   A[y] = keep; }
```

In hindsight we see that the code required to copy $merged[0..k-1]$ back into $A[from..from+k-1]$ is duplicated. So **merge** is improved as follows by altering the code of the if statement:

```
1   if (i != p+1) { // first park A[i..p] into A[..to]
2     r = to; s = p;
3     while (s >= i)
4       A[r--] = A[s--]; }
5   r = from; s = 0;
6   while (s < k)
7     A[r++] = merged[s++];
```

Note:

- The copying back of the merged array $merged[0..k]$ can be replaced with the copying of a few links if the $A[]$ is stored as a linked list.

Quicksort—another $O(n \log n)$ sort

- Quicksort was published in the Computer Journal in 1962 by Tony Hoare. He invented the **case** statement.
‘Partition’ $A[from..to]$ into two parts: first select $pivot \in A[from..to]$ and save its *index*. Next, move all elements $< pivot$ to the left of $pivot$ and all the others to its right. Finally place $pivot$ at its correct position p in the final list, then sort $A[from..p - 1]$ and $A[p + 1..to]$ similarly.

```
1 public void quickSort(int A[], int from, int to) {  
2   // in : A[from..to] array of elements in any order  
3   int p;  
4   if (to < 0  from >= to) return; // nothing to do  
5   if (from - to == 1) // A[from..to] == 2  
6       if (A[from] < A[to]) return;  
7       swap(A, from, to);  
8       return; }  
9   else {  
10      p = partition (A, from, to);  
11      quickSort (A, from, p-1);  
12      quickSort (A, p+1, to); } }
```

partition—the $O(n)$ heart of quicksort

‘Partition’ $A[from..to]$ into two parts: first select a $pivot \in A[from..to]$ and save its *index*. Next, move all elements $< pivot$ to the left of $pivot$ and all the others to its right. Finally place $pivot$ at its correct position p in the final list.

```
1 public int partition(int A[], int from, int to) {  
2     int i, j, pivot = A[from], lower=from+1, upper=to;  
3     while (lower < upper) {  
4         while (A[lower] < pivot && lower < upper)  
5             lower++;  
6         while (pivot <= A[upper] && lower < upper)  
7             upper--;  
8         if (lower < upper) swap(A, lower, upper); }  
9     swap(A, from, lower - 1);  
10    return upper; }
```

Quicksort—time complexity, $T(n)$

- If *pivot* always partitions the $A[from..to]$ into two equal parts then $T(1) = 0$, $T(2) = 1$, and $T(n) = 2T(n/2) + O(n)$ implying that $T(n) = n \log n$. Of course this hardly ever happens.
- If the one side of *pivot* is nearly always empty then there will be $O(n)$ partitioning steps taking $O(n)$ each time, giving $O(n^2)$.
- This happens when choosing the *pivot* as $A[from]$ and when $A[from..to]$ is almost sorted. Choosing the *pivot* randomly will reduce the likelihood of this happening.
- But $T(n) = O(n^2)$ in the *worst* case.
- What is the average case?
- Suppose that every $A[i] \in A[from..to]$ is equally likely to be chosen as the pivot, then the running time of quicksort, if i is the i -th smallest element, is $T(n) = n - 1 + T(i - 1) + T(n - i)$.

Quicksort—average running time

- Suppose that every $A[i] \in A[from..to]$ is equally likely to be chosen as the pivot, then the running time of quicksort, if i is the i -th smallest element, is $T(n) = n - 1 + T(i - 1) + T(n - i)$.
- To calculate the average, consider:

$$T(n) = n - 1 + T(1 - 1) + T(n - 1)$$

$$T(n) = n - 1 + T(2 - 1) + T(n - 2)$$

$$T(n) = n - 1 + T(3 - 1) + T(n - 3)$$

...

$$T(n) = n - 1 + T(i - 1) + T(n - i)$$

...

$$T(n) = n - 1 + T(n - 1) + T(n - n)$$

Now sum these values and divide by n .

Quicksort—average running time

- Given that each element was uniform randomly chosen the average running time is:

$$T(n) = n - 1 + T(1 - 1) + T(n - 1)$$

$$T(n) = n - 1 + T(2 - 1) + T(n - 2)$$

$$T(n) = n - 1 + T(3 - 1) + T(n - 3)$$

...

$$T(n) = n - 1 + T(i - 1) + T(n - i)$$

...

$$T(n) = n - 1 + T(n - 1) + T(n - n)$$

$$nT(n) = n(n - 1) + \sum_{i=1}^n (T(i - 1) + T(n - i)), \text{ so that}$$

$$T(n) = n - 1 + \frac{1}{n} \sum_{i=1}^n (T(i - 1) + T(n - i))$$

$$= n - 1 + \frac{1}{n} \sum_{i=1}^n T(i - 1) + \frac{1}{n} \sum_{i=1}^n T(n - i)$$

$$= n - 1 + \frac{2}{n} \sum_{i=0}^{n-1} T(i) \quad \text{A so-called recurrence with full history.}$$

$$= O(n \log n)$$

Closed form of a recurrence with full history

$T(n) = c + \sum_{i=1}^{n-1} T(i)$, where c is a constant and $T(1)$ is given, is called a *recurrence relation with full history*, solved by eliminating the history.

First observe that $T(n+1) = c + \sum_{i=1}^n T(i)$

Their difference is $T(n)$, i.e. $T(n+1) - T(n) = T(n)$, giving $T(n+1) = 2T(n)$.

Now $T(2) = T(1) + c$, so by telescoping,

$$\begin{aligned} T(n+1) &= 2T(n) \\ &= 2^2 T(n-1) \\ &= 2^3 T(n-2) \\ &\quad \dots \\ &= 2^{n-2} T(3) \\ &= 2^{n-1} T(2) \\ &= 2^{n-1} (T(1) + c) \end{aligned}$$

Now consider the recurrence relation with full history

$$T(n) = n - 1 + \frac{2}{n} \sum_{i=0}^{n-1} T(i).$$

Closed form of $T(n) = n - 1 + \frac{2}{n} \sum_{i=0}^{n-1} T(i)$

$$T(n) = n - 1 + \frac{2}{n} \sum_{i=0}^{n-1} T(i), \quad n \geq 2 \quad (1)$$

$$T(n+1) = (n+1) - 1 + \frac{2}{n+1} \sum_{i=0}^n T(i), \quad n \geq 2 \quad (2)$$

Multiply Equation(1) by n and Equation(2) by $n+1$ giving:

$$nT(n) = n(n-1) + 2 \sum_{i=0}^{n-1} T(i), \quad n \geq 2 \quad (3)$$

$$(n+1)T(n+1) = n(n+1) + 2 \sum_{i=0}^n T(i), \quad n \geq 2 \quad (4)$$

Subtract Equation (3) from Equation (4):

Closed form of $T(n) = n - 1 + \frac{2}{n} \sum_{i=0}^{n-1} T(i)$

$$nT(n) = n(n-1) + n \sum_{i=0}^{n-1} T(i), \quad n \geq 2 \quad (3)$$

$$(n+1)T(n+1) = n(n+1) + 2 \sum_{i=1}^n T(i), \quad n \geq 2 \quad (4)$$

Equation (4)–(3):

$$\begin{aligned} (n+1)T(n+1) - nT(n) &= n(n+1) - n(n-1) + 2T(n) \\ &= 2n + 2T(n), \quad n \geq 2 \end{aligned}$$

Giving,

$$\begin{aligned} (n+1)T(n+1) &= (n+2)T(n) + 2n, \\ T(n+1) &= \frac{n+2}{n+1}T(n) + \frac{2n}{n+1}, \quad n \geq 2 \end{aligned}$$

$\frac{2n}{n+1} < 2$, because $\frac{n}{n+1} < 1$,

$$T(n+1) \leq \frac{n+2}{n+1}T(n) + 2, \quad n \geq 2$$

$$\text{and thus } T(n) \leq 2 + \frac{n+1}{n}T(n-1), \quad n \geq 1$$

Closed form of $T(n) = n - 1 + \frac{2}{n} \sum_{i=0}^{n-1} T(i)$

Once again we telescope:

$$\begin{aligned} T(n) &\leq 2 + \frac{n+1}{n} T(n-1), \quad n \geq 1 \\ &= 2 + \frac{n+1}{n} \left[2 + \frac{n}{n-1} \left[2 + \frac{n-1}{n-2} \left[\cdots \frac{4}{3} \right] \right] \right] \\ &= 2 \left[1 + \frac{n+1}{n} + \frac{n+1}{n} \frac{n}{n-1} + \frac{n+1}{n} \frac{n}{n-1} \frac{n-1}{n-2} + \cdots \right. \\ &\quad \left. \cdots + \frac{n+1}{n} \frac{n}{n-1} \frac{n-1}{n-2} \cdots \frac{5}{4} \frac{4}{3} \right] \\ &= 2 \left[1 + \frac{n+1}{n} + \frac{n+1}{n-1} + \frac{n+1}{n-2} + \cdots + \frac{n+1}{3} \right] \\ &= 2(n+1) \left[1 + \frac{1}{n} + \frac{1}{n-1} + \frac{1}{n-2} + \cdots + \frac{1}{3} \right] \end{aligned}$$

Closed form of $T(n) = n - 1 + \frac{2}{n} \sum_{i=0}^{n-1} T(i)$

$$\begin{aligned} T(n) &\leq 2 + \frac{n+1}{n} T(n-1), \quad n \geq 1 \\ &= \dots \\ &= 2 \left[1 + \frac{n+1}{n} + \frac{n+1}{n-1} + \frac{n+1}{n-2} + \dots + \frac{n+1}{3} \right] \\ &= 2(n+1) \left[1 + \frac{1}{n} + \frac{1}{n-1} + \frac{1}{n-2} + \dots + \frac{1}{3} \right] \\ &= 2(n+1) \left[H(n+1) - \frac{3}{2} \right] \end{aligned}$$

Where $H(n) = 1 + 1/2 + 1/3 + \dots 1/n$ is the harmonic series,
 $H(n) = \ln n \gamma + O(1/n)$, where Euler's constant $\gamma = 0.57721566 \dots$
The closed formula for $T(n)$ is


$$T(n) \leq 2(n+1)(\ln n + \gamma - 1.5) = O(n \log n)$$

Pattern matching: $p[0..m-1] \in T[0..n-1]$

- A smaller *pattern*, p , must be found in a larger *text*, T , and its *position*, $i \in T$, must be reported or on failure -1 is returned.
- When programming a matching algorithm, great care must be taken at the boundaries:
 1. If the *pattern* is empty $p = ""$ then $i = 0$ should be returned.
 2. If the *text* is empty then failure is always reported, $i = -1$.
 3. If $|p| > |T|$, i.e. $p.length() > T.length()$, then the match fails.
 4. Test that a pattern can match $T[0..m-1]$.
 5. Test that a pattern can match $T[n-m..n-1]$.
 6. Test that the pattern can match somewhere in $T[1..n-2]$, i.e. it does not fall on a boundary.

Naive, greedy or Brute Force (BF) pattern matching


- Find the *pattern*, p , in the *text*, T , and return its *position* or the return -1 on failure.
- The *brute force (BF)* method matches characters in p one-by-one with every character in T until p covers a matching part of T .



Here $p[0..j-1]$ has matched $T[i..i+j-1]$ but $p[j]$ and $T[i+j]$ do not match.

- Next: $i++$; $j = 0$; and try to match the pattern $p[0..m-1]$ with the substring $T[i..i+m-1]$ from the text, until a match is found.
- In the worst case most of p will match parts of T and mismatches only occur near the end of p . In this case at $O(m)$ matches will occur $O(n)$ times when the pattern is found near the end of T .
- Worst case for BF matching is $O(nm)$.

Naive, greedy or Brute Force (BF) pattern matching




Check if $p[j] :: T[i+j]$ match. In the figure
 $p[0..j-1]$ has matched $T[i..i+j-1]$ but $p[j] \neq T[i+j]$.
Try matching $p[0]$ and $T[i+1]$ next.

```
1 int matchBF (String p, String T) {
2   int m = p.length(); if (m == 0) return 0;
3   int n = T.length();
4   int i = 0, found = -1;
5
6   while (found < 0 && i <= n-m) {
7     int j = 0;
8     while (j < m && p[j] == T[i+j])
9       j++;
10    if (j == m)
11      found = i;
12    i++; }
13   return found; }
```

Naive, greedy or Brute Force (BF) pattern matching

—a different view: match $p[j] :: T[i]$



Check if $p[j] :: T[i]$ match. In this figure
 $p[0..j-1]$ has matched $T[i-j..i-1]$ but $p[j] \neq T[i]$.
Try $p[0] :: T[i-j+1]$ next.

```
1 int matchBF (String p, String T) {
2   int m = p.length(); if (m == 0) return 0;
3   int n = T.length();
4   int i = 0, j = 0;
5   do {
6     if (p[j] == T[i])
7       if (j == m - 1) return i - j;
8     else { i++; j++; }
9   } else {
10    i = i - j + 1;
11    j = 0; }
12  } while (i < n - m);
13  return -1; }
```

Naive, greedy or brute Force pattern matching—complexity

Worst case for BF matching is $O(nm)$.

In the *average case* mismatches occur near the middle of pattern, if not earlier, and the pattern is matched near the middle of the text if not sooner. This does not improve the timing by much more than half $O(m)$ repeated half $O(n)$ times. The average case is still $O(nm)$ but the constants are smaller.

Next we discuss three methods that dramatically improve this time complexity by devising ways of avoiding redundant comparisons.

1. *Rabin-Karp* uses hashing to slide $\mathbf{p}$ over $\mathbf{T}$.
2. *Boyer-Moore* determines if and where the character $\mathbf{T}[\mathbf{i}]$ is present in the pattern $\mathbf{p}[\mathbf{0} \dots \mathbf{m}-\mathbf{1}]$ to decide where to place $\mathbf{p}$ next.
3. *Knuth-Morris-Pratt* matches $\mathbf{p}$ with itself to determine where to move $\mathbf{p}$ next.

RK, BM and KMP matching—briefly

Three methods that improve the time complexity:

Rabin-Karp-matching compares $\text{hash}(p) :: \text{hash}(T[i..i+m-1])$. Using a “rolling” hash economically to calculate $\text{hash}(T[i+1..i+m])$ from $\text{hash}(T[i..i+m-1])$, an $O(n)$ method ensues.


Boyer-Moore-matching preprocesses p to find the last occurrence of each character. It matches from $p[m-1]$ backwards through to $p[0]$. On non-matching $p[j]$ and $T[i]$, if $T[i]$ does not occur in p , then $p[0]$ is moved to next match $T[i+1]$ —ignoring characters in $T[i-j+1..i-1]$ —otherwise p is moved putting $p[m-1]$ over $T[i']$ where $i' = i + m - \min(j, \text{lastInP}[T[i]]);$, next, match $p[m-1] :: T[i']$.

In the worst case *BM* is likely to scan at most n characters in T .

Knuth-Morris-Pratt-matching preprocesses p by first matching p with itself storing information about where to move $p[0]$ next when it gets a mismatch at $T[i]$ — $p[0..j-1] == T[i-j..i-1]$.

Rabin-Karp matching (RK)


- $|p[0..m-1]| = m$ and $|T[i..i+m-1]| = m$.
If $\text{hash}(p) == \text{hash}(T[i..i+m-1])$, then p and $T[i..i+m-1]$ are very likely to be equal.
- *Rabin-Karp* compares $\text{hash}(p)$ with $\text{hash}(T[i..i+m-1])$.
- Since Rabin-Karp uses a “rolling” hash to calculate $\text{hash}(T[i+1..i+m])$ from $\text{hash}(T[i..i+m-1])$, an $O(n)$ method ensues.
- To explain rolling hash, consider patterns of only decimal digits.



$T = 27581463145918463717346715471874$ First get $pHash = 31459$, then one-by-one match it with each $THash_i$:

$27581_0, 75814_1, 58146_2, 81463_3, 14631_4, 46314_5, 63145_6, 31459_7$. Since $pHash$ matches with $THash_7$ the position 7 is returned.

Rabin-Karp—rolling hash for decimal digits



$T = 27581463145918463717346715471874$

$$\begin{aligned} pHash &= 31459 \\ THash_0 &= 27581 \\ THash_1 &= 75814 = (THash_0 - 20000) \times 10 + 4 \\ THash_2 &= 58146 = (THash_1 - 70000) \times 10 + 6 \\ THash_3 &= 81463 = (THash_2 - 50000) \times 10 + 3 \\ THash_4 &= 14631 = (THash_3 - 80000) \times 10 + 1 \\ THash_5 &= 46314 = (THash_4 - 10000) \times 10 + 4 \\ THash_6 &= 63145 = (THash_5 - 40000) \times 10 + 5 \\ THash_7 &= 31459 = (THash_6 - 60000) \times 10 + 9 \\ THash_8 &= 14591 = (THash_7 - 30000) \times 10 + 1 \end{aligned}$$

In general

$$THash_{i+1} = (THash_i - \lfloor THash_i / 10000 \rfloor 10000) \times 10 + (int)T[i + m]$$

and in Java this could be coded as

```
1 THash = (THash - (THash/10000)*10000) * 10
2   + (int)T.charAt(i+m)-(int)'0';
```

Rabin-Karp—rolling hash for characters

Let $|T| = m$, and $h_i \equiv \text{hash}(T[i..i + m - 1])$. Given h_i calculate h_{i+1} .

Design of a hashing function for Unicode characters is based on the code for digits

1. The contribution of the leftmost character to h_i :

$$h_{\text{leftmost digit}} = \lfloor h_i / 10^{m-1} \rfloor \times 10^{m-1}$$

2. Deduct it from h_i

$$\text{hash}(T[i + 1..i + m - 1]) = h_i - h_{\text{leftmost digit}}$$

3. Shift this one character left

$$\text{hash}(T[i + 1..i + m - 1]) \times 10$$

4. Add the character $T[i + m]$.

$$\begin{aligned} h_{i+1} &= \text{hash}(T[i + 1..i + m]) \\ &= \text{hash}(T[i + 1..i + m - 1]) \times 10 + T[i + m] \end{aligned}$$

Rabin-Karp—rolling hash for characters

Let $|T| = m$, and $h_i \equiv \text{hash}(T[i..i + m - 1])$. Given h_i calculate h_{i+1} .

The hash function for Unicode characters is based on the code for digits with the base changed from 10 to b .

1. The contribution of the leftmost character to h_i :

$$h_{\text{leftmost digit}} = \lfloor h_i / b^{m-1} \rfloor \times b^{m-1}$$

2. Deduct it from h_i

$$\text{hash}(T[i + 1..i + m - 1]) = h_i - h_{\text{leftmost digit}}$$

3. Shift this one character left

$$\text{hash}(T[i + 1..i + m - 1]) \times b$$

4. Add the character $T[i + m]$.

$$\begin{aligned} h_{i+1} &= \text{hash}(T[i + 1..i + m]) \\ &= \text{hash}(T[i + 1..i + m - 1]) \times b + T[i + m] \end{aligned}$$

The biggest Unicode character has the value $2^{16} - 1$.

Using base $b = 2^{16}$ makes the above method work. *There is a snag.*

Rabin-Karp—rolling hash for characters

Basing the hash function for Unicode characters by using the base $b = 2^{16}$ has some snags:

1. The contribution of the leftmost character to h_i :

$$h_{\text{leftmost digit}} = \lfloor h_i / b^{m-1} \rfloor \times b^{m-1}$$

The value of $b^{m-1} = 2^{16(m-1)}$ uses ... bits. So when $m > \dots$ it will overflow. This is remedied by using modular arithmetic. Every arithmetic operation must be calculated modulo some large prime that is less than 2^{16} .

2. Deduct it from h_i

$$\text{hash}(T[i + 1..i + m - 1]) = h_i - h_{\text{leftmost digit}}$$

There is a new snag.

The result may become negative. It must first be made positive.

Rabin-Karp—rolling hash for characters

There is a new snag.

The result may become negative. It must first be made positive. This is done by adding **modHash** to the number as many times as needed to make it positive. Suppose that $h < 0$, then

3. Shift this one character left

$$\text{hash}(T[i + 1..i + m - 1]) \times b$$

4. Add the character $T[i + m]$.

$$\begin{aligned} h_{i+1} &= \text{hash}(T[i + 1..i + m]) \\ &= \text{hash}(T[i + 1..i + m - 1]) \times b + T[i + m] \end{aligned}$$

The biggest Unicode character has the value $2^{16} - 1$.

Using base $b = 2^{16}$ makes the above method work.

Rabin-Karp—code

```
1 public static int modHash = 65521;
2 public static final int base
3     = (Character.MAX_VALUE + 1) % modHash;
4
5 static int firstHash(String s) {
6     int m = s.length();
7     int j, h = 0;
8     for(j = 0; j < m; j++){
9         h = mod(h * base);
10        h = mod(h + s.charAt(j));
11    }
12    return h;
13 }
```


Rabin-Karp—code

```
1 static int hash(String s) {
2     return firstHash(s);
3 }
4
5 static int baseLeftmost (String T) {
6     int baseLeft = 1;
7     int i;
8
9     for (i = 1; i < T.length(); i++)
10         baseLeft = mod(baseLeft*base);
11     return baseLeft;
12 }
13 }
```

Rabin-Karp—code

```
1 static int matchRK(String p, String T){
2   int m = p.length();
3   int n = T.length();
4   if (m > n)
5     return -1;
6   int pHash = hash(p);
7   int THash = hash((String)T.subSequence(0,m));
8   int baseLeft = baseLeftmost(p);
9   int i;
```

Rabin-Karp—code

```
1 // Get the subsequence=m starting at 0
2
3 for (i = 0; i + m <= n; i++) {
4     if (THash == pHash) {
5         if (stringCompare(
6             (String)T.subSequence(i, i+m), p))
7             return i;
8     }
9     if (i == n-m)
10        return -1;
11    // deduct contribution of left character
12    THash = mod(THash
13        - mod(baseLeft * T.charAt(i)));
14    // Shift left
15    // add character at right
16    THash = mod(mod(THash*base)
17        + T.charAt(i+m)) ;
18 }
19 return -1;
20 }
```

Rabin-Karp—code

```
1 static boolean stringCompare(String p,  
2                               String T) {  
3     int m = p.length();  
4     int n = T.length();  
5  
6     int i;  
7  
8     if (m != n)  
9         return false;  
10    for (i = 0; i < m; i++)  
11        if (p.charAt(i) != T.charAt(i))  
12            return false;  
13    return true;  
14 }
```

Rabin-Karp—code

```
1 static int mod (int top) {  
2 int modulo = top % modHash;  
3  
4     if (modulo < 0)  
5         return modulo  
6             + modHash * (1 + (-modulo/modHash));  
7     else  
8         return modulo;  
9  
10 }
```

Rabin-Karp matching—rolling hash

- Normally the patterns may get too big to be stored as a unique number. Apply modulo arithmetic when hashing. Use a large prime number `modHash` to `mod` every arithmetic result.
- Calculate the hash roughly as follows:

```
1 int firstHash(String s) {
2   int m = s.length();
3   int j, h=0;
4   for(j = 0, j < m, j++)
5     h = (h+mod(s.charAt(j)*base));
6   return h;
7 }
8
9 int nextHash(char left, char right,
10             int THash) {
11   THash = mod(THash - mod(left*base^(m-1)));
12   return mod(THash*base + right);
13 }
```


Rabin-Karp matching—pseudo code

```
1 int rabinKarpMatch(String p, String T){
2   int m = p.length();
3   int n = T.length();
4   int pHash = hash(p[0..m-1]);
5   int THash = hash(T[0..m-1]);
6
7   for (i = 0; i < n; i++) {
8     if (THash == pHash)
9       // verify equality of patterns
10      if (T[i..i+m-1] == p)
11        return i;
12     THash = hash(T[i+1..i+m]);
13   }
14   return -1;
15 }
```

Rabin-Karp matching (RK)—complexity

- Since every character in the *text* file must be used in the hashing formula, at worst all characters need to be scanned, taking $O(n)$.
- The pattern must be hashed, $O(m)$.
- The worst case time is thus $O(m + n)$.
- This is at least m times faster than the BF method.


Boyer-Moore matching (BM)



If the letter a never appears in p , the pattern may be moved on, placing $p[0]$ over $T[i+1]$ and $p[m-1]$ over $T[i+m]$ —*ignoring* the characters in $T[i-j+1..i-1]$.

Boyer-Moore-matching preprocesses p to find the last occurrence of each character. It matches from $p[m-1]$ backwards through to $p[0]$. On non-matching $p[j]$ and $T[i]$, if $T[i]$ does not occur in p , then $p[0]$ is moved to next match $T[i+1]$ —ignoring characters in $T[i-j+1..i-1]$ —*otherwise* p is moved putting $p[m-1]$ over $T[i']$ where $i' = i + m - \min(j, 1 + \text{lastInP}(T[i]))$; , next, match $p[m-1]::T[i']$. In the worst case *BM* can degenerate to $O(mn + |\Sigma|)$ but it is likely to scan less than n characters in T .

Boyer-Moore matching (BM)



- For p , calculate the table `lastInP` and place $p[0]$ over $T[0]$, i.e. set $i = m-1$ and $j = m-1$.
- Compare $p[j]::T[i]$, i.e. compare $p[m-1]::T[m-1]$. If equal, then $j--$ and $i--$, until a match of p at $j == 0$.
- If $p[i]::T[i]$ are not equal: Look up the entry in `lastInP(T[i])`. It is -1 if $T[i] \notin p$.
- If $T[i]$ does not appear in the pattern p then move i forward: $i' = i + m$.
- If the last appearance of $T[i] \in p$ is at k then move i forward: $i' = i + m - \min(j, 1+k)$.
- Repeat while $(i < n)$.

Boyer-Moore—getLast(p, lastInP)

```
1 static void getLast(String p, int lastInP[]) {
2 // pre: p[0..m-1] pattern for which
3 // lastInP must be created
4 // post: Boyer-Moore's lastInP_vector
5 //
6
7 int _j;
8 int _s=_256;
9 int _m=_p.length();
10
11 for_(j=0;_j<s;_j++)
12     lastInP[j]=_1;
13 for_(j=0;_j<m;_j++)
14     lastInP[(int)(p.charAt(j))]=_j;
15 }
```

Boyer-Moore—matchBM(p, T)

Start with the usual tests on the bounds **m** and **n** and set up **lastInP**.

```
1 static int matchBM(String p, String T) {
2 // pre: p[0..m-1] pattern to match
3 // T[0..n-1] text in which to find pattern
4 // post: returns index of first matched position
5
6 int m = p.length();
7 int n = T.length();
8
9     if (m == 0) return 0;
10    else if (n == 0) return -1;
11    else if (m > n) return -1;
12
13 int s = 256;
14 int last[] = new int [s];
15 int i; int j;
16
17    getLast(p, last);
```

Boyer-Moore—matchBM(p, T)

The code puts i , so that $p[0]$ lies at $T[i-m+1]$.

```
1  i = j = m - 1;
2  do {
3      if (p.charAt(j)==T.charAt(i))
4          if (j == 0 )
5              return i; // first match
6          else {
7              i--;
8              j--;
9          }
10     else {
11         i = i + m // jump step
12             - Math.min(j, 1+last[T.charAt(i)]);
13         j = m - 1;
14     }
15 } while (i < n );
```

Boyer-Moore matching—complexity


- BM must set up `lastInP[]`.
This costs $|\text{lastInP}[]| = |\Sigma|$.
Then it starts comparing from the back of `p`.
Unless `p` is covering itself in `T`,
or if the $|p| \ll |T|$ or $|p| \ll |\Sigma|$,
most comparisons are going to be unequal.

So, on *average*, the complexity is $O(n + \Sigma) = O(n)$ since Σ is a constant.


Boyer-Moore matching—complexity

- Many jumps might be m long, so the complexity could become $O(n/m)$.
e.g. $\text{p} = \text{"yyyyyyyyy"}'$, i.e. m ‘y’s, and no ‘x’s
 $\text{T} = \text{"xxxxxxxxxxxxxxxxxxxxxyyyyyyyy"}'$, $|\text{T}| = n$.
- The *worst case* for BM appears when
 $\text{p} = \text{"yxxxxxxxxxxx"}'$, i.e. a ‘y’ and $m - 1$ ‘x’s and $\text{T} = \text{"xxxxxxxxxxxxxxxxxyxxxxxxxxxxx"}'$, $|\text{T}| = n$.
BM degenerates to $O(mn + |\Sigma|)$
- Generally, BM is likely to scan less than n characters in T .

Knuth-Morris-Pratt matching—preview



Knuth-Morris-Pratt-matching (KMP) preprocesses p by first matching p with itself storing information about where to move $p[0]$ next when it gets a mismatch at $p'[i]$: We know $p[0..j-1] == p'[i-j..j-1]$.



In the matching piece KMP finds a *prefix* of $p \equiv p[0..k]$ that matches a *suffix* of $p' \equiv p'[i-k-1..i-1]$.

- When there *is a match*, KMP always sets the next value of j to $j++$ and steps i to $i++$.
- When there *is no match*, j becomes $\text{next}[j] \equiv k$ and if $j \geq 0$ then i stays the same and KMP avoids redundant comparisons in $p'[i-j..i-1]$, otherwise when $j < 0$ then $j = 0$ and $i++$.

KMP matching—examples

Let i run in T and j run through p , comparing $p[j]$ with $T[i]$ and putting $i++$; $j++$; when they are equal.


T $\begin{array}{c} 0 \qquad i \qquad \qquad \qquad n-1 \\ \boxed{\text{band} \text{bandana one bandanatwobandanas}} \end{array}$
 p $\begin{array}{c} 0 \qquad j \\ \boxed{\text{band} \text{anas}} \\ m-1 \end{array}$

$p[0..3] \equiv T[0..3]$ but $p[j=4] \equiv 'a' \neq T[i=4] \equiv 'b'$, Since $p[0..3] \equiv \text{'band'}$ does not match with any substring of itself, next put $j=0$ — i stays the same—comparing $p[0..] :: T[4..]$. KPM ‘knows’ $p[0] = \text{'band'}$ and safely jumps over it.


T $\begin{array}{c} 0 \qquad 4 \qquad \qquad 10 \ i \ 12 \qquad \qquad \qquad n-1 \\ \boxed{\text{bandbandana} \text{one bandanatwobandanas}} \end{array}$
 p $\begin{array}{c} 0 \qquad j \\ \boxed{\text{bandana} \text{s}} \\ m-1 \end{array}$

Now $p[0..6] \equiv \text{'bandana'}$ matches $T[4..10]$ and j becomes 0 putting $p[0..]$ over $T[11..]$.


Knuth-Morris-Pratt matching (KMP)



Subsequently j hobbles forward one-by-one until `bandana` is matched at $T[14..20]$. The mismatch at $p[7]='s'$ with $T[21]='t'$ causes p to jump forward to try matching $p[0..]$ over $T[21..]$



Knuth-Morris-Pratt matching (KMP)



When $p[0]$ mismatches both i and j are incremented:

$p[j=0] \equiv 'b'$ mismatches $T[i=21] \equiv 't'$ and then

$p[j=0] \equiv 'b'$ mismatches $T[i=22] \equiv 'w'$ and then

$p[j=0] \equiv 'b'$ mismatches $T[i=23] \equiv 'o'$ and then

$p[j=0] \equiv 'b'$ matches $T[i=24] \equiv 'b'$, and eventually the pattern

$p[0..7] \equiv \text{'bandanas'}$ matches $T[24..31]$.

The value of $\text{next}[0] = -1$ and all the other values of $\text{next}[j]$ here are 0 because 'bandanas' has no suffixes that match prefixes of itself.

The following example has substrings $p[0..k]$, $k > 0$, which have prefixes that match their own suffixes.

Consider the pattern **ababbababaa**.

Knuth-Morris-Pratt matching (KMP)

The pattern **ababbababaa** contains substrings $p[0..k]$, $k > 0$, with prefixes that match their own suffixes.

j	next[j]
0	a -1 -
1	b 0 a-
2	a 0 ab-
3	b 1 aba-: $p[0..0] \equiv a \equiv p[2..2]$
4	b 2 abab-: $p[0..1] \equiv ab \equiv p[2..3]$
5	a 0 ababb-
6	b 1 ababba-: $p[0..0] \equiv a \equiv p[5..5]$
7	a 2 ababbab-: $p[0..1] \equiv ab \equiv p[5..6]$
8	b 3 ababbaba-: $p[0..2] \equiv aba \equiv p[5..7]$
9	a 4 ababbabab-: $p[0..3] \equiv abab \equiv p[5..8]$
10	a 3 ababbababa-: $p[0..2] \equiv aba \equiv p[7..9]$

If the pattern **p=ababbababaa** mismatches the text

T=ababbababababababababababababab at p[j=10]≡‘a’ !=
T[i=10]≡‘b’, the next[j] value is 2 and j must be set to j =
next[j]+1≡3, leaving i the same.

Knuth-Morris-Pratt matching (KMP)

The pattern **abababababa** contains substrings $p[0..k]$, $k > 0$, with prefixes that match their own suffixes.

0	a	-2	-	
1	b	-1	a-	
2	a	-1	ab-	
3	b	0	aba-	$p[0..0] \equiv a \equiv p[2..2]$
4	a	1	abab-	$p[0..1] \equiv ab \equiv p[2..3]$
5	b	2	ababa-	$p[0..2] \equiv aba \equiv p[2..4]$
6	a	3	ababab-	$p[0..3] \equiv abab \equiv p[2..5]$
7	b	4	abababa-	$p[0..4] \equiv ababa \equiv p[2..6]$
8	a	5	abababab-	$p[0..5] \equiv ababab \equiv p[2..7]$
9	b	6	ababababa-	$p[0..6] \equiv abababa \equiv p[2..8]$
10	a	7	ababababab-	$p[0..7] \equiv abababab \equiv p[2..9]$

If the pattern $p = \text{abababababa}$ mismatches the text

$T = \text{ababababab} \underline{\text{b}} \text{abababababababababababababab}$ at $p[j=10] \equiv 'a' \neq T[i=10] \equiv 'b'$, the $\text{next}[j]$ value is 7 and j must be set to $j = \text{next}[j] + 1 \equiv 8$, leaving i the same.

Knuth-Morris-Pratt matching (KMP)

- In KMP the text T is processed forward.
- Since all of $p[0 \dots j-1]$ has been matched, none of the corresponding characters in $T[i-j \dots i-1]$ need ever be matched again.
- Sometimes the character at $T[i]$ which mismatched with character $p[j]$ may be repeatedly matched with different characters of the pattern p .
- The next placement of p over T depends entirely on what has already been matched in p .
- This information is applied to decide which character of p must next be matched with $T[i]$.

Knuth-Morris-Pratt matching—recap

- This information is applied to decide which character of p must next be matched with $T[i]$.
- k is calculated to place $p[k+1]$ over $T[i]$ where $p[0..k]$ is the *greatest* prefix in $p[0..j-1]$ that matches a suffix $p[j-k-1..j-1]$. Note that length $|p[0..k]| = k + 1 = |p[j-k-1..j-1]|$.
- We call the table of k 's `next[]`.

Knuth-Morris-Pratt matching—recap

- The longest prefix-suffix pair is placed over one another first.
- If there is a match: $i++;$ $j++;$, else
- If there is a mismatch a next j that is always smaller than the previous one is calculated using $j = \text{next}[j] + 1$.
- Note that i is only advanced when there is a match or when j becomes 0 .
- We next construct the `next[]` table.
- When there is a mismatch of $p[j]$ with $T[i]$ the value of `next[j]` tells us to advance j to `next[j]+1`.
- Usually i stays the same but can advance by 1 .

KMP—constructing `next[]`

- When `p[j]` and `T[i]` match, simply increment both `i` and `j`, using `i++; j++;`.
- We will now consider mismatches at various values of `j ∈ [0..m-1]`.
- At a mismatch of `p[0]`, simply put `i++; j++;`.
- We will construct `next[]` for Manber's example from Udi Manber [Section 6.7, pp. 148–155, *Introduction to Algorithms A Creative Approach*, Addison-Wesley Publishing Co., Reading, MA., 1989]

Note that the **ababbababaa** example is this example in slightly disguised form.

KMP—constructing next[]

Consider a mismatch at $p[0]$. Since no information has been gained, simply $i++;$ $j++;$.

T $\boxed{x \ y \ x \ y \ y \ x \ y \ x \ y \ x \ x}$
p $\boxed{- \ y \ x \ y \ y \ x \ y \ x \ y \ x \ x}$

T $\boxed{x \ y \ x \ y \ y \ x \ y \ x \ y \ x \ x}$
p $\boxed{x \ - \ x \ y \ y \ x \ y \ x \ y \ x \ x}$

On mismatching $p[1]$ —above right— $p[0..0]$ is known, but it does not help so $j=0$ and i remains the same.

After a mismatch at $p[2]$ —below left—the substring $p[0..1]$ is known, but no prefix equals any suffix so again $j=0$ and i remains the same.

T $\boxed{x \ y \ x \ y \ y \ x \ y \ x \ y \ x \ x}$
p $\boxed{x \ y \ - \ y \ y \ x \ y \ x \ y \ x \ x}$

T $\boxed{x \ y \ x \ y \ y \ x \ y \ x \ y \ x \ x}$
p $\boxed{x \ y \ x \ - \ y \ x \ y \ x \ y \ x \ x}$

The mismatch at $p[3]$ means that we have covered $p[0..2] \equiv "xyx"$. This has a prefix x - that equals a suffix $-x$ so now $j=1$ and i remains the same.

j	0	1	2	3	4	5	6	7	8	9	10
p	x	y	x	y	y	x	y	x	y	x	x
next[j]	-1	0	0	1	.	.	.	.	.	.	.

KMP—constructing next[]

When mismatching at $p[4]$, $p[0..3] \equiv \text{"xyxy"}$ It has a prefix xy- that equals a suffix -xy , so now $j=2$ and i remains the same.

T

x	y	x	y	y	x	y	x	x
---	---	---	---	---	---	---	---	---

T

x	y	x	y	y	x	y	x	x
---	---	---	---	---	---	---	---	---

On mismatching $p[5]$ —above right— $p[0..4] \equiv \text{"xyxyy"}$, has no prefix-suffix pair, so $j=0$ and i static.

After a mismatch at $p[6]$ —below left—the substring $p[0..5] \equiv \text{"xyxyyx"}$ has the prefix-suffix pair x- , -x and $j=1$ and i is static.

T

x	y	x	y	y	x	y	x	x
---	---	---	---	---	---	---	---	---

T

x	y	x	y	y	x	y	x	x
---	---	---	---	---	---	---	---	---

The mismatch at $p[7]$ means that we have covered $p[0..6] \equiv \text{"xyxyyyxy"}$. This has a prefix xy- that equals a suffix -xy so now $j=2$ and i remains the same.

j	0	1	2	3	4	5	6	7	8	9	10
p	x	y	x	y	y	x	y	x	y	x	x
next[j]	-1	0	0	1	2	0	1	2	.	.	.

KMP—constructing next[]

When mismatching at $p[8]$, $p[0..7] \equiv \text{"xyxyyxxyx"}$ It has a prefix xyx- that equals a suffix -xyx , so now $j=3$ and i remains the same.

T x y x y y x y x y x x
p x y x y y x y x x x

T x y x y y x y x y x x
p x y x y y x y x y x

If $p[9] \neq T[9]$ —above right— $p[0..8] \equiv \text{"xyxyyxxyxy"}$, This has a prefix xyxy- that equals a suffix -xyxy , so now $j=4$ and i remains the same.

After the mismatch at $p[10]$ —below—the substring

$p[0..9] \equiv \text{"xyxyyxxyxxy"}$ has the prefix-suffix pair xyx- , -xyx , so $j=3$ and i is static.

T x y x y y x y x y x x
p x y x y y x y x y x

The table has now been completed.

j	0	1	2	3	4	5	6	7	8	9	10
p	x	y	x	y	y	x	y	x	y	x	x
next[j]	-1	0	0	1	2	0	1	2	3	4	3

KMP—buildNext

Pseudocode for buildNext.

```
1 void buildNext(String p, int [] next) {
2   int m = p.length();
3   int i, j;
4
5   next[0] = -2;
6   next[1] = -1;
7   for (i = 2; i < m; i++) {
8     j = next[i-1]+1;
9     while (j >=0 && p[i-1] != p[j]){
10       j = next[j]+1; }
11     next[i] = j; }
12   for (j = 0; j < m; j++)
13     next[j]++; }
```

Pseudocode for buildNext

```
1 int matchKMP(String p, String T) {  
2     int m = p.length();  
3     int n = T.length();  
4  
5     if (m == 0) return 0;  
6     else if (n == 0) return -1;  
7     else if (m > n) return -1;  
8  
9     int i, j, next[] = new int [m+1];  
10  
11     buildNext(p, next);  
12     ...
```

KMP—KMPmatch

```
1  ...
2  buildNext(p, next);
3  i = 0;
4  j = 0;
5  while (i < n) {
6      if (p[j] == T[i]){
7          i++;
8          j++; }
9      else {
10         j = next[j];
11         if (j < 0) {
12             j = 0;
13             i++; } }
14     if (j == m)
15         return i - m; }
16 return -1; }
```

- Calculating `buildNext[]` takes $O(m)$. Matching the pattern using the `next[]` table takes $O(n)$. KMP runs in $O(m + n)$.
- One character in `T[i]` may be matched with many characters from the pattern `p`. How many?
- When there is a mismatch with `T[i]` another character from `p` is matched if `next[j] >= 0`.
- Suppose the first mismatch involved `p[j]`, since each evaluation of `next[j]` leads us to a smaller index $j \in [0..j-1]$, so we can only j times. However, to reach `p[j]` we must have gone forward j times without any backtracking.
- If we assign the costs of backtracking to the forward moves then we at most double the number of forward moves.
- There are only n forward moves so the number of comparisons is $O(n)$.


Preprocess the *text* T —not the *pattern* p

- In Rabin-Karp, Boyer-Moore, and Knuth-Morris-Pratt pattern matching the pattern p is *preprocessed* to speed up searching for its occurrence in a text T .
- If many searches $p \in T$ are done, such as searching for the presence of a word p in a dictionary T , then *preprocessing* the text T may help to speed up queries, especially if T is large or when many different queries p are made.
- A *trie* is a very fast, dense structure that allows searches in time proportion to the *pattern* size, i.e. $O(m)$.

Standard tries

The standard trie for a set of strings T is an ordered tree such that:


- Each node but the root is labelled with a character.
- The children of a node are alphabetically ordered.
- The paths from the leaves to the root yield the strings of T .
- e.g. standard trie for the set of strings
 $T = \{bear, bell, bid, bull, buy, sell, stock, stop\}$



Analysis of standard tries

A standard trie uses $O(n)$ space and supports searches, insertions and deletions in time $O(md)$, where:

- n total size of the strings in T , and m size of the string parameter of the operation, while $d = |\Sigma|$ size of the alphabet.
- Using binary search at each node for the next letter, searches, insertions and deletions may be done in $O(m \log d)$ time.
- By using hashing or direct lookup this is further reduced to $O(m)$.



Word matching with a trie


s	e	e		a		b	e	a	r	?		s	e	l	l		s	t	o	c	k	!	
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23

s	e	e		a		b	u	l	l	?		b	u	y		s	t	o	c	k	!	
24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46

b	i	d		s	t	o	c	k	!		b	i	d		s	t	o	c	k	!	
47	48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63	64	65	66	67	68


h	e	a	r		t	h	e		b	e	l	l	?		s	t	o	p	!	
69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	

- We insert some of the words of the text into a trie.
- Each leaf stores the occurrences of the associated word in the text.




Compressed tries

- A compressed trie has internal nodes of degree at least two.
- It is obtained from a standard trie by compressing chains of *redundant* nodes.




Compact Representation

- Compact representation of a compressed trie for an array of strings:
- Stores at the nodes ranges of indices instead of substrings
- Uses $O(s)$ space, where s is the number of strings in the array
- Serves as an auxiliary index structure




Suffix trie

- The suffix trie of a string T is the compressed trie of all the suffixes of T .



Analysis of suffix tries

- Compact representation of the suffix trie for a string T of size n from an alphabet of size d .
- Uses $O(n)$ space.
- Supports arbitrary pattern matching queries in T in $O(dm)$ time, where m is the size of the pattern.
- Can be constructed in $O(n)$ time.



Text compression with Huffman's code

- Given a text T reproduce it in a lossless way using less memory.

T is a piece of text in which we count the characters to find their frequencies.

- Huffman scheme counts the number of times characters appear in T .
- Pseudo code for getting the frequencies of the characters T in follows:

Text compression with Huffman's code

Ch		a	b	c	d	e	f	h	i	k	n	o	r	s	t	u	v
F	9	5	1	3	7	3	1	1	1	1	4	1	5	1	2	1	1

- By combining pairs of these frequencies, with the lowest first, and then adding their combinations to the frequency list, and removing the originals a binary tree is produced from which an optimal code can be written down.
- The compressed code may then rapidly be produced from the binary tree.

Huffman code—adapted from GT page 566.


Note that the text I have used differs slightly from that of GT page 566.

Char.		a	b	c	d	e	f	h	i	k	n	o	r	s	t	u	v
Freq.	9	5	1	3	7	3	1	1	1	1	4	1	5	1	2	1	1


- In order to find the lowest frequency characters first, an obvious method is to: sort the characters by frequency; combine the top pair; remove the top two characters but insert their combination back into the list at its correct position.
- The very first pair with the lowest frequency could appear at the lowest node in the final tree.
- Subsequent combinations fill the tree from the lowest level.
- The last combination is hooked into the root of the tree.

Huffman coding—combining the lower nodes


First, combine nodes that are pairs of 1s into 2-nodes.



Then combine the last 1-node with a 2 node, and start combining pairs of 2 nodes.




Reorder the list and combine the remaining 2 nodes. Combine the pair of 3-leaves.




Huffman coding—reorder and combine

Simply reorder the list.




Combine lowest nodes to the next level.



Note that there are only five nodes left to combine.

Huffman coding—combine the 1s


After combining 7 + 8 and 9 + 10, we get:



Next we will rearrange the nodes—high frequencies to the left.

Huffman coding—rearrange the forest of trees from highest to lowest

The nodes rearranged—high counts to the left, lowering to the right.




There are only three nodes left.

We will combine the lowest two, resort and again combine the lowest two to create a single rooted tree from the forest.


Huffman coding—combine and sort

Now combine the 13-tree and the 15-tree into a 28-tree, and resort the remaining forest of two trees.




Huffman code—finally combine and label

Combine the last two trees and label each left branch of the tree with a 0 and each right branch with a 1.




Huffman code—finally combine and label

We rearrange the tree such that the levels of each character can be seen more easily.



From the figure it follows that the number of digits needed to represent each character is the same as its depth in the Huffman tree.

Huffman code in tabular form



Char.		a	b	c	d	e	f	h	i
Freq.	9	5	1	3	7	3	1	1	1
Code	11	100	00110	0110	010	0111	001110	001111	001000

Char.	k	n	o	r	s	t	u	v
Freq.	1	4	1	5	1	2	1	1
Code	001001	0000	001010	101	001011	00010	000110	000111

"a fast runner need never be afraid of the dark" coded is:

1001100111 0100001011 0001011101 0001100000 0000011110
 1110000011 1011101011 0000011100 0111011110 1110011001
 1111100001 1101011000 0100001011 0010100011 1011000100

0111101111 1010100101 001001

Huffman code for T

Char.	□	a	b	c	d	e	f	h	i
Freq.	9	5	1	3	7	3	1	1	1
Code	11	100	00110	0110	010	0111	001110	001111	001000

Char.	k	n	o	r	s	t	u	v
Freq.	1	4	1	5	1	2	1	1
Code	001001	0000	001010	101	001011	00010	000110	000111

"a fast runner need never be afraid of the dark" coded is:

```
1001100111 0100001011 0001011101 0001100000 0000011110
1110000011 1011101011 0000011100 0111011110 1110011001
1111100001 1101011000 0100001011 0010100011 1011000100
```

```
0111101111 1010100101 001001
```

The Huffman code has 176 bits. Since $|\mathbf{T}| = 46$ and it has an alphabet of 16 characters a fixed length code of 4 bits per character is needed to encode this string, i.e. $46 \times 4 = 184$.

This is *not* the optimal Huffman code for this string. Why not?

Huffman code for T

Char.	␣	a	b	c	d	e	f	h	i
Freq.	9	5	1	3	7	3	1	1	1
Code	11	100	00110	0110	010	0111	001110	001111	001000

Char.	k	n	o	r	s	t	u	v
Freq.	1	4	1	5	1	2	1	1
Code	001001	0000	001010	101	001011	00010	000110	000111

"a fast runner need never be afraid of the dark" coded is:

```
1001100111 0100001011 0001011101 0001100000 0000011110
1110000011 1011101011 0000011100 0111011110 1110011001
1111100001 1101011000 0100001011 0010100011 1011000100
```

```
0111101111 1010100101 001001
```

The Huffman code has 176 bits. Since $|T| = 46$ and it has an alphabet of less than 16 characters a fixed length code of 4 bits per character is needed to encode this string, i.e. $46 \times 4 = 184$.

This is *not* the optimal Huffman code for this string. Why not?

Since our code caters for some of characters that have different counts from T, e.g. **b**, **c**:3, **d**:7, **e**:3, and **f**:1, than this string has, the code cannot be optimal. H.W. Calculate the optimal Huffman code for T.

Huffman with sorting and insertion into the sorted list

- Sorting the m characters of the text according to their frequencies may be done in $O(m \log m)$ time.
- The position of each insertion is found in $O(\log m)$ time. How? But there are this does not really help because insertions entail moving the rest of the list down—on average $m/2$ such items are shifted. So this part takes at worst $O(m)$ time for each of m entries, i.e. it takes $O(m^2)$ time.
- So the Huffman tree is built in $O(m \log m) + m^2$ which yields $O(m^2)$ time using the sort-and-insert method.
- This time can be improved.

Huffman code using a heap

- Instead of repeated sorting and insertion into the sorted list, a heap should rather be used.
- The $\langle \text{character(s)}, \text{frequency} \rangle$ entries are placed into a heap with the entries with the lowest frequency appearing on the top of the heap.
- The top two entries are removed from the heap and combined into a new entry.
- The new entry is then inserted into the heap at its appropriate position.
- Building the heap *bottom-up* takes $O(m)$ time. Doing m insertions into a heap takes $O(m \log m)$ time.
- The total cost of constructing the Huffman tree is thus $m \log m$.

Huffman code explained


Consider the set $\{a, b, c, d, e, f, r\}$ with the frequencies and probabilities for each letter:

Character	a	b	c	d	e	f	r	All
Frequency	2	4	5	7	15	29	60	122
Probability	0.016	0.033	0.041	0.057	0.123	0.238	0.492	1.000


The total of the frequencies is 122.

The letters are sorted according to counts or probability—smallest first. The probability is frequency divided by the total frequency.

The two nodes with the lowest counts are combined:



Insert the combined node $[a+b : 6]$ after node $[c : 5]$ and before node $[d : 7]$.




Huffman coding


Starting with



form node $[c+(a+b) : 11]$ by combining combine nodes $[c : 5]$ and node $[a+b : 6]$:




Put node $[c+(a+b) : 11]$ in its correct place:




Huffman coding



Form node $[d+(c+(a+b)) : 18]$ by combining combine nodes $[d : 7]$ and node $[c+(a+b) : 11]$:





Put node $[d+(c+(a+b)) : 18]$ in its correct place:




Huffman coding—repeat steps



Huffman coding—repeat steps




Huffman coding—last combination



Huffman coding—enter coding

Tag each left leg of the tree with a **0** and each right leg with a **1**.




Read off the codes for each letter in the tree:

Character	a	b	c	d	e	f	r
Frequency	2	4	5	7	15	29	60
Probability	0.016	0.033	0.041	0.057	0.123	0.238	0.492
Code	111110	111111	11110	1110	110	10	0

What does **1111110111110111101101110** represent?

Huffman coding—enter coding

A Huffman tree from which the levels of each character may be seen.



The code in tabular form.

Character	a	b	c	d	e	f	r
Frequency	2	4	5	7	15	29	60
Probability	0.016	0.033	0.041	0.057	0.123	0.238	0.492
Code	111110	111111	11110	1110	110	10	0

Divide-and-conquer methods


- In general, divide-and-conquer methods *divide* a problem with a set of N inputs X_i where $i \in [0..N)$ into k equal or almost equal parts.
- Each of these parts is then in turn divide-and-conquer-ed.
- Finally the results are combined.

```
1 def divide_and_conquer(X, low, high):
2     if high - low < small:
3         quickly resolve problem for small  $N$  in  $O(1)$  time
4         return result
5     else:
6         divide  $X$  into smaller parts  $P_1, P_2, \dots, P_k, k \geq 1$ 
7         return combine(divide_and_conquer( $P_1$ ),
8                           divide_and_conquer( $P_2$ ),
9                           ...,
10                          divide_and_conquer( $P_k$ ))
```

- The resulting timing depends largely on the time taken by **combine**.
- If $T(\text{combine}) = O(N)$,
$$T(N) = kT(N/k) + O(N) = O(N \log_k N)$$

An $O(N^2)$ method for 1-D closest pair

Given a set of N points in a straight line—use the X -axis for simplicity—find the closest pair of points.




- The problem can be solved brute force by running through all the pairs in $O(N^2)$ time as follows

```
1  shortestDistance = |X[1] - X[2]|
2  indexJ = 1
3  indexK = 2
4  for j in [1..N-1]:
5      for k in [j+1..N]:
6          difference = |X[j] - X[k]|
7          if difference < shortestDistance:
8              shortestDistance = difference
9              indexJ = j
10             indexK = k
```

An $O(N \log N)$ divide-and-conquer method for 1-D closest pair

Given a set of N points on the X -axis for find the closest pair of points.



The problem can be solved in $O(N \log N)$ time using divide-and-conquer.

```
1 def closestPairs(X, low, high):
2     if high - low == 2: return |X[high] - X[low]|
3     elif high == low: return 1.79768e+308
4     else:
5         m = median(X, low, high)
6         L = closestPairs(X, low, m)
7         R = closestPairs(X, m+1, high)
8          $\delta = \min(L, R)$ 
9         straddle = minimum (X, m- $\delta$ , m+ $\delta$ )
10    return min(L, R, straddle)
```

It takes time $T(N) = 2T(N/2) + O(N)$, i.e.,
 $T(N) = O(N \log N)$.

The longest maximum subvector problem

In a list or vector X of N integers find the maximum sum of any contiguous subvector.¹³

-66	-15	87	91	32	5	-47	47	4	-77
0	1	2	3	4	5	6	7	8	9

In the list X above the subvector $X[2..8]$ has the maximum sum 219.


Clearly if all the numbers are negative then the maximum contiguous subvector (MCS) is simply the largest number. In this case we define the MCS as 0.

We know how to solve this problem in $O(N)$ time but we are now interested in using divide-and-conquer to solve it in the slower time of $O(N \log N)$

¹³This example comes from pages 77–81, “Programming Pearls”, Jon Bentley, ACM Press and Addison-Wesley, New York, 2000.

A divide-and-conquer method for MCS


First, we divide the given array $X[]$ into two parts, **Left** and **Right**:



and then find the MCS **mLeft** in **Left** and the MCS **mRight** in **Right**.



There is still a snag—the MCS might straddle both **Left** and **Right**.



The MCS is the subvector that gives the biggest of these three values.

The final algorithm is (1) $\text{middle} = N/2$,

(2) Get MCS, **mLeft** of $X[\text{low}..\text{middle}]$, using

$\text{mLeft} = \text{MCS}(X, \text{low}, \text{middle})$, and get the MCS, **mRight** of $X[\text{middle}+1..N]$, using $\text{mRight} = \text{MCS}(X, \text{middle}+1, N)$.

More difficult is getting the straddling MCS, **mStraddle**.

Calculating the straddling MCS

The idea of the straddling MCS is first to calculate the MCS, called ‘maxsofarLeft’ that ends at $X[\text{middle}]$ but starts somewhere in **Left**—running backwards, and also calculating the MCS called ‘maxsofarRight’ that starts at $X[\text{middle}+1]$ and ends somewhere in **Right**. The sum of these two values is **mStraddle**.

```
1 int maxStraddle(low, high)
2     middle = (low + high) / 2
3     sum = 0; maxsofarLeft = 0
4     for i in [low..middle]
5         sum += X[middle - i + 1]
6         if sum > maxsofarLeft
7             maxsofarLeft = sum
8     sum = 0; maxsofarRight = 0
9     for i in [middle+1..high]
10        sum += X[i]
11        if sum > maxsofarRight
12            maxsofarRight = sum
13    return maxsofarLeft + maxsofarRight
```

A divide-and-conquer method for MCS

```
1 int MCS(low, high)
2     if low > high
3         return 0
4     if low == high
5         return max(0, X[low])
6     middle = (low + high) / 2
7     mLeft = MCS(X, low, middle)
8     mRight = MCS(X, middle+1, high)
9     mStraddle = maxStraddle (X, low, high)
10    return max (mLeft, mRight, mStraddle)
```

Suppose ‘MCS’ takes time $T(N)$ for an input of N . Then ‘mLeft’ takes time $T(N/2)$ and ‘mRight’ takes time $T(N/2)$. ‘maxStraddle()’ takes $O(N)$ because, at worst, it scans the entire input once. We conclude that

$$\begin{aligned} T(N) &= 2T(N/2) + O(N) \\ &= O(N \log N) \end{aligned}$$

Divide-and-conquer methods— $O(N \log N)$

Suppose the divide-and-conquer method has time complexity $T(N)$ when it has an input of size N .

Divide-and-conquer methods entail:

- (1) *dividing* the problem into equal parts,
 - (2) *solving each part* of the problem separately, and
 - (3) *joining* the results somehow.
1. *Dividing* the problem in two, has the low cost of $O(1)$. All that needs to be done is $N = N/2$.
 2. *Solving each half* of the problem separately costs $2T(N/2)$.
 3. *Joining the two results* is the critical part. Try to make this depend *linearly* on N . If it is sublinear all the better. Usually this part entails looking at every input once. In mergesort the two sorted halves are merged in time cN , making the cost of merge sort, $T(N) = 2T(N/2) + cN$.
 4. In MCS straddle passes through all N elements once, giving the same formula.

Divide-and-conquer methods— $O(N \log N)$ —I

Solving the recursive equation for merge sort and MCS

$$T(N) = 2T(N/2) + cN.$$

is quite easy.

$$\begin{aligned} T(N) &= 2T(N/2) + cN \\ &= 2[2T(N/4) + cN/2] + cN \\ &= 2^2T(N/2^2) + cN + cN \\ &= 2^2[2T(N/2^3) + cN/2^2] + cN + cN \\ &= 2^3T(N/2^3) + cN + cN + cN \\ &= 2^3T(N/2^3) + 3cN \\ &\vdots \\ &= 2^kT(N/2^k) + kcN \end{aligned}$$

Continue until $N/2^k = 1$.

$$\begin{aligned} T(N) &= 2T(N/2) + cN \\ &\vdots \\ &= 2^k T(N/2^k) + kcN \end{aligned}$$

Continue until $N/2^k = 1$.

Since there is nothing to do when the input is of size 1, it follows that $T(1) = 0$, so we set $N/2^k = 1$ giving $N = 2^k$. This happens when $k = \log_2 N$ and

$$\begin{aligned} T(N) &= T(1) + kcN \\ &= kcN \\ &= cN \log_2 N \\ &= O(N \log N) \end{aligned}$$

An $O(n)$ method for MCS(X , low, high)

This method starts at $X[\text{low}]$ and scans through all the elements up to $X[\text{high}]$ keeping book of the maximum contiguous subvector found so far.



The maximum sum in the first i elements is either the maximum in the first $i - 1$ elements, called 'maxSoFar' or that of the subvector that ends in $X[i]$, called 'maxToHere'.

'maxToHere' contains the biggest sum *ending* at $X[i - 1]$. Considering item $X[i]$ take the biggest of $\text{maxToHere} + X[i]$ and 0.

```
1 int MCS(low, high)
2     maxSoFar = 0.0
3     maxToHere = 0.0
4     for i in [low, high]
5         maxToHere = max(maxToHere + X[i], 0.0)
6         maxSoFar = max(maxSoFar, maxToHere)
7     return maxSoFar
```

The 0–1 Knapsack Problem

Given any set S of n items where item i has a weight, $w_i > 0$, and has a benefit, $b_i > 0$, such that the total benefit is

$$B = \sum_{i \in S} b_i \text{ while } \sum_{i \in T} w_i < C,$$

where C is the maximum weight that is called the constraint or the capacity of the container.

If the number of objects is always an integer, this is called the 0–1 knapsack problem.

Let T denote the subset of items selected then the objective is to

$$\text{maximize } \sum_{i \in T} b_i \text{ while } \sum_{i \in T} w_i < C.$$

The 0–1 knapsack problem: example

Given a set S of n items where item i has a weight, $w_i > 0$, and has a benefit, $b_i > 0$, then select some items that maximize the total benefit B , with weight of at most $C = 9$.

	1	2	3	4	5
weight	4	2	2	6	2
benefit	R20	R3	R6	R25	R80

The solution is

$T = \{S_1, S_3, S_5\}$, with weight $C = 8$.

0–1 Knapsack algorithm: Try I

Let $S = \{S_k\}_{k=1}^n$ be a set of items and let
 $B[k] = \{\text{best selection from } S_k\}$.

This problem does not have *subproblem optimality*, for example the set

$$S = \{(3, 2), (5, 4), (8, 5), (4, 3), (10, 9)\}$$

of (benefit, weight) pairs with weight constrained by $C \leq 20$

Best selection from S_4 :

$$B[4] = \{(3, 2), (5, 4), (8, 5), (4, 3)\}$$

Best selection from S_5 :

$$B[5] = \{(3, 2), (5, 4), (8, 5), (10, 9)\}$$

Since we cannot derive $B[5]$ from $B[4]$ the problem does not possess *subproblem optimality*.

0–1 Knapsack algorithm: Try II

Let $S = \{S_k\}_{k=1}^n$ be a set of items and let

$B[k, C]$ = best selection from S_k with weight at most C .

This problem *has subproblem optimality*.

$$B[k, C] = \begin{cases} B[k-1, C] & \text{if } w_k > C \\ \max\{B[k-1, C], B[k-1, C - w_k] + b_k\} & \text{otherwise} \end{cases}$$

i.e., the best subset of S_k with weight at most C is

either the best subset of S_{k-1} with weight at most C

or the best subset of S_{k-1} with weight at most $C - w_k$ plus item k .

0–1 Knapsack problem: algorithm

$$B[k, C] = \begin{cases} B[k-1, C] & \text{if } w_k > C \\ \max\{B[k-1, C], B[k-1, C - w_k] + b_k\} & \text{otherwise} \end{cases}$$


- Recall the definition of $B[k, C]$
 - Since $B[k, C]$ is defined in terms of $[k-1, *]$, we can use two arrays of instead of a matrix
 - Running time: $O(nC)$.
 - Not a polynomial-time algorithm since C may be large
 - This is a *pseudo-polynomial* time algorithm
- b_i —a positive “benefit”
 w_i —a positive “weight”

0-1 Knapsack problem: algorithm

Listing 1: 0-1 Knapsack problem: algorithm

```
1 Algorithm 01Knapsack(S, C):  
2 In: set S of n items with benefit b_i  
3   and weight w_i; maximum weight C  
4 Out: benefit of best subset of S with  
5   weight at most C  
6 let A and B be arrays of length C + 1  
7 for w = 0 to C do  
8   B[w] = 0  
9 for k = 1 to n do  
10  copy array B into array A  
11  for w = w_k to C do  
12    if A[w-w_k] + b_k > A[w] then  
13      B[w] = A[w-w_k] + b_k  
14 return B[C]
```


Matrix multiplication




The ij th element of the $m \times n$ matrix $\mathbf{C}$, the product of the $m \times p$ matrix $\mathbf{A}$ and the $p \times n$ matrix $\mathbf{B}$ is:

$$\begin{aligned} C_{ij} &= A_{i1}B_{1j} + A_{i2}B_{2j} + \dots + A_{ip}B_{pj} \\ &= \sum_{k=1}^p A_{ik}B_{kj} \end{aligned}$$


$\mathbf{C}$ is calculated by running through all the possible combinations of i and j using two nested **for** loops.

Matrix multiplication



```
1 void matMultiply (double [][] A, double [][] B,  
2                   double [][] C, int m, p, n) {  
3   int i, k, j; double Cij;  
4  
5   for (i=1; i<=m; i++)  
6     for (j=1; j<=n; j++){  
7       Cij = A[i][1]*B[1][j];  
8       for (k=2; k<=p; k++)  
9         Cij += A[i][k]*B[k][j];  
10      C[i][j] = Cij; } }
```

Matrix multiplication



Suppose $\{\mathbf{A}_i\}_{i=1}^n$ is a chain of compatible matrices with dimensions: $d_i \times d_{i+1}$.
We want to compute:

$$\mathbf{C} = \mathbf{A}_1 \times \mathbf{A}_2 \times \dots \times \mathbf{A}_n$$

This is known as a chain product.

The dimension of the final product $\mathbf{C}$ is $d_1 \times d_{n+1}$.

Matrix Chain Products

Let $\{\mathbf{A}_\alpha\}_{\alpha=1}^n$ be a chain of compatible matrices with dimensions: $d_\alpha \times d_{\alpha+1}$.
We want to compute:

$$\mathbf{C} = \mathbf{A}_1 \times \mathbf{A}_2 \times \dots \times \mathbf{A}_n$$

The products are done pairwise: e.g. $\mathbf{P} = \mathbf{A}_\alpha \times \mathbf{A}_{\alpha+1}$ where the elements of $\mathbf{P}$ are

$$\mathbf{P}_{ij} = \sum_{k=1}^{d_{\alpha+1}} \mathbf{A}_{\alpha ik} \mathbf{A}_{\alpha+1 kj}$$

The final dimensions of $\mathbf{P}$ are $d_\alpha \times d_{\alpha+2}$.

The problem is: Given the dimensions $\{d_\alpha\}$, how should the corresponding matrices be paired?

$\mathbf{B}$ is 3×100 , $\mathbf{C}$ is 100×5 , $\mathbf{D}$ is 5×5 .

$(\mathbf{B} * \mathbf{C}) * \mathbf{D}$ takes $1500 + 75 = 1575$ ops

$\mathbf{B} * (\mathbf{C} * \mathbf{D})$ takes $1500 + 2500 = 4000$ ops

The Enumeration Approach is $O(4^n)$

Matrix Chain-Product Algorithm:

- Try all possible ways to parenthesize

$$\mathbf{A}_1 \times \mathbf{A}_2 \times \dots \times \mathbf{A}_n$$

- Calculate number of operations for each one
- Pick the one that is best
- Running time:
- The number of parenthesizations is equal to the number of binary trees with n nodes
- This is exponential.
- This is the *Catalan* number, which is almost 4^n .
- This algorithm takes too long.

A Greedy Approach

- Idea α : repeatedly select the product that uses (up) the most operations.
- Counter example:

A is 10×5 , **B** is 5×10 , **C** is 10×5 , **D** is 5×10 .

Greedy idea α gives

$(\mathbf{A} \times \mathbf{B}) \times (\mathbf{C} \times \mathbf{D})$,

which takes

$500 + 1000 + 500 = 2000$ operations.

But the alternative bracketing yields:

$\mathbf{A} \times ((\mathbf{B} \times \mathbf{C}) \times \mathbf{D})$ takes

$500 + 250 + 250 = 1000$ operations.

Another Greedy Approach

- Idea β : repeatedly select the product that uses the fewest operations.

- Counter example

A is 101×11

B is 11×9

C is 9×100

D is 100×99

- Greedy idea β gives

$\mathbf{A} \times ((\mathbf{B} \times \mathbf{C}) \times \mathbf{D})$, which takes

$109989 + 9900 + 108900 = 228789$ ops., but

$(\mathbf{A} \times \mathbf{B}) \times (\mathbf{C} \times \mathbf{D})$ takes

$9999 + 89991 + 89100 = 189090$ ops.

- The greedy approach β also does not yield the optimal value.

A Recursive Approach

Define subproblems: Find the *best* parenthesization of $\mathbf{A}_i \times \mathbf{A}_{i+1} \times \dots \times \mathbf{A}_j$

Let $N_{i,j}$ denote the number of operations done by this subproblem.

The optimal solution for the whole problem is $N_{1,n}$.

Subproblem optimality: The optimal solution can be defined in terms of optimal subproblems.

- There has to be a final multiplication, i.e., the root of the expression tree, for the optimal solution. Suppose the final multiplication is at index i :

$$(\mathbf{A}_1 \times \dots \times \mathbf{A}_i) \times (\mathbf{A}_{i+1} \dots \times \mathbf{A}_n)$$

- Then the optimal solution $N_{1,n}$ is the sum of two optimal subproblems, $N_{1,i}$ and $N_{i+1,n}$ *plus* the time for the final multiplication.
- If the global optimum did not have these optimal subproblems, we could define an even better *optimal* solution.

A Characterizing Equation


- Then the optimal solution $N_{1,n}$ is the sum of two optimal subproblems, $N_{1,i}$ and $N_{i+1,n}$ plus the time for the final multiplication.
- The global optimal has to be defined in terms of optimal subproblems, depending on where the final multiplication appears.
- Let us consider all possible places for the final multiplication: Recall that $\mathbf{A}_i$ is a $d_i \times d_{i+1}$ dimensional matrix. So, a characterizing equation for $N_{i,j}$ is the following:

$$N_{i,j} = \min_{i \leq k \leq j} [N_{i,k} + N_{k+1,j} + d_i d_{k+1} d_{j+1}]$$

Note that subproblems are not independent—the subproblems overlap.

Visualizing dynamic programming

$$N_{i,j} = \min_{i \leq k \leq j} [N_{i,k} + N_{k+1,j} + d_i d_{k+1} d_{j+1}]$$



First, fill in the $N_{i,i} = 0, \forall i \in [1..n-1]$.

Get $N_{i,j}$ from previous entries in row i and column j .

Filling in each entry in the N table takes $O(n)$ time.

Getting actual parenthesization can be done by remembering k for each N entry.

Total run time: $O(n^3)$. Space requirements n^2 .

A Dynamic Programming Algorithm

Since subproblems overlap, we can't use recursion. Instead, we construct optimal subproblems *bottom-up*. Start with $N_{i,i} = 0, \forall i \in [1..n]$; then do length 2,3, ... subproblems. Running time: $O(n^3)$.

```
1 Algorithm matrixChain(S):
2 # pre: sequence S of n matrices to be multiplied
3 # post: number of operations in an optimal
4 # parenthesization of S
5
6 for i in range(1, n):
7     N[i][i] = 0
8 for b in range(1, n):
9     for i in range(n-b):
10         j = i+b
11         N[i][j] = +infinity
12         for k in range(1, j):
13             N[i][j] = min(N[i][j], N[i][k] + N[k+1][j] +
14                           d[i]*d[k+1]*d[j+1])
```

Dynamic programming in general

- Applies to a problem that appears to need a lot of time, possibly exponential, provided we have:
- *Simple subproblems*: the subproblems can be defined in terms of a few variables, such as j , k , l , m , and so on.
- *Subproblem optimality*: the global optimum value can be defined in terms of optimal subproblems
- *Subproblem overlap*: the subproblems are not independent, but instead they overlap: hence they should be constructed bottom-up.

Substrings

- A substring of a character string $x_0x_1x_2 \dots x_{i_n-1}$ is a string of the form $x_{i_0}x_{i_1}x_{i_2} \dots x_{i_k}$, where $i_{j+1} = i_j + 1$ and $i_0 \geq 0$ and $i_k \leq i_{n-1}$.
- Example String: "ABCDEFGH IJK"
- Substring: "ABCDEFGH IJK"
- Substring: "DEFGH"
- Substring: ""
- Not substring: "DCEFGH IJ"

Subsequences—Section 11.5.1

- A subsequence of a character string $x_0x_1x_2 \dots x_{i_n-1}$ is a string of the form $x_{i_0}x_{i_1}x_{i_2} \dots x_{i_k}$, where $i_j < i_{j+1}$.
- Example String: "ABCDEFGHIIJK"
- Subsequence: "ACEGJIK"
- Subsequence: "DFGHK"
- Not subsequence: "DAGH"
- Subsequence, but also a substring: "DEFGHI"
- A subsequence is *not the same* as substring.

The Longest Common Subsequence (LCS)

- Given two strings X and Y , the longest common subsequence (LCS) problem is to find a longest subsequence common to both X and Y
- "ABCDEFGH" and "XZACKDFWGH" have "ACDFG" as a longest common subsequence
- Note that a *subsequence* is not a *substring*:
"ZACK" is not only a *substring* but it is also a *subsequence* of "XZACKDFWGH" a subsequence.
- Every *substring* of T is a *subsequence* of T .

The Longest Common Subsequence (LCS)—applications

- LCS has applications to DNA similarity testing—alphabet is $\{\text{A, C, G, T}\}$
- LCS is used by spelling checkers to find the best word to fit your spelling mistake.
- LCS is used by `diff` in Linux/Unix to find differences between two files.

A Poor Approach to the LCS Problem

A Brute-force solution:

- Enumerate all subsequences of X
- Test which ones are also subsequences of Y
- Pick the longest one.

Analysis:



- If X is of length n , then it has 2^n subsequences
- This is an exponential-time algorithm

LCS solved with Dynamic Programming

- Define $L[i, j]$ to be the length of the longest common subsequence of $X[0..i]$ and $Y[0..j]$.
- Allow for -1 as an index, so $L[-1, k] = 0$ and $L[k, -1] = 0$, to indicate that the null part of X or Y has no match with the other.

Then we can define $L[i, j]$ in the general case as follows:

- if $x_i = y_j$, then $\backslash \backslash$ add this match
$$L[i, j] = L[i - 1, j - 1] + 1$$
- if $x_i \neq y_j$, then $\backslash \backslash$ no match here
$$L[i, j] = \max\{L[i - 1, j], L[i, j - 1]\}$$



An LCS Algorithm—in pseudocode

$$L_{i,j} = \begin{cases} \text{for } x_i = y_j & L_{i-1,j-1} + 1 \\ \text{otherwise} & \max\{L_{i-1,j}, L_{i,j-1}\} \end{cases}$$

L	-1	0	1	2	3			j	...	n-1	
-1	0	0	0	0	0	0	0	0	0	0	0
0	0										
1	0										
2	0										
3	0										
	0										
i	0										
	0										
...	0										
n-1	0										

LCS Algorithm—closer to Java

Listing 2: LCS Java-like algorithm— $O(N^2)$

```
1 char [] void LCS(char X[], int n, char Y[], int m){
2 //pre: Strings X of length n, and Y of length m.
3 //post: For i in [1..n-1], j in [1..m-1]:
4 // L[i,j] = length of LCS of X[0..i] and Y[0..j]
5 int L[][] = new int [n][m];
6     for (i=1; i < n; i++)
7         L[i][0] = 0;
8     for (j=0; j < m; j++)
9         L[0][j] = 0;
10    for (i=1; i < n; i++)
11        for (j=1; j < m; j++)
12            if (x[i] == y[j])
13                L[i][j] = L[i-1][j-1] + 1;
14            else
15                L[i][j] = max(L[i-1][j], L[i][j-1]);
```


Listing 3: LCS—backtracking

```
1 // Backtrack for longest common subsequence
2 int length = L[n-1][m-1];
3 char lcs[] = new char[length+1];
4     for (i = n-1, j = m-1; i > 0 && j > 0 && length > 0; )
5         if (x[i] == y[j]) {
6             lcs[length--] = x[i--];
7             j--; }
8         else if (L[i-1][j] > L[i][j-1])
9             i--;
10        else
11            j--;
12
13    return lcs; }
```

Analysis of LCS Algorithm

- The simple loops take $O(m)$ and $O(n)$ each.
- Outer nested loop takes $O(n)$ and repeats the inner loop $O(m)$ times.
- The body of the inner loop takes $O(1)$, i.e., it is independent of n or m .
- The backtracking is $O(1)$ per step and is at most $\max(O(n), O(m))$, therefore it can be ignored.
- So, the total running time is $O(nm)$ and the space required is also $O(nm)$.
- The result is in $L[n, m]$ from which the subsequence may be recovered.

Traversal of trees




Searching for a node, given its key, in a *balanced* binary search tree (BST), tends to be logarithmic— $O(\log n)$.

L indicates a move to the **Left** subtree. **R** moves to the **Right** subtree and **N** denotes a visit to the **Node** itself.

The **LNR**—in-order—traversal of the above tree is:

4, 7, 8, 10, 17, 19, 21, 22, 23, 25, 27, 29, 35, 37, 41, 47.

Traversal of trees



The 6 standard orders of traversal when searching for a node in such trees are: **NLR**—pre-order, **LNR**—in-order, **LRN**—end-order or post-order, **NRL**, **RNL** and **RLN**.


The **LNR**—in-order—traversal of the above tree is:

4, 7, 8, 10, 17, 19, 21, 22, 23, 25, 27, 29, 35, 37, 41, 47.

Which order will visit all the nodes in the opposite order?

Traversal of trees—LRN

The following tree *is* a BST.



The **LNR**—in-order traversal of the above tree is:

4, 7, 8, 10, 17, 19, 21, 22, 23, 25, 27, 29, 35, 37, 41, 47.


Which ordering traverses the keys from the largest to the smallest?

What order does **LRN** yield?

What order does **NLR** yield?

Tree traversal—examples using a non BST


The following tree *is not* a BST.



- The **LNR**—in-order traversal of the above tree is:
3, 2, 4, 1, 6, 5, 7.
- The **RNL**—RL-in-order traversal of the above tree:
7, 5, 6, 1, 4, 2, 3.
- The **NLR**—pre-order traversal of the above tree:
1, 2, 3, 4, 5, 6, 7.
- The **RLN**—RL-end-order or post-order traversal of the above tree is:
7, 6, 5, 1, 4, 3, 2.

Traversal of trees—examples using a BST

This tree *is* a BST.



- The **LNR**—in-order traversal of the above tree is:
1, 2, 3, 4, 5, 6, 7.
- The **RNL**—RL-in-order traversal of the above tree:
7, 6, 5, 4, 3, 2, 1.
- The **NLR**—pre-order traversal of the above tree is:
4, 2, 1, 3, 6, 5, 7.
- The **LRN**—end-order or post-order traversal of the above tree is:
1, 3, 2, 5, 7, 6, 4.

Tree traversal—breadth-first and depth-first

- Another approach is to order the nodes according methods that lump the children of a node together.

TRAVERSETREE

insert the root node into list.

do


visit and *remove* the head of list:

visit by inserting head's children in the list.

repeat until the list is depleted.

- The list uses a
 1. **stack** for **depth-first** traversal.
 2. **queue** for **breadth-first** order.

Tree traversal—depth-first



DEPTHFIRSTTREETRAVERSAL

push the root node onto stack.

do

visit and *remove* the top of the stack:


visit by pushing top's children onto the stack.

repeat until the stack is empty.

The **depth-first** traversal of the above tree:

4, 2, 1, 3, 6, 5, 7.

Tree traversal—breadth-first



Using a **queue** instead of a *stack* turns the algorithm into **breadth-first traversal**.

BREADTHFIRSTTREETRAVERSAL

add the root node into the queue.

do

visit and *remove* the front of the queue:


visit by adding front's children into the queue.

repeat until the queue is empty.

The **breadth-first** traversal of the above tree:

4, 2, 6, 1, 3, 5, 7.

Tree traversal—depth-first recursively




DEPTHFIRSTTRAVERSAL

```
void DFS (tree  $T$ , node  $v$ ) {  
    mark the node  $v$ .  
    call prelude( $v$ );  
    for all edges  $\{(v, z) | z \in \text{childOf}(v)\}$   
        if ( $z$  is unmarked) DFS ( $T$ ,  $z$ );  
        call postlude( $v, z$ );}
```

The **depth-first** traversal of the above tree:
4, 2, 1, 3, 6, 5, 7.

Graphs—edges, nodes, degree

A **graph** $G = (V, E)$ is a set V of **vertices** or **nodes** and a set E of **edges**. Each edge connects two vertices. A **self-loop** connects a node to itself.



The edges represent the bridges and the nodes represent land in Euler's Königsberg bridges problem:

Is there a path where each bridge is traversed exactly once?

Edges may be **directed** or **undirected**.

Each vertex has a **degree**, $d(v)$ is number of edges incident to v .

$d(A) = 5$ and $d(B) = 3$.

In a directed graph the **outdegree** is the number of edges that leave the node, and **indegree** is the number of vertices that enter the node.

Graphs—path, cycle, reachable, tree

A **path** from n_1 to n_k is a sequence of nodes $n_1, n_2, \dots, n_k$ that are connected by edges

$$(n_1, n_2), (n_2, n_3), \dots, (n_{k-1}, n_k).$$

In a **simple path** each node appears only *once* and u is **reachable** from v if there is a path from u to v .

A **circuit** or **cycle** is a path whose first and last vertices are the same.

The **undirected form** of a directed graph $G = (V, E)$ is the same graph without directions on the edges.

A graph in its undirected form is called **connected** if there is a path from any node to any other node.


A **forest** is graph that does not contain a cycle.

A **tree** is a connected forest and a **rooted tree** is a directed tree with one distinguished node called the **root** such that all its edges point away from the root.

Graphs—Eulerian graphs

Euler's Königsberg bridges problem:

Which graphs can be traversed with a circuit that visits each edge exactly once?






In an **Eulerian graph** the degree of EVERY NODE MUST BE EVEN.

Each vertex has a **degree**, $d(v)$ is number of edges incident to v . $d(A) = 5$, $d(B) = 3$, $d(C) = 3$, $d(D) = 3$.

So the Euler's Königsberg bridges problem is **not** Eulerian.

Graphs—spanning tree, subgraph

A **subgraph** of a graph $G = (V, E)$ is a graph $H = (U, F)$ such that $U \subseteq V$ and $F \subseteq E$.









Graphs—types I

A graph G may be
undirected or *directed*,

unweighted or *weighted*,

unlabelled or *labelled*,









Graphs—types II

A graph G may be
cyclic or *acyclic*, i.e. a DAG or directed, acyclic graph,


sparse or *dense*,

simple or *non-simple*—with multi-edges and loops,




Graphs—types III

A graph G may be *unconnected* or *connected* .



A directed graph G may be *cyclic* or *acyclic*.

Graphs—types IV




Graphs—spanning tree, subgraph

A **spanning forest** of an undirected graph G is a subgraph of G that is a forest and that contains all the vertices of G .

A **vertex-induced subgraph** of a graph $G = (V, E)$ is a subgraph $H = (U, F)$ such that $U \subseteq V$ and F consists of all the edges of E both of whose vertices belong to U .

A **bipartite graph** is a graph whose nodes can be divided into two sets such that all edges connect nodes from one set to nodes in the other set.




Graph traversal—depth-first

The idea is to begin by visiting some start or **root** node, marking it and then calling DFS using each neighbour one-by-one as the root.

Depth-first search (DFS) in an undirected simple graph starts at a node we call the **root**.

```
void DFS (graph  $G$ , node  $v$ ) {  
    mark the node  $v$ .  
    call  $prelude(v)$ ;  
    for all edges  $\{(v, z) | z \in$   
        neighbourOf( $v$ ) $\}$   
        if ( $z$  is unmarked)  
            DFS ( $G$ ,  $z$ );  
    call  $postlude(v, z)$ ;  
}
```



Every node is visited once.


Graph traversal—depth-first iteratively

Here, again, the idea is to visit every node once.

Depth-first search (DFS) in an undirected simple graph starts at a node we call the **root**.

Mark the node visited and consider all its neighbours by *stacking* them.

Continue by visiting the top node in the stack, marking it, stacking its neighbours if they have not been marked. Continue until the stack is empty.



Graph traversal—depth-first

Depth-first search (DFS) in an undirected simple graph starts at a node we call the **root**.


push the root node onto stack.

do


visit, mark and *remove* the top of the stack:

visit the node by working with it and by pushing top's unmarked neighbours onto the stack.

repeat until the stack is empty.



Graph traversal—depth-first



Some depth-first traversals of this graph are:

a, b, e, h, k, j, i, f, c, d, g and

a, d, f, i, j, k, g, e, h, c, b.

Are these paths both depth-first traversals?

Digraph—depth-first traversal

Depth-first search (DFS) in an directed simple graph starts a node we call the **root**.


push the root node onto stack.

do

visit, mark and *remove* the top of the stack:


visit the node by working with it and by pushing top's unmarked reachable neighbours onto the stack.

repeat until the stack is empty.



Depth-first order


Directed acyclic graph—Topological sorting



Let $\vec{G}$ be a digraph with n nodes. The edges above represent **precedence**, ' $\prec$ ', such that $a \prec b, b \prec e, a \prec d, d \prec e, d \prec f, a \prec f, a \prec c, c \prec f, \dots$

The figures show two different topological orderings of the same graph.

Directed acyclic graph—Topological sorting



Let $\vec{G}$ be a digraph with n nodes. The edges above represent **precedence**, ' $\prec$ ', such that $a \prec b, b \prec e, a \prec d, d \prec e, d \prec f, a \prec f, a \prec c, c \prec f, \dots$

A **topological ordering** of $\vec{G}$ is an ordering $v_1, \dots, v_n$ of the vertices of $\vec{G}$ such that for every edge (v_i, v_j) of $\vec{G}$, $i < j$, i.e. any directed path in $\vec{G}$ traverses nodes in precedence order.

A directed graph $\vec{G}$ has a topological ordering if and only if it is acyclic.

Directed acyclic graph—Topological sorting

```
1 void topologicalSort (graph G) {  
2   // pre: acyclic digraph  
3   // post: topological ordering  
4   S = initially empty stack;  
5   for all u in G.nodes() do  
6     if (inDegree(u) == 0)  
7       S.push(u);  
8   i = 1;  
9   while !S.isEmpty() do  
10    u = S.pop();  
11    i++;  
12    for all outgoingEdge(u,w) do  
13      inDegree(w)--;  
14      if (incounter(w) == 0)  
15        S.push(w); }
```


Shortest paths

Let $G = (V, E)$ be a weighted graph where V is a set of nodes or vertices, E is a set of weighted nodes.

Each edge may be labelled with a *length* or *weight*.

The length between two nodes v_i and v_{i+1} is indicated by $\text{length}(v_i, v_{i+1})$.

The *length* of a path

$$P = ((v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k))$$

in the graph is defined as

$$\text{length}(P) = \sum_{i=0}^{k-1} \text{length}(v_i, v_{i+1})$$

The *distance*, $d(v, u)$ from vertex v to u in a graph G is the shortest path length from v to u , if it exists.

Shortest paths

When no path exists from v to u we write

$$d(v, u) = \infty.$$

Negative weights are possible but can lead to unforeseen problems.
We will regard the distance as a positive metric.

Single-source shortest path

Dijkstra's algorithm for an undirected simple graph finds the shortest path from a node v to every other node in the graph.

Each vertex $w \in G$ is marked with a label $w.SP$ which is the shortest path (v, w) attained so far.

Initially $v.SP = 0$ and $w.SP = \infty$, for each $w \neq v$.

We use a heap H to store all the nodes according to their shortest-path distance from v . Initially v is at the top of the heap H .

Now consider node w at the top of the heap and its nearest neighbour is entered onto the heap.

Dijkstra's algorithm for shortest paths

```
1 void shortestPaths (graph G, node v) {  
2   // pre:  $G = (V, E)$  is a weighted directed graph,  $v$  is the source node.  
3   // post: for each  $w$ ,  $w.SP$  is the shortest path  $(v, w)$ .  
4   i = 0;  
5   for w in G do  
6     w.mark = false;  
7     w.SP = infty;  
8     Heap[++i]=w;  
9   v.mark = true; v.SP = 0; Heap[0]=v;  
10  // Enter all nodes  $w$  into a heap  $H$  according to  $w.SP$ ;  
11  while w in Heap && w.mark == false do  
12    get w with minimal w.SP off heap;  
13    w.mark = true;  
14    for all edges  $(w, z)$  && z.mark == false do  
15      if  $(w.SP + \text{length}(w, z) < z.SP)$   
16        z.SP = w.SP + length(w, z);  
17        update z in Heap; }
```

Dijkstra's algorithm for shortest paths

Updating the length of a path **z.SP** takes $O(\log m)$ comparisons, where m is the size of the heap.

There are $|V|$ deletions from the heap and at most $|E|$ updates.

So there are $O(|E| \log |V|)$ comparisons in the heap, giving a running time of $O((|E| \log |V|) \log |V|)$


Putting $|V| = n$ and $|E| = m$, the running time is $O((m \log n) \log n)$

Minimal spanning trees (MST)

- A *forest* is a loop-free graph.
- A *tree* is a connected forest.
- A subgraph *spans* a graph if it contains all the vertices of the graph.

Minimal spanning trees (MST)

Recall that a spanning tree is a tree that connects all the vertices V of a graph $G = (V, E)$ using edges from E .



Several of 7 edges can be switched, so that the MST need not be unique.

A minimal spanning tree (MST) uses edges optimally such that the sum of their weights is a minimum.

The total cost is $8 + 5 + 7 + 6 + 7 + 9 + 4 + 6 + 3 + 3 = 58$.

The MST is unique when all the weights are different.

There will be $n-1$ edges in the MST for a graph with n nodes.

Minimal spanning trees (MST)

The three algorithms we will consider are called

- Otakar Borůvka's algorithm (1926)
Used to set up electrical network in Moravia.
- Joseph Kruskal algorithm (1956) and the
- Prim-Jarník algorithm (1930, 1957)

There are some others too.

Borůvka's algorithm for MSTs

```
1 //Pre: Connected weighted graph  $G = (V, E)$ ,  $n = |E|$ 
2 //Post: Minimal spanning tree in  $T$ 
3  $T = (V, \emptyset)$  // all the vertices of  $G$  but no edges
4 while number of edges in  $T < n - 1$  do
5      $E = \{\}$  // an empty set of edges
6     for each component  $C$  in  $T$  do
7          $S = \{\}$  // an empty set of edges
8         for each vertex  $v$  in  $C$  do
9             add minimal edge  $(v, u)$  where  $u \notin C$  to  $S$ 
10        add the cheapest edge in  $S$  to  $E$ 
11    add edges  $E$  of to  $T$ .
```

If $v = |V|$ is the number of vertices in the graph G . Since there are $e = |E|$ edges the outer loop of Borůvka's algorithm clearly takes $O(e)$ steps. The inside loops take at most $v/2, v/2^2, v/2^3$, etc., steps, i.e., $\log v$ steps, so the algorithm runs in time $O(e \log v)$.

Borůvka's algorithm for MSTs

An alternate statement of Borůvka's algorithm for MSTs is

```
1  Given  $G = (V, E)$ 
2   $T =$  graph consisting of  $V$  with no edges
3  while  $T$  has  $< n-1$  edges do
4      for each connected component  $C$  of  $T$  do
5           $e =$  minimum cost edge  $(v, u)$  so that  $v$  in  $C$  and  $u$  not in  $C$ 
6           $T = T$  union  $\{e\}$ 
```

The union-find algorithm

The union-find algorithm deals with adding edges to a set of v vertices or nodes and then finding out whether two nodes are connected.

- **union** command connects two nodes
e.g. **union**(u , v) links node u to node v if they are not already connected.
- **find** or **isConnected** determines if two nodes are linked. e.g.
isConnected(u , v) is a predicate that checks if node v is linked to node u .

The union-find algorithm

Consider the disconnected graph with no edges and eight nodes

Before any operations



`isConnected(a, f)  $\Rightarrow$  false`

`union(b, f)`



`isConnected(a, f)  $\Rightarrow$  false`

`union(a, b)`



`isConnected(a, f)  $\Rightarrow$  false`

`union(a, e)`



`isConnected(a, f)  $\Rightarrow$  true`

The union-find algorithm

`isConnected(u, v)` is an *equivalence relation* since it is:


reflexive: `isConnected(u, u)`,

symmetric: `isConnected(u, v)  $\equiv$  isConnected(v, u)`, and

transitive: if `isConnected(u, v)` and `isConnected(v, r)` then `isConnected(u, r)`.

Equivalence relations divide sets into equivalence classes.

The graph depicted by



has three equivalence classes $\{a, b, e, f\}$, $\{c\}$, $\{d, g, h\}$.

These classes consist of connected components.

A Java API for a union-find data type

A data structure for union-find where the size of the graph/set N can be very large and the number of operations M can be very large.

The `isConnected` and `union` commands can appear in any order.

```
1 public class UF
2 UF(int N) // initialize union-find data structure with N objects [0..N - 1]
3 void union(int u, int v) // add connection between u and v
4 boolean isConnected(int u, int v) // u and v in same component?
5 int find(int u) // component identifier for u ∈ [0..N - 1]
6 int count() // number of components
```

The nodes `p` and `q` are connected if and only if `id[p]` equals `id[q]`, thus and all connected components have the same `id[.]` value.

The code for `find`

```
1 public int find(int p) {
2     return id[p]; }
```

Code for **union** and **isConnected**

The code for union

```
1 public void union(int p) {  
2     int pID = find(p),  
3     qID = find(q);  
4     if (pID == qID) return;  
5  
6     for (int i = 0; i < id.length; i++)  
7         if (id[i] == pID) id[i] = qID;  
8     count--; }
```

Kruskal's algorithm for MSTs

- 1 Repeat the following step as many times as possible:
- 2 Among the edges of G not yet chosen,
- 3 choose the shortest edge which does not form any loops,
- 4 with those edges already chosen.

Clearly the set of edges eventually chosen forms a spanning tree of G , and in fact it forms a minimal spanning tree.¹⁴

¹⁴Joseph B. Kruskal, Jr. *On the shortest spanning subtree of a graph and the traveling salesman problem*, Proceedings of the American Mathematical Society, Vol. 7, No. 1 (Feb., 1956), pp. 48–50.

Longest maximum subvector problem solved in 5 ways

In a list or vector X of N integers find the maximum sum of any contiguous subvector.¹⁵

-66	-15	87	91	32	5	-47	47	4	-77
0	1	2	3	4	5	6	7	8	9

In the list X above the subvector $X[2..8]$ has the maximum sum 219.

Clearly if all the numbers are negative then the maximum contiguous subvector (MCS) is simply the largest number. In this case we define the MCS as 0.

¹⁵This example comes from pages 77–81, “Programming Pearls”, Jon Bentley, ACM Press and Addison-Wesley, New York, 2000.

An $O(n^3)$ method for MCS

In our first attempt we simply calculate *all* possible subvectors $X[\text{low}..\text{high}]$ and find the greatest.

```
1 int max_n^3 (x)
2     n = len(x)
3     maxsofar = 0
4     for low in [0,n)
5         for high in [low,n)
6             sum = 0
7             for r in [low,high)
8                 sum += x[r]
9                 if (sum > maxsofar)
10                     maxsofar = sum
11 return maxsofar
```

Clearly three nested loops yield an $O(N^3)$ method.

There is a simple way to reduce this to only two loops by preserving partial sums.

An $O(n^2)$ method for MCS—1

In our second attempt again calculate all the subvectors $X[\text{low}..\text{high}]$ but use partial sums to determine them and find the greatest.

```
1 int max_n^2(x)
2     n = len(x)
3     maxsofar = 0
4     for low in [0,n)
5         sum = 0
6         for high in [low,n)
7             sum += x[r]
8             if (sum > maxsofar)
9                 maxsofar = sum
10    return maxsofar
```

There are now only two nested loops clearly making it an $O(N^2)$ method.
There is another way of using only two loops.

Another $O(n^2)$ method for MCS—2

In our third attempt we first calculate the cumulative sums for for each possible range of numbers subvectors $X[0..i]$ where where $i \in [0..N]$


```
1 int max_n^2(x)
2     n = len(x)
3     sumTo[0] = 0
4     for i in [0,n)
5         sumTo[i] = sumTo[i-1] + X[i]
6     maxsofar = 0
7     for low in [0,n)
8         for high in [low,n)
9             sum = sumTo[high] - sumTo[low-1]
10            // sum of all elements in X[low..high)
11            if (sum > maxsofar)
12                maxsofar = sum
13 return maxsofar
```

With two nested loops, this is still $O(N^2)$.

The next method is subquadratic.

A divide-and-conquer method for MCS


First, we divide the given array $X[]$ into two parts, **Left** and **Right**:



and then find the MCS **mLeft** in **Left** and the MCS **mRight** in **Right**.



There is still a snag—the MCS might straddle both **Left** and **Right**.



The MCS is the subvector that gives the biggest of these three values.

The final algorithm is (1) $\text{middle} = N/2$,

(2) Get MCS, **mLeft** of $X[\text{low}..\text{middle}]$, using

$\text{mLeft} = \text{MCS}(X, \text{low}, \text{middle})$, and get the MCS, **mRight** of $X[\text{middle}+1..N]$, using $\text{mRight} = \text{MCS}(X, \text{middle}+1, N)$.

More difficult is getting the straddling MCS, **mStraddle**.

Calculating the straddling MCS

The idea of the straddling MCS is first to calculate the MCS, called ‘maxsofarLeft’ that ends at $X[\text{middle}]$ but starts somewhere in **Left**—running backwards, and also calculating the MCS called ‘maxsofarRight’ that starts at $X[\text{middle}+1]$ and ends somewhere in **Right**. The sum of these two values is **mStraddle**.

```
1 int maxStraddle(low, high)
2     middle = (low + high) / 2
3     sum = 0; maxsofarLeft = 0
4     for i in [low..middle]
5         sum += X[middle - i + 1]
6         if sum > maxsofarLeft
7             maxsofarLeft = sum
8     sum = 0; maxsofarRight = 0
9     for i in [middle+1..high]
10        sum += X[i]
11        if sum > maxsofarRight
12            maxsofarRight = sum
13    return maxsofarLeft + maxsofarRight
```

A divide-and-conquer method for MCS

```
1 int MCS(low, high)
2     if low > high
3         return 0
4     if low == high
5         return max(0, X[low])
6     middle = (low + high) / 2
7     mLeft = MCS(X, low, middle)
8     mRight = MCS(X, middle+1, high)
9     mStraddle = maxStraddle (X, low, high)
10    return max (mLeft, mRight, mStraddle)
```

Suppose ‘MCS’ takes time $T(N)$ for an input of N . Then ‘mLeft’ takes time $T(N/2)$ and ‘mRight’ takes time $T(N/2)$. ‘maxStraddle()’ takes $O(N)$ because, at worst, it scans the entire input once. We conclude that

$$\begin{aligned} T(N) &= 2T(N/2) + O(N) \\ &= O(N \log N) \end{aligned}$$

Divide-and-conquer methods— $O(N \log N)$

Suppose the divide-and-conquer method has time complexity $T(N)$ when it has an input of size N .

Divide-and-conquer methods entail:

- (1) *dividing* the problem into equal parts,
- (2) *solving each part* of the problem separately, and
- (3) *joining* the results somehow.

1. *Dividing* the problem in two, has the low cost of $O(1)$. All that needs to be done is $N = N/2$.
2. *Solving each half* of the problem separately costs $2T(N/2)$.
3. *Joining the two results* is the critical part. Try to make this depend *linearly* on N . If it is sublinear all the better. Usually this part entails looking at every input once. In mergesort the two sorted halves are merged in time cN , making the cost of merge sort, $T(N) = 2T(N/2) + cN$.
4. In MCS straddle passes through all N elements once, giving the same formula.

Divide-and-conquer methods— $O(N \log N)$ —I

Solving the recursive equation for merge sort and MCS

$$T(N) = 2T(N/2) + cN.$$

is quite easy.

$$\begin{aligned} T(N) &= 2T(N/2) + cN \\ &= 2[2T(N/4) + cN/2] + cN \\ &= 2^2T(N/2^2) + cN + cN \\ &= 2^2[2T(N/2^3) + cN/2^2] + cN + cN \\ &= 2^3T(N/2^3) + cN + cN + cN \\ &= 2^3T(N/2^3) + 3cN \\ &\vdots \\ &= 2^kT(N/2^k) + kcN \end{aligned}$$

Continue until $N/2^k = 1$.

Divide-and-conquer methods— $O(N \log N)$ —II

$$\begin{aligned} T(N) &= 2T(N/2) + cN \\ &\vdots \\ &= 2^k T(N/2^k) + kcN \end{aligned}$$

Continue until $N/2^k = 1$.

Since there is nothing to do when the input is of size 1, it follows that $T(1) = 0$, so we set $N/2^k = 1$ giving $N = 2^k$. This happens when $k = \log_2 N$ and

$$\begin{aligned} T(N) &= T(1) + kcN \\ &= kcN \\ &= cN \log_2 N \\ &= O(N \log N) \end{aligned}$$

An $O(n)$ method for MCS(X , low, high)

This method starts at $X[\text{low}]$ and scans through all the elements up to $X[\text{high}]$ keeping book of the maximum contiguous subvector found so far.



The maximum sum in the first i elements is either the maximum in the first $i - 1$ elements, called 'maxSoFar' or that of the subvector that ends in $X[i]$, called 'maxToHere'.

'maxToHere' contains the biggest sum *ending* at $X[i - 1]$. Considering item $X[i]$ take the biggest of $\text{maxToHere} + X[i]$ and 0.

```
1 int MCS(low, high)
2     maxSoFar = 0.0
3     maxToHere = 0.0
4     for i in [low, high]
5         maxToHere = max(maxToHere + X[i], 0.0)
6         maxSoFar = max(maxSoFar, maxToHere)
7     return maxSoFar
```


A review of Trees, lattices and graphs

Graphs, trees and lattices are ubiquitous in computation. Under the cover of finite state machines graphs can compute almost anything.

- searching states spaces of games such as draughts, chess, $n \times n$ tiles, etc.
- optimizing reticulation of water, power, telecoms
- designing optimal VLSI chips
- navigation of your GPS, shipping
- proving theorems
- representing knowledge
- currency arbitrage transactions
- spelling checkers
- optimizing typesetting layout
- traffic robot control
- texture mapping

Trees, lattices and graphs


- A *tree* is an acyclic *graph*.
- The edges are usually referred to as *branches*, and the terminal nodes without any sub branches are called *leaves*.
- The branches may be *directed* or *undirected*.
- A directed tree has a unique *root*.
- A directed binary tree.



- If a node has more than one parent it is in a network or graph.

Trees, lattices and graphs


- Edges may be labelled with *weights*.
- A *connected* graph has links to all its nodes.



- A *lattice* is a directed, connected, acyclic graph.
- A tree with a single path from its root to its only leaf is a *degenerate* tree.

Decision trees


- Decision trees or decision lattices are useful for classification.
A part of “guess the animal” follows.



- Classification by a binary decision tree works very fast—in $\log_2(n)$ time when it is reasonably balanced.

State and problem spaces


- A *state space* is a graph that is used to map all the possible states of a system. It shows all the possible transitions between states.
- The following is a state space for a vending machine—*circa* 1990—that accepts ‘quarter’ (25 cents) ‘nickel’ (5 cents) coins.



- The machine is greedy—it accepts 55 cents or more.
- This state space is a *finite-state machine*.

Finite-state machine (FSM)


- A FSM that checks for the strings 'while', 'write' and 'begin'



- Note that not all the errors should be trapped by the 'Error' state.
- This kind of FSM can be used as the basis of a very efficient spelling checker. Once it has been set up it accepts a word present in the dictionary in time proportional to the word's length independent of the size of the dictionary.
- A state diagram can be used to represent the solutions to a problem. This is known as a problem space.
- Some states correspond to answers or solutions—there may be many solution states.

Problem and state spaces


The monkey and bananas problem is a classic example: only a few instructions are needed: jump off couch, move to ladder, move ladder under bananas, climb ladder, and grab bananas.




Various orders of instructions may be followed.

Searching for optimal paths in a graph

- The *Travelling Salesman Problem (TSP)*:
Find the optimal closed path in which a salesman visits each node once and returns to the start. It is displayed as a network.




- This problem can be solved by a *search tree*. The solution is displayed as a tree. The path with the lowest cost is the winner.



And-or trees and goals

- An AND-OR lattice or AND-OR tree is useful for representing backward-chaining problems.



- It is possible to use AND, OR and NOT gates to represent AND-OR lattices with hardware.

Trees in theorem proving

- An unexpected use of a tree. Consider

Some programmers hate failures.

No programmer hates any success.

No failure is a success.

- We need these predicates:

$P(x)$ means x is a programmer.

$F(x)$ means x is a failure.

$S(x)$ means x is a success.

$H(x, y)$ stands for x hates y .

- Formulated in predicate logic:

(1) $\exists x[P(x) \wedge \forall y(F(y) \Rightarrow H(x, y))]$

(2) $\forall x[P(x) \Rightarrow \forall y(S(y) \Rightarrow \neg H(x, y))]$

(3) $\neg \forall y[F(y) \Rightarrow \neg S(y)]$


- These are called *well-formed formulas* or *wffs*.

Resolution theorem proving—I

- In *resolution theorem proving* we can show that formulas are unsatisfiable.
- The axioms are the clauses of the formula.
- In an inference rule, if clauses $(x \wedge \alpha)$ and $(\neg x \wedge \beta)$ are already present, then a new clause $(\alpha \wedge \beta)$ maybe added.
- *Unsatisfiability* is proven when an empty clause is derived.
- Every unsatisfiable formula has a resolution proof.
- Resolution proofs for some formulas may take exponential time.

Resolution theorem proving—II

- In *resolution theorem proving* some simple steps can be followed to turn the well-formed formulas step-by-step into a form that can then be used to decide the formulas using a tree.
- The *resolution refutation tree* for a proof of “no failure is a success” follows



- The clauses $\neg F(y) \vee H(a, y)$ and $F(b)$ are resolved by taking the b of $F(b)$ as an instance of y in $\neg F(y) \vee H(a, y)$ giving $H(a, b)$.
- The *nil* means that the formula is TRUE.
- Note that other trees are also possible.

Searching in graphs



We discuss the following methods to search for a goal in graphs. The aim is simply to find the goal.

- *Depth-first search*: good when blind alleys do not get too deep.
- *Breadth-first search*: number of branches is not too large.
- When goal distance can be estimated easily.
 - *Hill-climbing with depth-first search*: good choice is likely in the good-looking alternatives.
 - *Beam search*: good-looking paths are likely to lead to the goal at all levels.
 - *Best-first search*: good paths may look bad at shallow levels.
- *Branch-and-bound search*: expands the node at the end of the shortest path leading to an open node—it expands the least-cost partial path.

Searching in graphs

Winston's graph from Artificial Intelligence (1984), p. 89, illustrates a number of graph searching algorithms.


The goal is to reach node G starting at node S . The full *depth-first tree* follows.



There are four successful paths $S-G$. The shortest is $S-D-E-F-G$, with a length of 13.



Depth-first searching in graphs

Depth-first search dives down the branches of the search tree as far as it can go before backtracking to nodes on a higher level.



Using *depth-first search* the first successful path—of length 19—is found after visiting 8 nodes.

Searching in graphs





Using *depth-first search* the first successful path—of length 19—is found after visiting 8 nodes.

Using *breadth-first search* the first successful path—of length 13—is found after visiting 22 nodes. The optimal path has been found accidentally.

Notice that depth-first search does *not* find the optimal path. Neither is it fast. If the search tree has deep branches that do not contain the goal much time may be wasted.

Searching in graphs




1. Depth-first search

```
def DFS (graph, S, G):  
    found_goal = False  
    initialize stack to root node, S  
    while stack is not empty,  
        and not found_goal:  
        E = pop stack()  
        if E == G:  
            found_goal = True  
        else :  
            push E's children onto stack  
    return found_goal
```



Breadth-first searching in graphs

Breadth-first search sweeps across levels in the search tree of the graph, or network until it reaches the goal.



Using *breadth-first search* the first successful path—of length 13—is found after visiting 22 nodes.

Searching in graphs




2. Breadth-first search

```

def BFS (graph, S, G):
    found_goal = False
    initialize queue to root node, S
    while queue is not empty,
        and not found_goal:
        E = pop_queue()
        if E == G:
            found_goal = True
        else :
            add E's children to back of queue
    return found_goal


```

Searching in graphs—hill climbing





Hill climbing does a depth-first search guided by an “as the crow flies” estimate to the goal.

The nodes are sorted according to the distance left to the goal. At each step the shortest path remaining is chosen.



Searching in graphs—hill climbing



3. Depth-first search with hill climbing


```
def DFS (graph, S, G):  
    found_goal = False  
    initialize stack to root node, S  
    while stack is not empty,  
        and not found_goal:  
        E = pop_stack()  
        if E == G:  
            found_goal = True  
        else :  
            sort E's children—by shortest  
                estimated remaining distance  
            push these sorted E's children onto stack  
    return found_goal
```


Searching in graphs—hill climbing examples

Examples of hill-climbing problems

- Get to the top of a mountain in a fog where the visibility is two meters.
- Set the temperature controlled by a thermostat without markings.
- Improve the colour, tint, brightness, tuning of TV.


Put abstractly, this search maximizes or minimizes a function.



Hill climbing is not always easy

Hill-climbing algorithms encounter several obstacles

- The *foothill problem*—several local optima may be encountered before the global optimum is reached.




- The *plateau problem*—in a large plateau with isolated peaks it is not only hard to find a peak using hill climbing, but it is difficult to discover the highest peak, or lowest well.
- The *ridge problem*—when the focus lands on a sharp ridge there is only one direction which leads upwards, all the others guide the search in the wrong direction.

Hill climbing is not always easy

Hill-climbing algorithms encounter several obstacles

- The *foothill problem*
- The *plateau problem*—in a large plateau with isolated peaks it is not only hard to find a peak using hill climbing, but it is difficult to discover the highest peak, or lowest well.




- The *ridge problem*—when the focus lands on a sharp ridge there is only one direction which leads upwards, all the others guide the search in the wrong direction.

Hill climbing is not always easy

Hill-climbing algorithms encounter several obstacles



- The *foothill problem*
- The *plateau problem*
- The *ridge problem*. Consider the contour map below. When the focus lands on a sharp ridge, shown as a dotted line, there is only one correct direction—the **green** arrow—which leads upwards, *all* the others guide the search in the wrong direction.



Beam search

Beam search expands several paths and pursues only the best few

- The search progresses level-by-level like breadth-first search.
- Beam search of width w



Finding many paths to the goal

The idea is to find the best path to the goal by with guiding heuristics

- *British Museum* algorithm—finds all possible paths by
 - depth-first search or
 - breath-first

to search inefficiently until every solution is found

- branch-and-bound search—efficient
- discrete dynamic programming search—efficient
- A^* search—efficient

Branch-and-Bound Search

Branch-and-bound search successively *expands* the least-cost partial path eliminating unnecessary work.

Branch-and-bound search

- tracks of all contending partial paths
- the shortest path is extended first
- terminates when no path can be expanded to the goal that yields a shorter or more optimal complete path path.

Branch-and-Bound Search

- To conduct a branch-and-bound search, form a one-element queue consisting of a zero-length path that contains only the root node.
- Until the first path in the queue terminates at the goal node or the queue is empty,
- Remove the first path from the queue; create new paths by extending the first path to all the neighbours of the terminal node.
- Reject all new paths with loops.
- Add the remaining new paths, if any, to the queue.
- Sort the entire queue by path length with least-cost paths in front.
- If the goal node is found, announce success; otherwise, announce failure.

Adding Underestimates

- Guess distance remaining in general:

$$e(\text{total path length}) = d(\text{already traveled}) + e(\text{distance remaining})$$

Known estimate

- Underestimates:

$$u(\text{total path length}) = u(\text{already traveled}) + u(\text{distance remaining})$$

Known underestimate

For road map—straight-line distance is guaranteed to be an underestimate—this is the crux of the approach

- underestimate as close as possible to the true distance
- underestimate of close to zero is of little value

Branch-and-bound search with a lower-bound estimate

- To conduct a branch-and-bound search with a lower-bound estimate, form a one-element queue consisting of a zero-length path that contains only the root node.
- Until the first path in the queue terminates at the goal node or the queue is empty,
- Remove the first path from the queue; create new paths by extending the first path to all the neighbours of the terminal node.
- Reject all new paths with loops.
- Add the remaining new paths, if any, to the queue.
- Sort the entire queue by the sum of the path length and a lower-bound estimate of the cost remaining, with least-cost paths in front.
- If goal node is found, announce success; otherwise, announce failure.

The dynamic-programming principle

- objective: eliminate redundant partial paths the Dynamic-Programming Principle

The best way through a particular, intermediate place is the best way to it from the starting place,

+ followed by the best way from it to the goal. There is no need to look at any other paths to or from the intermediate place.

Branch-and-Bound Search and Dynamic-Programming

- To conduct a branch-and-bound search with dynamic programming, form a one-element queue consisting of a zero-length path that contains only the root node.
- Until the first path in the queue terminates at the goal node or the queue is empty,
- Remove the first path from the queue; create new paths by extending the first path to all the neighbours of the terminal node.
- Reject all new paths with loops.
- Add the remaining new paths, if any, to the queue. If two or more paths reach a common node, delete all those paths except the one that reaches the common node with the minimum cost.
- Sort the entire queue by path length with least cost paths in front.
- If goal node is found, announce success; otherwise, announce failure.

A* Search

- To conduct A* search, form a one-element queue consisting of a zero-length path that contains only the root node.
- Until the first path in the queue terminates at the goal node or the queue is empty,
 - Remove the first path from the queue; create new paths by extending the first path to all the neighbours of the terminal node.
 - Reject all new paths with loops. If two or more paths reach a common node, delete all those paths except the one that reaches the common node with the minimum cost.
 - Sort the entire queue by the sum of the path length and a lower-bound estimate of the cost remaining, with least-cost paths in front.
- If goal node is found, announce *success*; otherwise, announce *failure*.

Several Search Procedures Find the Optimal Path

- The British Museum procedure is good only when the search tree is small.
- Branch-and-bound search is good when the tree is big and bad paths turn distinctly bad quickly.
- Branch-and-bound search with a guess is good when there is a good lower-bound estimate of the distance remaining to the goal.
- Dynamic programming is good when many paths converge on the same place.
- The A^* procedure is good when both branch-and-bound search with a guess and dynamic programming are good.

Application: obstacle avoidance problem for robots

- General idea (compare Number Scrabble and Tic-Tac-Toe):
- redescribe the problem in a simpler representation solve the simpler problem
- redescribe the solution in the original representation
- find a new presentation: configuration-space obstacles
- turns: object-obstacle problems — $\rightarrow$ point-obstacle problems
- allows to use A^* search techniques

Find “One” Paths $\iff$ Find the “Shortest” Path

- why do we care?
- in which situations do we care?

Examples:

- main streets versus side streets (riding a bicycle to work, using a high functionality application)
- buy a can of (low-altitude) tennis balls
- find someone to repair a car
- search in the WWW (40 000 items returned) $\Rightarrow$ “optimal” solutions?

Summary

- The British Museum procedure is one of many search procedures oriented toward finding the shortest path between two points. The British Museum procedure relies on working out all possible paths.
- Branch-and-bound search usually saves a lot of time relative to the British Museum procedure. It works by extending the least-cost partial path until that path reaches the goal.
- Adding underestimates to branch-and-bound search improves efficiency. Deleting redundant partial paths, a form of dynamic programming, also improves efficiency. Adding underestimates and deleting redundant partial paths converts branch-and-bound search into A^* search.
- The configuration-space transformation turns object-obstacle problems into point-obstacle problems. So transformed, robot path-planning problems succumb to A^* search.

Induction

- We use *induction* to compose a procedure by assuming that we can do the program for $n - 1$ items.
- Then we show how to do the code for n items. We also need to know how to deal with one item.
- For example when coding the Hanoi towers problem, assume that you CAN do `hanoi(n-1, from, help, to)` and proceed to show how to program the code when there is one more item, giving n items altogether.
- So (1) use `hanoi(n-1, from, to, help)` to clear everything off the bottom, n th disk, then (2) move the bottom or n th disk, to the `to` tower, and, finally, (3) move the pile on the `help` tower to the `to` tower with `hanoi(n-1, from, help, to)`.
- We have derived the program AND proved that it is correct.

Induction—The towers of Hanoi

Using induction when coding the Hanoi towers problem, assume that you CAN do `hanoi(n-1, from, help, to)` and proceed to show how to program the code when there is one more item, giving n items altogether.

So (1) use `hanoi(n-1, from, to, help)` to clear the bottom, n th disk, then (2) move the bottom or n th disk, to the `to` tower, and, finally, (3) move the pile on the `help` tower to the `to` tower with `hanoi(n-1, from, help, to)`.

```
1 def hanoi(n, f, t, h):
2     if n == 1:
3         print ("Move_from_%s--->_%s---wait_at_%s" % (f, t, h))
4     else:
5         hanoi(n-1, f, h, t)
6         print ("Move_from_%s--->_%s---wait_at_%s" % (f, t, h))
7         hanoi(n-1, h, t, f)
```

The program has been derived AND proved correct by employing induction.

Algebraic and numerical algorithms

- Evaluating a polynomial at a point
- Horner's rule
- Exponentiation
- Greatest common divisor
- Polynomial multiplication

Evaluating a polynomial at a point

Given a sequence of numbers $a_n, a_{n-1}, \dots, a_2, a_1, a_0$ and a real number x calculate the value of the polynomial

$$P_n(x) = a_n x^n + a_{n-1} x^{n-1} + \dots + a_2 x^2 + a_1 x + a_0$$

- Note that the problem involves $n + 2$ numbers.
- Let us try to formulate an inductive way to code the algorithm

Inductive hypothesis

p_0 : Basis: $P_0(x) = a_0$.

p_k : Assume that we know how to calculate $P_k(x)$. We must show how to derive the calculation for p_{k+1} using $P_k(x)$ as a start.

$$P_{k+1}(x) = a_{k+1} x^{k+1} + P_k(x)$$

Evaluating a polynomial at a point

We will rewrite the induction hypothesis slightly to make the coding more transparent.

Inductive hypothesis

p_0 : The basis of the induction is $P_0(x) = a_0$.

p_{n-1} : Assume that we know how to calculate $P_{n-1}(x)$. We must show how to derive the calculation for p_n using $P_{n-1}(x)$ as a start.

$$P_n(x) = P_{n-1}(x) + a_n x^n$$

This leads to our program in a very natural way

```
1 def P(n, a, x):
2     if n == 0:
3         return a[0]
4     else:
5         return P(n-1, a, x) + a[n]*x**n
```

This algorithm requires n additions, multiplications and exponentiations.

Evaluating a polynomial at a point

Using the same

Inductive hypothesis

p_0 : The basis of the induction is $P_0(x) = a_0$.

p_{n-1} : Assume that we know how to calculate $P_{n-1}(x)$. We must show how to derive the calculation for p_n using $P_{n-1}(x)$ as a start.

$$P_n(x) = P_{n-1}(x) + a_n x^n$$

we now code the program using iterative code

```
1 def P(n, a, x):  
2     polynomial = a[0]  
3     for k in range(1, n+1):  
4         polynomial = polynomial + a[k]*x**k  
5 return polynomial
```

The iterative algorithm requires n additions, multiplications and exponentiations—the same as that of the recursive algorithm.

Evaluating the powers of x

Now we alter our induction hypothesis slightly to get rid of some of the redundant calculations of powers of x . If we were to use the value of x^{n-1} in order to calculate x^n we could save $\frac{(n-1)n}{2}$ multiplications needed for calculating the exponents of x .

Inductive hypothesis for the exponents

p_0 : The basis of the induction is $Xpower(0) = 1$.

p_{n-1} : Assume that we know how to calculate $Xpower(n-1)$.

$$Xpower(n) = x \times Xpower(n-1).$$

This does not make it quite clear how to organize calculating $P_n(x)$.

A possible approach would be to calculate all the exponents of x in one go and store them in an array.

But there is a better way.

Rewrite the induction hypothesis.

Horner's rule for evaluating a polynomial at a point

Rewrite the induction hypothesis using the fact that we *know* how to calculate the exponents of x . We will start at coefficient a_n , and extend our polynomial to include a_{n-1} , a_{n-2} , until we reach a_0 .

Inductive hypothesis

p_0 : The basis of the induction is $P'_0(x) = a_n$ or $P'_1(x) = a_n x + a_{n-1}$.

p_{n-1} : Assume that we can calculate $P'_{n-1}(x)$ where

$$P'_{n-1}(x) = a_n x^n + a_{n-1} x^{n-1} + \cdots + a_2 x^2 + a_1 x.$$

The calculation for p_n using $P'_{n-1}(x)$ as a start is

$$P'_n(x) = x \cdot P'_{n-1}(x) + a_0$$

$P'_n(x) \equiv P_n(x)$. This leads to our program in a very natural way

```
1 def P(n, a, x):  
2     polynomial = a[n]  
3     for k in range(n-1, -1, -1):  
4         polynomial = x*polynomial + a[k]  
5     return polynomial
```

This algorithm requires n additions, and n multiplications, whereas our naive algorithm took an extra $\frac{(n-1)n}{2}$ multiplications for the x powers.

Exponentiation— x^n

It seems trivial to compute x^n using the inductive idea that we can assume p_0 : that we know how to calculate x^0 —it is always 1, and assume p_{n-1} : that we know how to calculate x^{n-1} and can then easily show how to perform p_n : $x^n = x \cdot x^{n-1}$.

```
1 def power(x, n):
2     if n == 0:
3         return 1
4     else:
5         return x * power(x, n-1)
```

The non-recursive code derived from the induction argument

```
1 def power(x, n):
2     pow = 1
3     for k in range(n):
4         pow = x*pow
5     return pow
```

In both programs there are n multiplications—this can be improved.

Exponentiation— x^n —faster

Consider

$$\begin{array}{ll} 2^{256} = 2^{128} \times 2^{128}, & 2^{128} = 2^{64} \times 2^{64}, \\ 2^{64} = 2^{32} \times 2^{32}, & 2^{32} = 2^{16} \times 2^{16}, \\ 2^{16} = 2^8 \times 2^8, & 2^8 = 2^4 \times 2^4, \\ 2^4 = 2^2 \times 2^2, & 2^2 = 2 \times 2. \end{array}$$

This takes 8 multiplications, whereas naive exponentiation takes $256 = 2^8$ multiplications, i.e. instead of n multiplications the calculation takes $\log_2 n$ multiplications.

This raises some questions.

- (1) How do we do this for any power of 2?
- (2) What is going on? Can the base be changed to x ?
- (3) How will we use this idea to speed up x^n ?

We will now address these questions.

Fast exponentiation—analysis of 2^{256}

Rewrite the term 2^{256} as $2^{100000000_2}$,

$$\begin{aligned} 2^{100000000_2} &= (2^{10000000_2})^2, \\ &= ((2^{1000000_2})^2)^2, \\ &= (((2^{100000_2})^2)^2)^2, \\ &= ((((2^{10000_2})^2)^2)^2)^2, \\ &= ((((((2^{1000_2})^2)^2)^2)^2)^2)^2, \\ &= (((((((2^{100_2})^2)^2)^2)^2)^2)^2)^2, \\ &= (((((((((2^{10_2})^2)^2)^2)^2)^2)^2)^2)^2, \\ &= ((((((((((2^{1_2})^2)^2)^2)^2)^2)^2)^2)^2)^2)^2. \end{aligned}$$

The last line says everything. It describes the meaning of $2^{100000000_2}$ exactly: first, there is a 2^1 —represented by the 1, and this is followed by 8 squarings.

We next consider $2^{19} = 2^{10011_2}$.

Fast exponentiation—analysis of 2^{19} and x^{19}

$$\begin{aligned} 2^{19} &= 2^{16+2+1} \\ &= 2^{10000_2+10_2+1_2} = 2^{10011_2} \\ &= 2^{10000_2} \times 2^{10_2} \times 2^{1_2} \\ &= (((2^2)^2)^2)^2 \times 2^2 \times 2 \end{aligned}$$

What if we alter the base 2 to x ?

$$\begin{aligned} x^{19} &= x^{16+2+1} \\ &= x^{10000_2+10_2+1_2} = x^{10011_2} \\ &= x^{10000_2} \times x^{10_2} \times x^{1_2} \\ &= (((x^2)^2)^2)^2 \times x^2 \times x \end{aligned}$$

1	0	0	1	1
4	3	2	1	0

 nineteen

The bit number usually corresponds to the power of 2 that it represents.

Next, we consider an algorithm to calculate x^n using these insights.

Fast exponentiation—calculating x^{19}

$n = 19 = 10011_2$. These bits are used one-by-one to calculate x^n .

$$\begin{aligned}x^{19} &= x^{10011_2} = x^{10000_2 + 10_2 + 1_2} = x^{10000_2} \times x^{10_2} \times x^{1_2} \\ &= (((x^2)^2)^2)^2 \times x^2 \times x\end{aligned}$$

Bit₄ = 1 contributes $1 \times x^{16} = (((x^2)^2)^2)^2$ to the result,
bit₃ = 0 contributes $0 \times x^8$, i.e. nothing, to the result,
bit₂ = 0 contributes $0 \times x^4$, i.e. nothing, to the result,
bit₁ = 1 contributes $1 \times x^2 = x^2$ to the result, and
bit₀ = 1 contributes $1 \times x^1 = x$ to the result.

Our calculation proceeds from the right as follows

Initialize the variables.	result = 1;	power = x
bit ₀ = $1 \times x^1 = x$	result *= power;	power *= power
bit ₁ = $1 \times x^2 = x^2$	result *= power;	power *= power
bit ₂ = $0 \times x^4$, i.e. nothing		power *= power
bit ₃ = $0 \times x^8$, i.e. nothing		power *= power
Bit ₄ = $1 \times x^{16} = (((x^2)^2)^2)^2$	result *= power;	power *= power

Fast exponentiation—calculating x^{39}

$$\begin{aligned}x^{39} &= x^{100111_2} = x^{100000_2+100_2+10_2+1_2} = x^{10000_2} \times x^{100_2} \times x^{10_2} \times x^{1_2} \\&= (((((x^2)^2)^2)^2)^2 \times (x^2)^2 \times x^2 \times x\end{aligned}$$

bit₅ = 1 contributes $1 \times x^{32} = (((((x^2)^2)^2)^2)^2$ to the result,
bit₄ = 0 contributes $0 \times x^{16}$, i.e. nothing, to the result,
bit₃ = 0 contributes $0 \times x^8$, i.e. nothing, to the result,
bit₂ = 1 contributes $1 \times x^4 = (x^2)^2$ to the result, and
bit₁ = 1 contributes $1 \times x^2 = x^2$ to the result.
bit₀ = 1 contributes $1 \times x^1 = x$ to the result.

Our calculation proceeds from the right as follows

Initialize the variables.

bit ₀ = 1 × $x^1 = x$	result = 1;	power = x
bit ₁ = 1 × $x^2 = x^2$	result *= power;	power *= power
bit ₂ = 1 × $x^4 = (x^2)^2$	result *= power;	power *= power
bit ₃ = 0 × x^8 , i.e. nothing		power *= power
bit ₄ = 0 × x^{16} , i.e. nothing		power *= power
bit ₅ = 1 × $x^{32} = (((((x^2)^2)^2)^2)^2$	result *= power;	power *= power

Fast exponentiation—algorithm for x^n

Initialize the multiplicand called **result**.

In each step: get the corresponding bit of **n**,
if it is **1**, multiply by the current **power**.

In every step double the **power**, and halve **n**.

Stop when **n** becomes zero.

```
1 def exponent (x, n):
2     result = 1
3     power = x
4     while n > 0:
5         if n % 2 == 1:
6             result *= power
7             power *= power
8             n //= 2
9     return result
```

The number of multiplications in the expansion

$x^n = x \times x \times \cdots \times x$, n times, is $O(n)$.

This algorithm requires only $O(\log n)$ multiplications.

Cryptography—RSA public-key encryption

Rivest, Shamir and Adleman—in 1978—developed their RSA public-key encryption scheme.

Assume that one needs to communicate securely with a participant **P** who has two keys— e for encryption and d for decryption.

Start with the product of two primes p and q , $n = pq$ and get a number d that does not divide the *totient* $\phi(n) = (p - 1)(q - 1)$

$$\begin{aligned} n &= p \cdot q \\ d &\nmid \phi(n) = (p - 1)(q - 1) \\ e \cdot d &\equiv 1 \pmod{\phi(n)} \\ E_P(M) &= M^e \pmod{n} = C \\ D_P(C) &= C^d \pmod{n} = M \end{aligned} \tag{5}$$

One can prove that Equation 5 guarantees that

$$\begin{aligned} D_P(E_P(M)) &= D_P(C) = M \\ \text{i.e. } ((M^e) \pmod{n})^d \pmod{n} &= M \end{aligned}$$

RSA public-key encryption

- Suppose p and q are two very large prime numbers with $n = p \cdot q$
- The participant $\mathbf{P}$ selects another very large integer d such that it has no common divisor with $\phi(n) = (p - 1)(q - 1)$, i.e. $d \nmid \phi(n)$
- Next a number e is calculated such that $e \cdot d \equiv 1 \pmod{\phi(n)}$
- e will be the *encryption* key and d will be the *decryption* key
- The values of e and n are made public—hence “public-key encryption”—by $\mathbf{P}$
- The values of d , p and q are kept secret by $\mathbf{P}$ so that others can encrypt messages that only $\mathbf{P}$ can decrypt easily.

RSA public-key encryption

- Both the encryption and decryption are done by repeated exponentiation by either e or d
- Anybody that knows the public keys can encrypt a message M into a coded secret message C to send to $\mathbf{P}$ using
$$E_P(M) = M^e \bmod n = C$$
- Since both n and e are public, everyone can encrypt a message M and send the coded message C to $\mathbf{P}$
- Since $\mathbf{P}$ is the only one who knows d , only $\mathbf{P}$ can decrypt the coded message C back into the plain text M using
$$D_P(C) = C^d \bmod n = M$$
- RSA public-key encryption has made *secure communication* possible

RSA public-key encryption in review

$$\begin{aligned}n &= p \cdot q \\d &\nmid (p-1)(q-1) \\e \cdot d &\equiv 1 \bmod (p-1)(q-1) \\E_P(M) &= M^e \bmod n = C \\D_P(C) &= C^d \bmod n = M\end{aligned}\tag{6}$$

One can prove that Equation 6 guarantees that

$$D_P(E_P(M)) = D_P(C) = M$$

RSA encryption is very safe because knowing only n and e , it is extremely tedious to calculate any one of d , p or q since no efficient algorithm exists for factorizing the product of very large primes.

RSA is very efficient since both encryption and decryption can be done in $O(\log n)$ multiplications and divisions of numbers less than n .

Fast exponentiation—algorithm for $x^e \bmod n$

Since $p \cdot q \bmod n = (p \bmod n \cdot q \bmod n) \bmod n$

the modulo n operation can be applied at every step of the exponentiation

This keeps the size of the products bounded to n in any step

```
1 def powermod (x, e, n):
2     result = 1
3     power = x % n
4     while e > 0:
5         if e % 2 == 1:
6             result = ((result % n) * (power % n)) % n
7             power = (power * power) % n
8             e //= 2
9     return result
```

The RSA encryption algorithm

```
1 def rsa(x, e, n):
2     global primes
3     l = 2500;
4     p = primes[l]; q = primes[l+17]
5     # p = 22307; q = 22511
6     n = p*q; phin = (p-1)*(q-1) # Euler's phi(n)
7
8     # Find any odd integer d such that
9     # d does not divide perfectly into (p-1)*(q-1)
10
11     d =
12     e = n // d
13
14     # Check that e*n == 1 mod phi(n)
15     . . .
```

The details have been left to the reader.

Greatest common divisor—Euclid's algorithm

If $k = \text{GCD}(n, m)$, then k divides both n and m , and all other numbers that divide n and m are smaller than k ;

k is the *greatest common divisor* of n and m , put otherwise

$$\text{GCD}(n, m) = \max\{k : k \mid n \text{ and } k \mid m\}$$

Euclid noticed that if k divides both n , and m , i.e. $k \mid n$, and $k \mid m$,

then $n = p \cdot k$, and $m = q \cdot k$,

thus $n - m = (p - q) \cdot k$ and $m - n = (q - p) \cdot k$,

so $k \mid n - m$, and $k \mid m - n$.

Finally, doing repeated subtractions, i.e. division, $k \mid n \bmod m$.

```
1 def gcd(n,m):  
2     while m != 0:  
3         m, n = n % m, m  
4     return n
```


Polynomial multiplication— $O(n^2)$

Let $P(x) = p_{n-1}x^{n-1} + \dots + p_1x + p_0$

and $Q(x) = q_{n-1}x^{n-1} + \dots + q_1x + q_0$,

Then $PQ = (p_{n-1}x^{n-1} + \dots + p_1 + p_0)(q_{n-1}x^{n-1} + \dots + q_1 + q_0) =$

$p_{n-1}q_{n-1}x^{2n-2} + \dots + (p_{n-1}q_{i+1} + p_{n-2}q_{i+2} + \dots + p_{i+1}q_{n-1})x^{n+i} + \dots + p_0q_0.$

The number of multiplications and additions is clearly $O(n^2)$.

This can be reduced to $O(n \log n)$ the fast Fourier transform.

We first try using a divide-and-conquer algorithm.

Polynomial multiplication—divide-and-conquer

Let $P(x) = p_{n-1}x^{n-1} + \cdots + p_1x + p_0$ and $Q(x) = q_{n-1}x^{n-1} + \cdots + q_1x + q_0$,

Assuming that n is a power of 2, split each of P and Q into equal parts,

putting $P = P_1 + x^{n/2}P_2$ and $Q = Q_1 + x^{n/2}Q_2$, where

$P_1 = p_0 + p_1x + \cdots + p_{n/2-1}x^{n/2-1}$, $P_2 = p_{n/2} + p_{n/2+1}x + \cdots + p_{n-1}x^{n/2-1}$,

and

$Q_1 = q_0 + q_1x + \cdots + q_{n/2-1}x^{n/2-1}$, $Q_2 = q_{n/2} + q_{n/2+1}x + \cdots + q_{n-1}x^{n/2-1}$.

Thus $PQ = (P_1 + x^{n/2}P_2)(Q_1 + x^{n/2}Q_2) =$

$P_1Q_1 + (P_1Q_2 + P_2Q_1)x^{n/2} + P_2Q_2x^n$.

pause Each of these products P_iQ_j can be further halved inductively until the solution is reached.

Since there are 4 products involved at each step, the total number of operations $T(n)$ required is given by

$T(n) = 4T(n/2) + O(n)$, $T(1) = 1$.

The solution of this recurrence relation is $O(n^2)$ —so this method has yielded no improvement.

Polynomial multiplication—divide-and-conquer

In order to improve on $O(n^2)$ we must solve the problem using less multiplications. Consider the multiplication table

$\times$	P_1	P_2
Q_1	A	C
Q_2	B	D

$$\begin{aligned} E &= (P_1 + P_2)(Q_1 + Q_2) \\ &= P_1Q_1 + P_1Q_2 + P_2Q_1 + P_2Q_2 \\ &= A + (B + C) + D \end{aligned}$$

In order to compute $A + (B + C)x^{n/2} + Dx^n$ we observe that $B + C$ can be derived from the product $E = (P_1 + P_2)(Q_1 + Q_2)$, i.e. $B + C = E - A - D$. Hence we need to calculate only three products of smaller polynomials. Adding $P_1 + P_2$, and $Q_1 + Q_2$ is linear in the number of their coefficients.

The recurrence relation for this algorithm is

$$T(n) = 3T(n/2) + O(n)$$

and its solution is $T(n) = O(n^{\log_2 3}) = O(n^{1.59})$ which is significantly faster than $O(n^2)$.

- What is science?
- How does it come about?
- Science is the unending pursuit of further but hopefully more accurate patiently accumulated knowledge and improved approximations of the intangible *Wirklichkeit*—reality.
- Computer scientists are more interested in representing, measuring and computing reality.
- Euclid based his geometry on a few simple axioms.
- Gottfried von Leibniz dreamed of a calculus of reason where problems are resolved with pen and paper. He said, “*Calculemus*”—Let us calculate.
- Hilbert profoundly influenced the axiomatization of mathematics and reinvented plane geometry with four axioms and the parallel postulate as a fifth axiom.

- What is computation?
- What can be computed? David Hilbert (1862–1943) said:
*Wir müssen wissen,
Wir werden wissen.*
- In Paris in 1900 Hilbert said, speaking about proof in mathematics:

“This satisfaction of the solution of every mathematical problem is a powerful motivation during our labours; we hear the constant call:

Here is the problem, find the solution.

One can find it by clear thinking; since in mathematics there is nothing that cannot be known.”

- Hilbert was reflecting what was widely believed by logicians and mathematicians at the beginning of the 20th Century.
- Was he right?

Decision problems—*Entscheidungsproblem*

- Determining the truth of statements in formal systems is known as the *Entscheidungsproblem* or the decision problem
- Kurt Gödel (1906–1978) *proved* it to be an impossible dream.
- For a system as simple as arithmetic over the integers he showed that it is impossible to decide every statement.
- In fact he proved that there are TRUE statements concerning arithmetic that cannot be proved correct in an axiomatic system.
- He was a PhD student when he discovered this.
- The result must have shocked his mentors—it stirred the foundations of mathematics.
- The mighty Hilbert was PROVED to be wrong.

- In 1930 it rocked the foundations of Hilbert's view of formal logic at a time when mathematics appeared to be the unalterable bastion of science while *quantum physics* was altering our understanding of the *Wirklichkeit*.
- Hilbert's dream of mechanizing maths the way Hilbert wanted to do it was smashed by Gödel
- How was this idea received?
- John Dawson states, “One of the most profound discoveries in the history of mathematics was assimilated promptly and without objection by Gödel's contemporaries.”
- The prevailing Hilbertian view was subsumed by a new view of science by Gödel where statements may be valid or invalid or might not be provable at all.

- Hilbert never uttered a word of criticism. How could he argue against Gödel proof?
- Only one of Gödel's peers made a severe criticism which he withdrew in writing a week later.
- Knowing what is impossible tells us a lot about what computers can and cannot do
- Gödel's showed that no algorithm can be written to read a statement about integer arithmetic and tell us whether that statement is TRUE or FALSE.

Computation—Kurt Gödel and others

- Gödel's showed that no algorithm can be written to read a statement about integer arithmetic and tell us whether that statement is TRUE or FALSE.
- Alonso Church (1903–1995) proved that in λ -calculus problems $\exists$ which had no algorithmic solution.
- Steven Kleene (1909–1994) used recursion theory.
- Emile Post (1897–1954) did it using productions.
- Alan Turing (1912–1954) designed a “machine” and proved that this machine could compute everything that was effectively computable.
- All these discoveries lead to the idea that these methods are equivalent in some sense.

Computation—Church-Turing thesis

- We can take a Turing machine and emulate Post's productions or program a Turing machine to emulate itself or Kleene's recursive system.
- We can emulate a Turing machine by using a C or Java program.
- These two statements known as the *Church-Turing thesis* are commonly believed
 1. All reasonable definitions of 'algorithm' which are so far known are equivalent.
 2. Any reasonable definition of 'algorithm' which anyone will ever make will turn out to be equivalent to the definitions that we know.
- The *Church-Turing thesis* merely states that we think we have a good intuitive grasp of the concept of devising instructions to perform some task.

Computation—Alan Turing

- Computers that interact with the world can already do in varying degrees many tasks that humans once thought made them unique.
- Alan Turing showed in 1936 that programmable machines could do anything that is effectively computable
- He also showed that no machine can be built to decide whether every given program would give a result and halt—we will sketch the proof of this later. This proves that there are uncomputable things
- This is disappointing, but it is a relief that this can at least be proved.
- There are still very many computable algorithms
- John von Neumann (1903–1957) *designed* a practical machine not unlike today's modern computers that could be built and used to execute algorithms.

Limitations of Computation

- Computers took us to the moon
- Science is virtually impossible without computers Will computers be able to think?
- Marvin Minsky: “There is no reason to suppose machines have any limitations not shared by man.”
- Computers are better at very many 'intellectual' activities
- That computers are not so hot on philosophy perhaps says more about philosophy than about the potential of computers.
- Computers can still anything that a mathematician can do. We have not yet implemented a mathematician.
- There is nothing inherently impossible about this.
- Computers have the advantage of immortality. Programs written for computers are persistent.

Limitations of Computation

- Douglas Lenat's EURISKO and CYC solve problems too complex for humans.
- Knowing what is impossible tells us a lot about what computers can and cannot do
- Alan Turing showed that no machine can be built to decide whether every given program would give a result and halt.
- This proves that there are uncomputable things
- James Watson about Craig Venter's genome sequencing machines: "It isn't science the machines could be run by monkeys."
- We will now discuss some uncomputable algorithms.

Turing's halting problem

- When a program does not halt, it is a ubiquitous problem for programmers to know if their program is in a closed loop or perhaps that the computation is very slow.
- So it would be useful to have another program that decides if the program actually will halt or loop interminably.
- This is called the *halting problem*.
- Our algorithm is a program that is given an arbitrary program **P** and the input data **D** for **P**.
- Let us call our program **halttester** which takes two arguments.
- So the call **halttester(P, D)** will yield '**OK**' if the program **P** halts on being run with data **D** or yields '**BAD**' if it does not.

Turing's halting problem

- `halttester(P, D)` can take any `P` or `D`.
- Construct a new program called `newHalttester(P)` with only one argument as follows:

```
1 value newHalttester(P) {  
2   // Checks if program P halts with input P  
3   return halttester(P,P); }
```

- If we assume that `halttester(P,D)` exists and thus `newHalttester(P)` can be programmed, then we can construct a new program `funny(P)` as follows:

```
1 value funny(P) {  
2   // Assume haltester exists  
3   if (newHalttester(P) == "BAD")  
4     halt;  
5   else  
6     while(1==1); }
```

Turing's halting problem



```
1 value funny(P) {  
2   // Assume haltester exists  
3   if (newHaltester(P) == "BAD")  
4     halt;  
5   else  
6     while(1==1); }
```

- What is the action of `funny(P)`?
- If `P` halts then `funny(P)` loops and if `P` loops then `funny(P)` halts.
- This is clearly impossible. `funny(P)` can neither halt or loop forever.
- Our assumption that `haltester(P, D)` exits is clearly incorrect because it leads to a contradiction.

Turing's halting problem: summary of proof

The proof of the impossibility of a halt tester is summarized below.

1. Assume that we can construct a program called `halttester`.
2. Use it to write another program named `funny` using `newHalttester`.
3. Show that `funny` has some impossible property.
4. Conclude that the assumption must be wrong.

Turing's halting problem: summary of proof

The proof of the impossibility of a halt tester is summarized below.

1. Assume that we can construct a program called `halttester`.
2. Use it to write another program named `funny` using `newHalttester`.
3. Show that `funny` has some impossible property.
4. Conclude that the assumption must be wrong.

Other noncomputable problems

- *The totality problem* A program P that halts on all inputs D is called *total*.
- *The equivalence problem* The idea of a program P that can compare two programs to see if they behave in exactly the same way is called the *equivalence problem*.
- *Rice's Theorem* There is no program to determine if another program actually achieves what it is intended to do.
- The idea of *partial computability*.
- On the other hand there is the *recursion theorem*—there is always a program that can print itself.

The totality problem

A program $P(D)$ that halts on *all* inputs D is called *total*.

Whether or not a program $P(D)$ halts on all inputs D is called the *totality problem*—the *totality problem* is noncomputable.

First write the program `funnypd(I)`.

```
1 value funnypd(I) { // input I is ignored
2   simulate P(D); }
```

Assume that the *totality problem* IS computable; code `halttester(P)`.

```
1 value halttester(P) {
2   // Assume isTotal can be programmed
3   if (isTotal(funnypd(I)))
4     return "OK";
5   else
6     return "BAD"; }
```

Asking if `funnypd(I)` is total is the same as asking whether $P(D)$ halts because we may code `halttester` using it. Since this is impossible our assumption is FALSE.

The equivalence problem 1/2

The idea of a program **P** that can compare two programs to see if they behave in exactly the same way is called the *equivalence problem*—it is not computable. Construct **funnyp(D)** and **simple(D)**:

```
1 value funnyp(D) { // outputs '13' if P(D) halts
2   simulate P(D);
3   display(13); }
4
5 value simple(D) { // outputs '13' regardless
6   display(13); }
```

Asking if **funnyp(D)** is equivalent to **simple(D)** amounts to asking if **P** is total.

If **P** halts on every input, then **funnyp(D)** will always output **13**, whereas if there is some input for which **P(D)** does not halt, then **funnyp(D)** will not halt and neither will it print out **13**.

The equivalence problem 2/2

Asking if `funnyp(D)` is equivalent to `simple(D)` amounts to asking if `P` is total.

If `P` halts on every input, then `funnyp(D)` will always output `13`, whereas if there is some input for which `P(D)` does not halt, then `funnyp(D)` will not halt and neither will it print out `13`. An algorithm for the equivalence problem will also give us an algorithm for the totality problem:

```
1 value total(P) {  
2   // Assume isEquivalent can be programmed  
3   if (isEquivalent(funnyp, simple, D))  
4     return "TOTAL";  
5   else  
6     return "notTOTAL"; }
```

But since it is impossible to program `total(P)` the assumption that we can program `isEquivalent` is FALSE.

Partial computability

The halting problem is called *partially computable* because when it halts we can tell that it has halted,

i.e. there is an algorithm that can say that it has halted when it halts but cannot say when it is in a loop.

The *totality* and *equivalence* problems are NOT partially computable.

This distinction relates to *proof* systems.

It is possible to show that problems are partially computable if and only if they have a proof system.

Gödel showed that arithmetic is *not* even partially computable.

Thus in any arithmetic proof system there are statements that are TRUE even though they cannot be proved.

Conclusion

On a philosophical note—

- Will computers ever be able to think?
- Explaining and automating how science discovers and accumulates new ideas and adds these truths and untruths falteringly to our increasing and persistent knowledge of the world is one road of many possible means to automate thinking.

Complexity

- Some problems are theoretically impossible to program. Other problems take interminably long.
- An $O(2^n)$ problem takes too long to run to completion using a super fast nanosecond computer that runs $O(1)$ problems in one nanosecond when n is as small as 100

$$\begin{aligned}2^{100}\text{ns} &= (2^{10})^{10}\text{ns} = 1024^{10}\text{ns} \\&\approx 1000^{10}\text{ns} \\&= 10^{30}\text{ns} = 10^{30-9}\text{seconds} = \frac{10^{21}}{3600}\text{hours} = \frac{10^{21}}{86400}\text{days} \\&= \frac{10^{21}}{31557600}\text{years} = \frac{10^{21-8}}{0.31557600}\text{years} \\&\approx 3 \times 10^{13}\text{years} \\2^{100}\text{ns} &= 3.16681 \times 10^{13}\text{years}\end{aligned}$$

- A *Julian year* has exactly 31557600 seconds or 365.25 days.

The fastest transistors

- A picosecond computer will take about 3×10^{10} years to do our 2^{100} problem. Feng and Hafez (2005) built a 604Ghz transistor that switches in about 2 picoseconds.
- Using a picosecond computer our problem will still be infeasible, since it will take about 3×10^{10} years. Feng and Hafez built a 604Ghz transistor in 2005. They switch at about 2 picoseconds.
- William Snodgrass, Milton Feng and Haled Hafez have improved the speed of their transistors to 845Ghz operating at $-55^{\circ}C$. At room temperature $25^{\circ}C$ it runs at 765Ghz. At the University of Illinois at Urbana-Champaign (2006, December 13). World's Fastest Transistor Approaches Goal Of Terahertz Device. ScienceDaily. Retrieved October 8, 2009, from <http://www.sciencedaily.com/releases/2006/12/061212091344.htm>

The fastest transistors

- With a frequency of 845 gigahertz, their latest device is approximately 300 gigahertz faster than transistors built by other research groups, and approaches the goal of a terahertz device.
- LOS ANGELES, May 19, 2009—The Defense Advanced Research Projects Agency (DARPA) has awarded Northrop Grumman Corporation (NYSE: NOC) phase 1 of the \$37-million Terahertz (THz) Electronics contract, U.S. Department of Defense and company officials announced today.

Northrop Grumman-developed transistor that holds the Guinness World Record for speed. The transistor provides much higher frequency and bandwidth capabilities. The fastest transistor based on a device that has a maximum frequency of oscillation well in excess of 1,000 GHz

Times for $O(2^n)$ calculations

Suppose a problem has time complexity $O(2^n)$.

Given n how long will a nanosecond machine take to execute such a problem?

n	2^n	Time taken
100	$2^{45} \times 24 \times 60 \times 60 \times 365.25 \times 10^9$	3.2×10^{13} years
95	$2^{40} \times 24 \times 60 \times 60 \times 365.25 \times 10^9$	10^{12} years
85	$2^{30} \times 24 \times 60 \times 60 \times 365.25 \times 10^9$	milliard years
75	$2^{20} \times 24 \times 60 \times 60 \times 365.25 \times 10^9$	million years
65	$2^{10} \times 24 \times 60 \times 60 \times 365.25 \times 10^9$	1000 years
55	$24 \times 60 \times 60 \times 365.25 \times 10^9$	Julian year
51	$24 \times 60 \times 60 \times 365.25/12 \times 10^9$	month
46	$24 \times 60 \times 60 \times 10^9$	day
42	$60 \times 60 \times 10^9$	hour
36	60×10^9	minute
30	10^9	second

Efficient algorithms and tractability

- In a sense we cannot hope to compute these exponential problems for large n .
- *Algorithms*—or programs—that are $O(n^k)$ or $O(P(n))$ where $P(n)$ is a polynomial, are called **efficient**.
- *Problems* that yield these programs are said to be **tractable**.
- This seems to be a weird definition because a program that runs in time n^{10} is not that efficient, but neither is the *linear* problem that runs in time $10^{10}n$.
- Most tractable problems tend have practical solutions, i.e. the size of k seldom exceeds 4. They are seldom worse than quadratic.
- We call the set of problems with polynomially bounded algorithms, **P**.

There are many problems for which no polynomial-time algorithm is known.
The set of problems with polynomial-time bounded algorithms is called **P**.
There are many problems for which no polynomial-time algorithm is known.
The situation turns out to be much worse.
There are two clear problems.

1. Do we have a polynomial-time algorithm for a given problem?
2. Can we prove that no such polynomial-time algorithm exists?

Next we try to characterize these problems.¹⁶

¹⁶Most of the notes on complexity are based on my paraphrasing of Chapter 11 of Udi Manber, Introduction to Algorithms a Creative Approach, 1989, Addison-Wesley. Manber's book is a must read for the serious algorithmist.

NP-complete problems—1

- What problems are known not to be in **P**?
- We suspect that many problems lie outside of **P**.
- How do we deal with problems outside of **P**?
- There is a class of problems called **NP-complete problems** are all equivalent¹⁷ in the sense that: *There exists an efficient algorithm FOR ANY ONE of these NP-complete problems*
 $\iff$ *There exist efficient algorithms FOR ALL NP-complete problems.*
- There is a widely held belief—no such proof is known—that there is no efficient algorithm for any NP-complete problem.
- These problems all have another interesting characteristic.
Finding the solution is hard, but once a solution has been found verifying it can be done in polynomial time.

¹⁷David Harel puts describes the NP-complete problems as follows:

“Each one’s fate is tightly bound to that of all the others. Either *all* NP-complete problems are tractable, or *none* of them are!

NP-complete problems—2

- Even if such an efficient algorithm exists for any NP-complete problem, then it must be rather complicated because it has eluded us ever since 1971 when Stephen Cook formulated the idea.
- Cook showed that the boolean satisfiability problem—called *SAT*—is NP-complete.
It is known as Cook's theorem.¹⁸
- Many thousands of problems have already been categorized as **NP**
- How do we show that a problem is NP-complete?
- The idea of **polynomial-time reduction** is used.

¹⁸It was independently proven by Leonid Levin and is also known as the Cook-Levin theorem.

“Knowing that a problem is *NP-complete* is a compelling reason to stop searching for an efficient algorithm.”¹⁹

¹⁹Page 452, Jon Kleinberg and Eva Tardos, Algorithm Design, Pearson 2005.

Decision problems and reduction

- A **decision problem** is an algorithm is like a predicate. It yields the answer yes or no.
- Most problems can be converted easily to decision problems.
- For example if we want to find the maximum number of matching elements in two separate lists we can write a decider that tests if more than k matches exist.
- If you have a decider for a given problem it is easy to derive an algorithm for the original problem.
- How could you find the actual value of the maximum using such a decider?
- Hint: repeatedly call the decider using $\dots$ to guide the direction of k .


Decision problems and reduction

- Let U be the set of all inputs to a decider. Deciders can be seen as **language recognizers**:
 $L = \{s \mid \text{decide that } s \in L\}.$
- $L \subseteq U$ is the set of all inputs which the decider accepted.
- L is known as the **language** corresponding to the problem.
- *Language* and *problem* are interchangeable terms.
- **Definition:** Let L_1 and L_2 be two languages from the input spaces U_1 and U_2 . L_1 is **polynomially reducible** to L_2 if there exists a *polynomial-time converter* that converts each input $u_1 \in U_1$ to another input $u_2 \in U_2$ such that $u_1 \in L_1$ if and only if $u_2 \in L_2$.
- The algorithm is polynomial in the size of the input u_1 . The size of u_2 is also polynomial in the size of u_1 .

Polynomial reduction

- The definition shows us how to convert one language—problem—into another.
- If we have an algorithm for L_2 , we can compose an algorithm for L_1 by **converting** its input u_1 to an input $u_2 \in U_2$ for L_2 , and simply use the decider for L_2 to see if that converted u_2 is in L_2 . This gives us a decider for L_1 .
- We have now proved that:
If L_1 is polynomially reducible to L_2 and there is a polynomial-time algorithm for L_2 , then there is a polynomial-time algorithm for L_1 .
- Reducibility is not symmetric. If L_1 is polynomially reducible to L_2 then the reverse does not necessarily apply.
- L_2 is regarded as the harder problem.
- If the reducibility works in both directions the languages are called **polynomially equivalent**.

An example of polynomial reduction—1




We are given a graph G with N nodes.

Construct a travelling salesman network T as follows. The nodes of T are precisely the nodes of G but edges are drawn between every two nodes, assigning a weight of 1 to an edge if it was present in G , and 2 if not.

T has a travelling salesman tour of length $N + 1$ or less precisely when G contains a Hamiltonian path.

We need to answer questions about the existence of Hamiltonian paths.

An example of polynomial reduction—2




To answer questions about the existence of Hamiltonian paths, take the input graph G and carry out the transformation to T . Then determine if T has a travelling salesman tour no longer than $N + 1$, where N is the number of nodes in G .

The answer to the first question on G is “yes” whenever the answer to the second question on T is “yes”.

The transformation only takes polynomial time.

An example of polynomial reduction—3



Answering questions about the existence of Hamiltonian paths, do the graph construction and consider the travelling salesman problem.

The transformation only takes polynomial time.

In terms of *tractability* the Hamiltonian path problem is no worse than the travelling salesman problem.

If travelling salesman problem has an efficient solution then so will the Hamiltonian path problem.

Satisfiability—SAT

Any boolean expression can be written in conjunctive normal form (CNF), e.g.

$$S = (x \vee y \vee \bar{z}) \wedge (\bar{x} \vee y \vee z) \wedge (\bar{x} \vee \bar{y} \vee \bar{z}),$$

or more readably written as

$$S = (x + y + \bar{z}).(\bar{x} + y + z).(\bar{x} + \bar{y} + \bar{z}).$$

S is said to be **satisfiable** if an assignment of 0s and 1s can be found for each of the variables x , y and z that makes $S = 1$.

The assignment $x = 0$, $y = 1$, and $z = 0$ makes $S = 1$.

Once we have an assignment it is easy to evaluate the boolean expression in polynomial time to show that it is 1.

A Turing machine and all its operations can be described as a boolean expression.

Satisfiability—the SAT problem is NP-complete

Cook's theorem

The SAT problem is NP-complete.

Proof

A Turing machine and all its operations can be described as a boolean expression.

S can be described by TM $M \iff$ the TM will terminate in its accept state for a given input.

This is not too easy to do, *but* constructing a boolean expression for any TM is no more than complicated that a polynomial number of steps the TM takes. Thus any NP algorithm can be described by an instance of a SAT problem.

Polynomial reduction—use on new problem

Now we can understand the common fate property.

Every NP-complete problem is reducible to every other one.

Thus the tractability of any one leads to the tractability of all the others.

To demonstrate that a new problem N is indeed NP-complete we must only establish a polynomial-time reduction to any single problem R that we know to be NP-complete, and we must find another NP-complete problem Q that can be polynomially reduced to N : so N cannot be any better than R and it cannot be any worse than Q .

$Q \in \mathbf{NP}$ is reducible to N , and N is reducible to $R \in \mathbf{NP}$ establishes the NP-completeness of N , i.e. $N \in \mathbf{NP}$.

Polynomial reduction— another general property

- If L_1 is polynomially reducible to L_2 and L_2 is polynomially reducible to L_3 then L_1 is polynomially reducible to L_3 .
- This is easy to achieve. Compose a converter from L_1 to L_2 and one from L_2 to L_3 . Since we use two polynomial reductions their composition is still polynomial, resulting in a polynomial-time conversion algorithm.
- When a problem cannot be be solved efficiently we try to reduce it to a known problem in **NP**.

NP-complete problems in one sentence

- There are classes of problems whose solutions can be tested in polynomial time, but they appear to possess only exact solutions that are exponential, e.g. SAT, travelling salesman, bin packing, LCS, timetabling, Hamilton cycle, vertex cover, job-shop scheduling etc.
- The book by M.R. Garey and David S. Johnson lists many NP-complete problems.
- _____

Intractable problems

- There are intractable problems—Hanoi.
- The towers of Hanoi problem takes $2^n - 1$ moves to complete.
- An n digit decimal odometer takes 10^n clicks to reset.
- An n -digit base- n odometer takes n^n clicks to reset.
- We use this idea to construct an $O(n^{n^n})$ odometer.

The base- n odometer problem is $O(n^n)$

The odometer is a counter using an n -tuple of numbers, $(d_0, d_1, \dots, d_n)$ where each $d_i \in [0..n)$, Initially the odometer is set to $(0, 0, \dots, 0)$

At each step 1 is added to the odometer using:

```
1 def inc (d, i): # d = [0]*n and i is an integer
2     n = len(d)
3     if i < n:
4         if d[i] < n - 1:
5             d[i] += 1; return 0
6         else:
7             d[i] = 0; return inc(d, i + 1)
8     else: return 1
```

```
1 def odometer(n):
2     d = [0]*n
3     result = 0
4     while result != 1:
5         result = inc (d, 0)
```

When $n = N^N$ the odometer is an $O(N^{N^N})$ method.

Useful mathematical facts—Logarithms

We write $\log$ for $\log_2$, and $\ln$ for $\log_e$ —some authors use $\lg$ for $\log_2$

$$\text{If } x = b^n, \text{ then } \log_b x = n.$$

$$\log_b 1 = 0, \log_b b = 1.$$

$$\log x \times y = \log x + \log y,$$

$$\log x/y = \log x - \log y,$$

$$\log_b x^n = n \log_b x, \text{ therefore } \log_b b^n = n, \text{ and}$$

$$\log_b x = \frac{\log_{10} x}{\log_{10} b} = \frac{\log_e x}{\log_e b} = \frac{\log_2 x}{\log_2 b} = \frac{\log x}{\log b} = \Theta(\log x),$$

$$\log_2 x = \frac{\log_{10} x}{\log_{10} 2} \approx \frac{\log_{10} x}{0.30102000} \approx \frac{\log_{10} x}{0.3},$$

$$\log_{10} x \approx 0.3 \times \log_2 x.$$

$$\text{Since } \log_b y \cdot \log_b x = \log_b x \cdot \log_b y,$$

$$\log_b x^{\log_b y} = \log_b y^{\log_b x},$$

$$x^{\log_b y} = y^{\log_b x}.$$

Summing logarithms

$$\begin{aligned}\sum_{r=1}^n \lfloor \log_2 r \rfloor &= (n+1) \lfloor \log_2 n \rfloor - 2^{\lfloor \log_2 n \rfloor + 1} + 2 \\ &= \Theta(n \log n).\end{aligned}$$

Sum of functions bounded by integral

If $f(x)$ is a monotonically increasing continuous function,

$$\sum_{r=1}^n f(r) \leq \int_1^{n+1} f(x) dx.$$

Useful mathematical facts—Bits and digits

The formula, $\log_2 x = \frac{\log_{10} x}{\log_{10} 2} \approx \frac{\log_{10} x}{0.30102000} \approx \frac{\log_{10} x}{0.3}$, shows how to calculate the number of bits x requires given $\log_{10} x$ —its number of digits.

The formula $\log_{10} x \approx 0.3 \times \log_2 x$ shows us how to get the number of digits represented by a given number of bits— $\log_2 x$ —of a number x .

e.g. The largest unsigned 24-bit number is $2^{24} - 1$,

$$0.3 \log_2(2^{24} - 1) \approx 0.3 \times 24 = 7.2 \text{ digits.}$$

So 24 bits accurately represent 7 digits.

On the other hand, by the same argument a 23 bit number can only represent numbers of $\lfloor 6.9 \rfloor = 6$ digits accurately.

This is where the *hidden 1 bit* of the 23-bit mantissa of the IEEE 754 Floating Point Standard is used to turn a 23 mantissa into a 24 bit mantissa by adding a hidden bit that is always 1 excepting when the number is zero.

Useful mathematical facts—Asymptotic analysis

Since we are making approximations of the time that algorithms take, we consider non-negative increasing functions.

Big- O notation— $T(n) = O(g(n))$ —“ $T(n)$ is of order $g(n)$ ”

$\iff$ there is a constant c such that $T(n) \leq cg(n)$, when n is big enough.
 $T(n)$ is asymptotically bounded from *above* by $g(n)$.

Omega Ω notation— $T(n) = \Omega(g(n))$ —“ $T(n)$ is at least of order $g(n)$ ”

$\iff$ there is a constant c such that $T(n) \geq cg(n)$, when n is big enough.
 $T(n)$ is asymptotically bounded from *below* by $g(n)$.

Theta Θ notation— $T(n) = \Theta(g(n))$ —“ $T(n)$ is of order exactly $g(n)$ ”

$\iff$ both $T(n) = O(g(n))$ and $T(n) = \Omega(g(n))$ hold.
 $T(n)$ is asymptotically *tightly* bounded by $g(n)$.

Useful mathematical facts—Asymptotic analysis, examples

e.g. Let $T(n) = 64n^3 + 32n^2 + 16n + 8$

$T(n)$ is $O(n^3)$, $O(n^4)$, $\dots O(n^{100})$, $\dots$, $\Omega(n^3)$, $\Omega(n^2)$, $\Omega(n)$, and $\Theta(n^3)$.

$T(n)$ is not $O(n)$, $\Omega(n^4)$, $\Theta(n)$, $\Theta(n^2)$.

Although $T(n)$ is $O(n^4)$ it is not $\Omega(n^4)$. So it is not $\Theta(n^4)$.

Small o notation

We write $f(n) = o(g(n))$ when $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0$.

It means that the function $f(n)$ is an order smaller than $g(n)$.

This implies that if $g(n) = o(T(n))$ then $O(T(n) + g(n)) = O(T(n))$,

e.g. $O(n^n + n^2 + n \log n + n + 10000000000) = O(n^n)$.

Asymptotically equal.

We write $f(n) \sim g(n)$ —“ $f(n)$ is asymptotically equal to $g(n)$ ”
 $\iff \lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 1$.

L'Hospital's Rule (1696)—was discovered by **Jacob Bernoulli** (1654–1705) but named after **Guilliaume L'Hospital** (1661–1704).

For functions f and g ,
if both $\lim_{x \rightarrow a} f(x) = 0$ and $\lim_{x \rightarrow a} g(x) = 0$,
or both $\lim_{x \rightarrow a} f(x) = \pm\infty$ and $\lim_{x \rightarrow a} g(x) = \pm\infty$,
and $\lim_{x \rightarrow a} \frac{f'(x)}{g'(x)}$ exists,
then $\lim_{x \rightarrow a} \frac{f(x)}{g(x)} = \lim_{x \rightarrow a} \frac{f'(x)}{g'(x)}$.

Useful mathematical facts— $n!$

How do $n!$ and n^n compare?

We have $n! = 1.2.3 \dots n < n.n.n \dots n = n^n$, such that $n!$ is clearly bounded from above by n^n , i.e. $n! = O(n^n)$ and $n^n = \Omega(n!)$.

James Stirling (1692–1770) gave the following approximation

$\lim_{n \rightarrow \infty} \frac{n!}{\sqrt{2\pi n} \left(\frac{n}{e}\right)^n} = 1$, or $n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n (1 + O(1/n))$, i.e.

$$\begin{aligned} \ln n! &= \ln \left[\sqrt{2\pi n} \left(\frac{n}{e}\right)^n (1 + O(1/n)) \right] \\ &= \ln \sqrt{2\pi} + 0.5 \ln n + n \ln n - n \ln e + \varepsilon \\ &= \Theta(n \ln n) \end{aligned}$$

R.W. Gosper, in 1975, gives a better approximation of $n!$ which yields reasonable values of $n!$ for all n with a small relative error, $\sqrt{2n + \frac{1}{3}\pi} n^n e^{-n}$

Summation

$$\sum_{r=1}^n r = \frac{n(n+1)}{2} = \Theta(n^2),$$

$$\sum_{r=1}^n r^2 = \frac{n(n+1)(2n+1)}{6} = \Theta(n^3),$$

$$\sum_{r=1}^n r^3 = \frac{n^2(n+1)^2}{4} = \Theta(n^4),$$

$$\sum_{r=1}^n r^k = \Theta(n^{k+1}).$$

$$\sum_{r=0}^n (b-1)b^r = b^{n+1} - 1,$$

e.g. $b = 10$, $n = 7$, $\text{sum} = 99999999 = 10^8 - 1 \approx 10^8$.

$$\sum_{r=-1}^n (b-1)b^r,$$

e.g. $b = 10$, $n = 7$, $\text{sum} = 99999999.9 = 10^8 - 0.1 \approx 10^8$.

$$\sum_{r=0}^n 2^{-r} = 1.111\dots 11_2 = 2 - 2^{-n},$$

$$\sum_{r=0}^n 2^r = 111\dots 11_2 = 2^{n+1} - 1,$$

$$\sum_{r=1}^n r2^r = 2 + (n-1)2^{n+1},$$

$$\sum_{r=1}^n r2^{-r} < 2.$$

COS 212 Examination Paper II
Tuesday 19th November at 8h30
in S1

Note:

I do not take any responsibility for informing you of venue or timetable changes.

Preparation for COS 212 Examination

- Understand and be able to describe and replicate parts of all practicals done during the year.
- Study the class notes and other handouts carefully.
- Understand tutorials, quizzes, homework.
- Be prepared for repeat questions out of old exam papers and class test papers—please note that I do not supply these. Beware that when old questions are repeated they invariably have subtle changes.
- Although this course is based on all your preceding courses in Computer Science, concentrate on this semester's work.
- Make a special effort to study and understand the chapter on analysis of algorithms—Chapter 4 in Goodrich and Tamassia 4th edition, or Chapter 3 of the 3rd edition. The same content is available in Section 1.4 of Sedgewick and Wayne, Pages 172–215.

Preparation for COS 212 Examination

- When answering questions in the examination

Read the questions properly before answering.

Students very often answer questions that they imagine have been asked.

Memo for examination

- Linear search—be able to give or describe the algorithm and its timing. Suppose 85% of the search takes place in the first 5% of the list, what is the average search time?

Can this linear search be fine tuned?

- Binary search—be able to give or describe algorithm and timing: e.g. How long does it take to find an element not in the list?
How many comparisons does it take to show that a given element IS in the list.
How many probes—comparisons—does it take to show that a given element IS NOT in the list.

- Heaps / priority queues.
 - Be able to explain the time complexity of building a heap top-down and building a heap bottom-up and code them.
 - Explain sorting with heap and explain its space and time complexity.
 - Be able to program sorting using a heap.
 - What algorithms do you know that use heaps to speed them up? Give a list.
 - Find k th smallest element. Will you use a heap or Hoare's partition algorithm? You should be able to program this problem using both methods.

Memo for examination

- Quicksort: partition, timing was dealt with in detail in class notes. Merge sort. Timing. Find the k -th smallest/largest element in a list without sorting—hint: use partition, or use a priority list built bottom-up.
- Boyer-Moore, Knuth-Morris-Pratt, Rabin-Karp versus brute force. Know all their timings. Know details of algorithms. Be able to fill in gaps in given programs.
- Tries
- Huffman encoding and decoding. How does Vayner encryption—based on Huffman coding—work. Speed of Huffman using heap versus not using a heap.
- Be able to describe efficient spelling checkers given constraints such as: “Using a preprocessed dictionary the checking algorithm must be independent of the size of the dictionary.” or “The program may not use a method that relies on hashing.” or “The dictionary and data must be processed at the same time.”

Memo for examination

- Longest common subsequence. What is dynamic programming? What other methods have you seen that use dynamic programming.
- Graphs: definitions. Depth-first search (DFS) for trees. Breadth-first search (BFS) for trees. DFS and BFS for graph traversal. Dijkstra's shortest path algorithm, Kruskal's algorithm and Prim-Jarník algorithm for minimal-spanning trees. Timings of each—you must be able to explain the timings.
- For graphs also be able to explain the datastructures to store them: edge list, node list, adjacency list, etc.
- Topological sorting.
- Be able to demonstrate all algorithms, given a small set of data.

Memo for examination

- Give and explain an $O(2^n)$, an $O(n^n)$, and an $O(n^{n^n})$ algorithm.
- What is $\log_b$ of any number given as a base b number in terms of its digits?
- What is the biggest unsigned/signed number that can be represented using 16, 32, 64 or 128 bits— give the answer without using a calculator.
- How long will the Towers of Hanoi run using 64 disks. Estimate without using a calculator. Be able to derive the time complexity $T(N) = 2^N - 1$ for the Towers of Hanoi.
- Understand and be able to derive and explain $T(N) = 2T(N/2) + cN = O(N \log N)$ —the time taken by divide and conquer algorithms.

Memo for examination—Computability and NP complexity

- Be able to prove that the halting problem is non-computable.
- Define totality, equivalence.
- What is meant by partial computability?
- What are proof systems?
- Give a class of algorithms that are considered efficient?
- When do we consider an algorithm intractable?
- What is meant by **NP** completeness?
- Understand and be able to explain what the practical implications are of knowing that an algorithm is **NP** complete.
- Give a list of at least 5 problems that are **NP** complete.
- Be able to explain how to go about proving that a given problem is **NP** complete.